

Ch 5. SQL: Queries, Constraints, Triggers

Sang-Won Lee

<http://icc.skku.ac.kr/~swlee>

SKKU VLDB Lab.

(<http://vldb.skku.ac.kr/>)

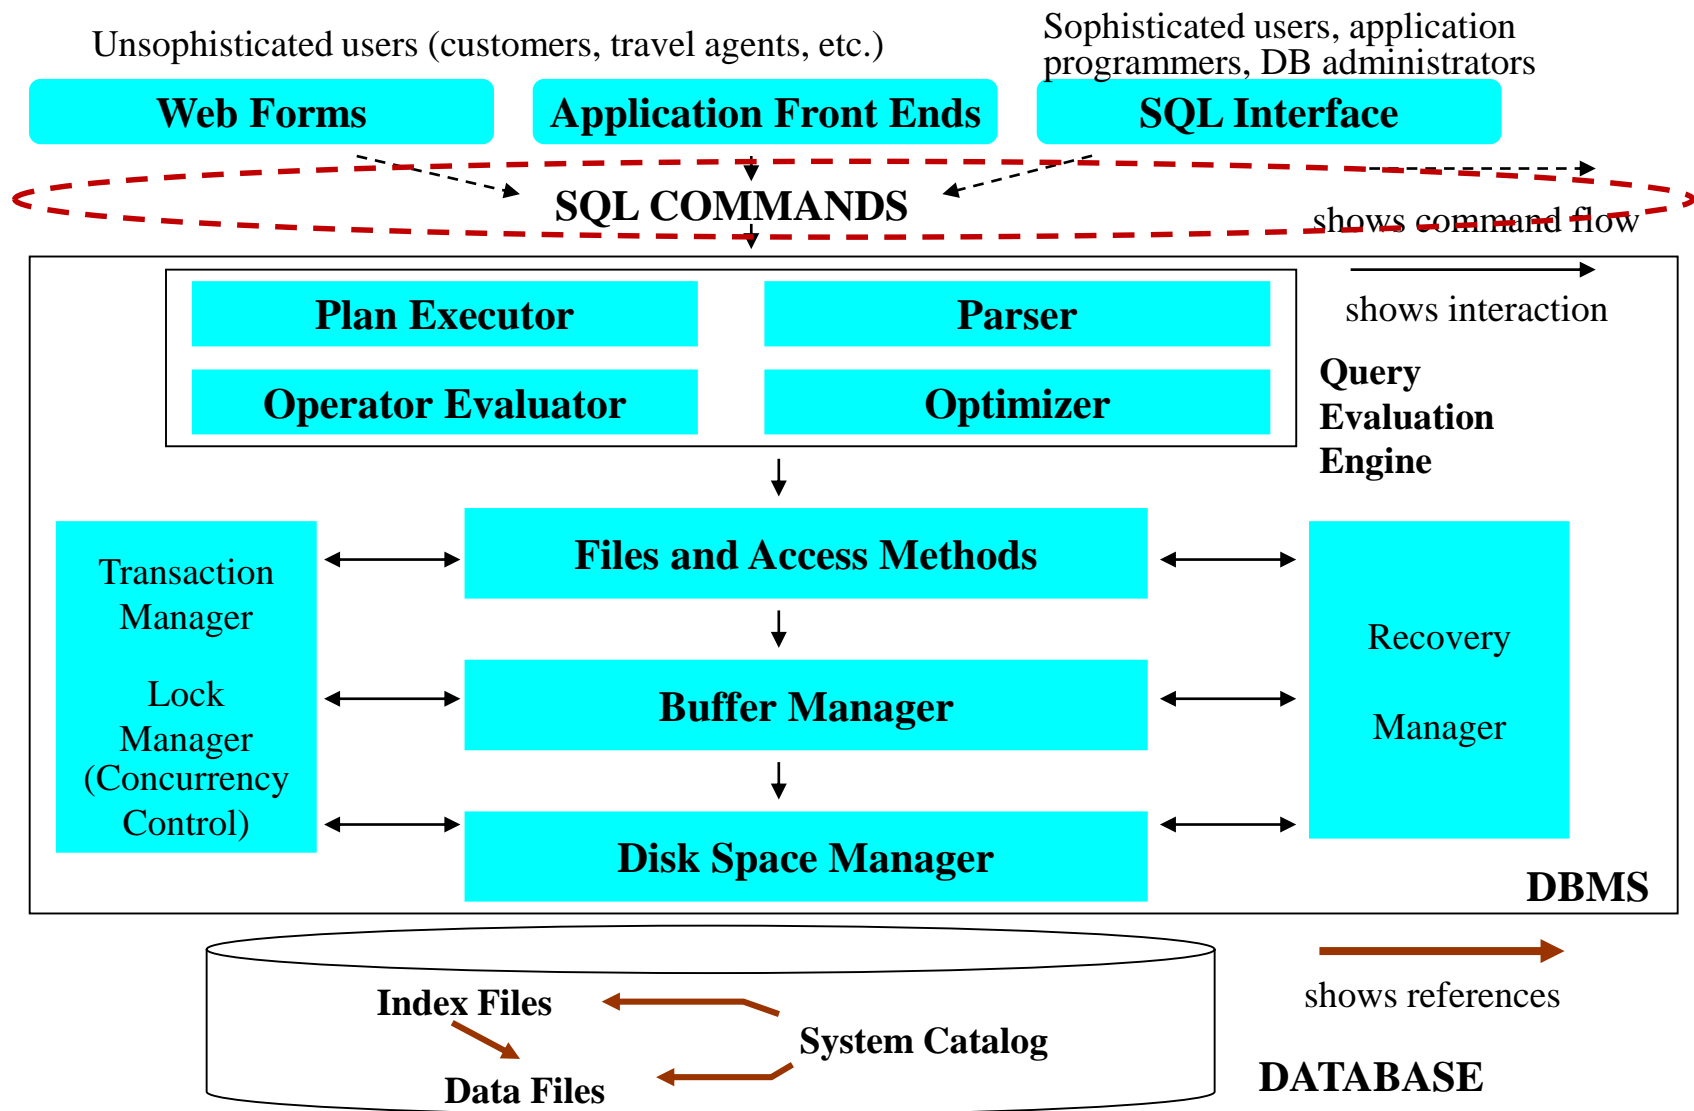
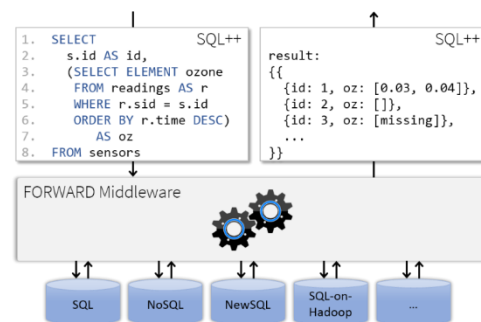


Figure 1.3 Anatomy of an RDBMS

SQL in Big Data Era

- What is SQL?
 - DB dead? (5 Years ago)
- The BIG data era has come
- Now everyone sells SQL
 - SQL-on-{Hadoop, NoSQL}
 - Google's cloud Spanner, BigQuery
 - Amazon's RDS, AuroraDB
 - Databricks's DeltaLake, Snowflake
- SQL becomes the lingua franca of the data management world in big data era?!



Overview

- DML: queries + insert/update/delete
- DDL(Data Definition Language)
 - Logical: table/view/ICs
 - Physical: index/partitions
- DCL(Data Control Language)
 - Trigger, Advanced ICs, Transaction Management, Security
- Misc.
 - Embedded/Dynamic SQL, Client-Server/Remote DB Access (e.g. ODBC/JDBC)
 - Advanced Features: OO, Logics(recursive queries), DSS & OLAP, Data Mining, Spatial/Temporal DB, Text & XML Data, Game, Social Network

** Refer to 5.1 for details


SQL Expressive Powers

1. Relational Algebra or Calculus
2. Aggregation / Grouping
3. Deductive Logics / Analytic Functions (Windowing)
4. Data Mining Features
5. **Machine Learning (e.g. + Linear Algebra?)**

SQL and Data Independence



- “Queries should be expressed in terms of high-level, nonprocedural concepts that are independent of physical representation. Selection of an algorithm for processing a given query could then be done by an optimizing compiler, based on the access paths available and the statistics of the stored data; if these access paths or statistics should later change, the algorithm could be re-optimized without human intervention.” ([SQL @ Encyclopedia of Database Systems](#))
- Design considerations of relational data model (and RDBMS) for data independence
 - 1) High-level language (declarative, non-procedural SQL), 2) value-based relationship, 3) content(or name)-based addressing (e.g. unlike (x,y) in Excel or C-like pointer, detpno =10, column C instead of in 3rd column), and 4) no assumptions about physical schema such as data layout (row-wise vs. column-wise) , existence of indexes and table partitioning
- SQL’s Benefits
 - Database application **development productivity**
 - Data Independence



5.2 Example Schemas

Sailors

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

M:N relationship
b/w Sailors and Boats

- Tables and queries from <http://www.cs.wisc.edu/~dbbook>

1. Relational Algebra or Calculus
2. Aggregation / Grouping
3. Deductive Logics / Analytic Functions (Windowing)
4. Data Mining Features

Basic SQL Query

```
SELECT [DISTINCT] target-list  
FROM relation-list  
WHERE qualification
```

- relation-list A list of relation names, possibly with a *range (or tuple) variable* after each name.
- target-list A list of attributes of relations in *relation-list*
- qualification Comparisons (attr op const or attr1 op attr2, where *op* is one of $<$, $>$, $=$, \leq , \geq , \neq) combined using logical connectives **AND**, **OR**, and **NOT**.
- **DISTINCT** is an optional keyword indicating that the answer should not contain duplicates.
 - **Note**: Set semantics in R.A. vs. multi-set(or bag) semantics in SQL
 - By default, duplicates are **NOT** eliminated in SQL.

Single Table Query

```
SELECT S.sid, S.sname, S.rating, S.age
FROM Sailors AS S /* S: tuple variable */
WHERE S.rating > 7
```

```
SELECT S.name, S.age
FROM Sailors (AS) S
```

```
SELECT DISTINCT S.name, S.age
FROM Sailors (AS) S
```

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5

sname	age
Dustin	45.0
Brutus	33.0
Lubber	55.5
Andy	25.5
Rusty	35.0
Horatio	35.0
Zorba	16.0
Art	25.5
Bob	63.5

sname	age
Dustin	45.0
Brutus	33.0
Lubber	55.5
Andy	25.5
Rusty	35.0
Horatio	35.0
Zorba	16.0
Horatio	35.0
Art	25.5
Bob	63.5

Conceptual Evaluation Strategy

- Semantics of an SQL query defined in terms of the following conceptual evaluation strategy:
 1. Compute **cross-product** of *relation-list*. (what if any relation is empty?)
 2. Discard resulting tuples if they fail *qualifications* (i.e., **NOT TRUE**).
 3. Delete attributes that are not in *target-list*.
 4. If DISTINCT is specified, eliminate **duplicate** rows.
- This strategy is probably the **least efficient** way to compute a query! An **optimizer** will find more efficient strategies to compute *the same answers*.

Example of Conceptual Evaluation

Q1: Find the names of sailors who reserved boat 103

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND R.bid=103
```

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/96
58	103	11/12/96

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

1. FROM

2. WHERE

3. SELECT

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>	<i>sid</i>	<i>bid</i>	<i>day</i>
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

A Note on Range Variables (or Tuple Variables)

- Really needed **only if** the same relation appears twice in the FROM clause (that is, self-join). The previous query can also be written as:

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND bid=103
```

OR

```
SELECT sname
FROM   Sailors, Reserves
WHERE  Sailors.sid=Reserves.sid
      AND bid=103
```

Tuple Variables

Sometimes we need to refer to two or more tuples in the same relation. To do so, we define several tuple variables for that relation in the from clause and use the tuple variables as aliases of the relation. The effect is exactly the same as was achieved by the range-statement in QUEL, so SQL, which appeared at first to be a “syntactically sugared” form of relational algebra, is now revealed to resemble tuple relational calculus.

- It is good style, however, to use **range (or tuple) variable** always!

Find the name of sailor who reserved boat 103

- SQL (by user)

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND bid=103
```

- Tuple Relational Calculus (TRC)

$$\{P \mid \exists S \in \text{Sailors} \ \exists R \in \text{Reserves} (R.\text{sid} = S.\text{sid} \wedge R.\text{bid} = 103 \wedge P.\text{sname} = S.\text{sname})\}$$

- Relational Algebra: 2 different expressions

1. DBMS translate SQL into relational algebra

$$\pi_{\text{sname}}(\sigma_{\text{bid}=103}(\text{Reserves} \bowtie \text{Sailors}))$$

2. Query optimizer (QO) looks for other algebra expressions that produce the same result
3. QO will find the second expression is likely to be less expensive because the sizes of intermediate relations are smaller, thanks to the early use of selection

$$\pi_{\text{sname}}((\sigma_{\text{bid}=103} \text{Reserves}) \bowtie \text{Sailors})$$

Q4: Find sailors who've reserved **at least one** boat

```
SELECT S.sid  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid
```

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/96
58	103	11/12/96

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

- Join can be interpreted as **existential quantifier**
- Would adding **DISTINCT** to this query make a difference?
- What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause? Would adding DISTINCT to this variant of the query make a difference?

5.2.2 Expressions and Strings

```
SELECT S.age, age1=S.age-5, 2*S.age AS age2  
FROM Sailors S  
WHERE S.sname LIKE 'B_%B'
```

- Illustrates use of arithmetic expressions and string pattern matching: *Find triples (of ages of sailors and two fields defined by expressions) for sailors whose names begin and end with B and contain at least three characters.*
- AS and = are two ways to name fields in result.
- LIKE is used for string matching. `'_'` stands for any one character and `'%'` stands for 0 or more arbitrary characters.
- Date and time format: e.g. Oracle (<http://infolab.stanford.edu/~ullman/fcdb/oracle/or-time.html>)

Regular Expressions in SQL

```
SELECT zip  
FROM zipcode  
WHERE REGEXP_LIKE(zip, '^[[:digit:]]')
```

ZIP

ab123

123xy

007ab

abcxy

- The increased importance of text data → Regular Expression

- See [here](#) for Regular Expression in Oracle 10g

✓ Here: http://docs.oracle.com/cd/B19306_01/appdev.102/b14251/adfns_regexp.htm

5.3 Union, Intersect, Except

- Another basic operators in R.A
 - UNION, DIFFERENCE
- ANSI SQL provides the following set operators
 - UNION, UNION ALL
 - INTERSECT, INTERSECT ALL
 - EXCEPT(or MINUS)
 - IN, ANY, ALL, EXISTS (covered in Section 5.4)

Q5: Find sid's of sailors who've reserved a red or a green boat

- UNION: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
- If we replace OR by AND in the first version, what do we get?
- Also available: EXCEPT (What do we get if we replace UNION by EXCEPT?)

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND (B.color='red' OR B.color='green')
```

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color='red'
UNION
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color='green'
```

Q6: Find sid's of sailors who've reserved a red and a green boat

- INTERSECT: Can be used to compute the intersection of any two *union-compatible* sets of tuples.
- Included in the SQL/92 standard, but some systems don't support it.
- Contrast symmetry of the UNION and INTERSECT queries with how much the other versions differ.

```
SELECT S.sid
FROM Sailors S, Boats B1, Reserves R1,
      Boats B2, Reserves R2
WHERE S.sid=R1.sid AND R1.bid=B1.bid
      AND S.sid=R2.sid AND R2.bid=B2.bid
      AND (B1.color='red' AND B2.color='green')
```

Key field!

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color='red'
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
      AND B.color='green'
```

5.4 Nested Queries(or Subqueries)

Q1: Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND bid=103
```



```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN (SELECT R.sid
                  FROM   Reserves R
                  WHERE  R.bid=103)
```

- A very powerful feature of SQL
 - a **WHERE** clause can itself contain an SQL query!
 - Also, **SELECT**, **FROM**, or **HAVING** clause can.
- To find sailors who've *not* reserved #103, use **NOT IN**.
- To understand semantics of nested queries, think of a ***nested loops*** evaluation: *For each Sailors tuple, check the qualification by computing the subquery.*
- Do the two queries above have **the same** semantic?
 - What if one sailor reserved two or more boats?

Multiply Nested Queries

Q2: Find names of sailors who have *(not)* reserved a *red* boat:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid (NOT) IN (SELECT R.sid
                      FROM Reserves R
                      WHERE R.bid IN ( SELECT B.bid
                                       FROM Boats B
                                       WHERE B.color = 'red'))
```

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5


<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
64	101	9/5/98
64	102	9/8/98
74	103	9/8/98

<i>bid</i>	<i>bname</i>	<i>color</i>
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Nested Queries with Correlation

Q1: Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
               FROM Reserves R
               WHERE R.bid=103 AND S.sid=R.sid)
```



- **EXISTS** is another set comparison operator, like IN.
- **Correlated vs. non-correlated**: illustrates why, in general, subquery must be re-computed for each Sailors tuple.
- If **UNIQUE** is used, and * is replaced by *R.bid*, finds sailors with at most one reservation for boat #103. (UNIQUE checks for duplicate tuples; * denotes all attributes. **Why** do we have to replace * by *R.bid*? **Because** “unique(select * ..” will return **True** even though more than one reservation for boat 103 exists!)

More on Set-Comparison Operators

- We've seen IN, EXISTS and UNIQUE. Also, NOT IN, NOT EXISTS and NOT UNIQUE.
- Also: *op* **ANY(SOME)**, *op* **ALL**
 - *op*: $>, <, =, \geq, \leq, \neq$
 - IN equivalent to =ANY; NOT IN equivalent to \neq ALL
- e.g. Find sailors whose rating is greater than that of some sailors called Horatio:

```
SELECT *  
FROM Sailors S  
WHERE S.rating > ANY (SELECT S2.rating  
                        FROM Sailors S2  
                        WHERE S2.sname='Horatio')
```

- What if inner query is empty?
False
- **ANY** \rightarrow **ALL**: **True**

Rewriting **INTERSECT** Queries Using **IN**

*Q6: Find sid's of sailors who've reserved both a red **and** a green boat:*

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
AND S.sid IN (SELECT S2.sid
              FROM Sailors S2, Boats B2, Reserves R2
              WHERE S2.sid=R2.sid AND R2.bid=B2.bid
                AND B2.color='green')
```

Who reserved Red Boat

And, further who has sids that is in Green Boat Reservers list

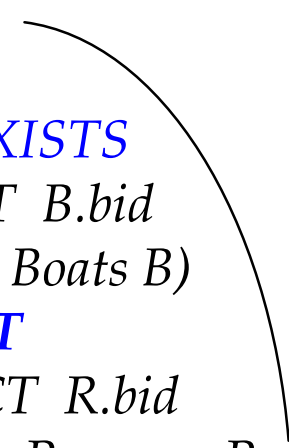
- Similarly, **EXCEPT** queries re-written using **NOT IN**.
- To find *names* (not *sid*'s) of Sailors who've reserved both red and green boats, just replace *S.sid* by *S.sname* in SELECT clause. (What about INTERSECT query?)

Division in SQL

(1)

Q9: Find sailors who've reserved *all boats*.

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
      ((SELECT B.bid
        FROM Boats B)
      EXCEPT
      (SELECT R.bid
        FROM Reserves R
        WHERE R.sid=S.sid))
```



- Let's do it the hard way, without **EXCEPT**:

(2)

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS (SELECT B.bid
                  FROM Boats B
```

Sailors S such that ..

```
                  WHERE NOT EXISTS (SELECT R.bid
```

there is no boat B without ..

```
                  FROM Reserves R
```

```
                  WHERE R.bid=B.bid
```

a Reserves tuple showing S reserved B

```
                  AND R.sid=S.sid))
```


5.5 Aggregate Operators

SQL Expressive Powers

1. Relational Algebra or Calculus
2. Aggregation / Grouping
3. Deductive Logics / Analytic Functions (Windowing)
4. Data Mining Features

- Significant extension of Relational Algebra

COUNT (*)
COUNT ([DISTINCT] A)
SUM ([DISTINCT] A)
AVG ([DISTINCT] A)
MAX (A)
MIN (A)

single column

```
SELECT COUNT (*)  
FROM Sailors S
```

```
SELECT AVG (S.age)  
FROM Sailors S  
WHERE S.rating=10
```

```
SELECT COUNT (DISTINCT S.rating)  
FROM Sailors S  
WHERE S.sname='Bob'
```

```
SELECT S.sname  
FROM Sailors S  
WHERE S.rating= (SELECT MAX(S2.rating)  
FROM Sailors S2)
```

```
SELECT AVG ( DISTINCT S.age)  
FROM Sailors S  
WHERE S.rating=10
```

Find name and age of the **oldest** sailor(s)

- The first query is **illegal**! (We'll look into the reason a bit later, when we discuss **GROUP BY**.)
- The third query is equivalent to the second query, and is allowed in the SQL/92 standard, but is **not supported in some systems** .

```
SELECT S.sname, MAX (S.age)
FROM Sailors S
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX (S2.age)
       FROM Sailors S2)
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX (S2.age)
       FROM Sailors S2)
      = S.age
```

Aggregate Operators: Any / All Alternatives

```
SELECT S.sname  
FROM Sailors S  
WHERE S.age > (SELECT MAX(S2.age)  
               FROM Sailors S2  
               WHERE S2.rating = 10)
```

```
SELECT S.sname  
FROM Sailors S  
WHERE S.age > ALL (SELECT S2.age  
                  FROM Sailors S2  
                  WHERE S2.rating = 10)
```

GROUP BY and HAVING

- So far, we've applied aggregate operators to all (qualifying) tuples. In many real-world scenarios, however, we want to apply them to each of several **groups** of tuples.
- Q31: Find the age of the youngest sailor *for each* rating level.
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (!):

```
For i = 1 .. 10  
    SELECT MIN (S.age)  
    FROM Sailors S  
    WHERE S.rating = i
```

- In general, we don't know **how many rating levels** exist, and **what the rating values** for these levels are!

Group By: Examples

- For each rating, find the average age of the sailors

```
SELECT S.rating, AVG (S.age)  
FROM Sailors S  
GROUP BY S.rating
```

- For each rating, find the age of the youngest sailor with age ≥ 18

```
SELECT S.rating, MIN (S.age)  
FROM Sailors S  
WHERE S.age  $\geq$  18  
GROUP BY S.rating
```

Queries With GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

- The ***target-list*** contains (i) attribute names (i.e. attribute list) (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
 - The attribute list (i) **MUST BE** a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group.
 - A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.

Conceptual Evaluation

1. The cross-product of *relation-list* is computed, tuples that fail *qualification* are discarded,
2. 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in grouping-list.
3. The *group-qualification* is then applied to eliminate some groups.(by HAVING clause)
4. One answer tuple per qualifying group is generated.

Example(1)

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
```

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
71	zorba	10	16.0
64	horatio	7	35.0
29	brutus	1	33.0
58	rusty	10	35.0

1. Form cross product

rating	age
1	33.0
7	45.0
7	35.0
8	55.5
10	35.0

2. Delete unneeded columns, rows;
form groups

3. Perform Aggregation



rating	age
1	33.0
7	35.0
8	55.0
10	35.0

Answer Table

Example(2) Q32: Find the age of the **youngest sailor with age ≥ 18 ,
for each rating with **at least 2 such** sailors**

```
SELECT S.rating, MIN (S.age)
FROM Sailors S
WHERE S.age  $\geq$  18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
9	35.0
3	25.5
3	63.5
3	25.5

1) After WHERE

rating	age
1	33.0
3	25.5
3	25.5
3	63.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0

2) After GROUP BY

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5
96	Frodo	3	25.5

Answer relation

rating	minage
3	25.5
7	35.0
8	25.5

3) After HAVING and AGG.

Q33: For each red boat, find # of reservations for this boat

```
SELECT B.bid, COUNT (*) AS scount
FROM Boats B, Reserves R
WHERE B.bid=R.bid AND B.color='red'
GROUP BY B.bid
```

What if a red boat has no reservation?

```
SELECT B.bid, COUNT (*) AS scount
FROM Boats B, Reserves R
WHERE R.bid=B.bid
GROUP BY B.bid
HAVING B.color = 'red'
```

ILLEGAL!! WHY??

Q34: Find the average age for each rating that has at least 2 sailors (of any age)

```
SELECT S.rating, AVG (S.age)  
FROM Sailors S  
GROUP BY S.rating  
HAVING COUNT(*) > 1
```

<i>rating</i>	<i>avgage</i>
3	44.5
7	40.0
8	40.5
10	25.5



<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	3	25.5
95	Bob	3	63.5
96	Frodo	3	25.5

**Q35: Find the average age of sailors with age > 18,
for each rating with at least 2 sailors (of any age)**

```
SELECT S.rating, AVG (S.age)  
FROM Sailors S  
WHERE S.age > 18  
GROUP BY S.rating  
HAVING 1 < (SELECT COUNT (*)  
              FROM Sailors S2  
              WHERE S.rating=S2.rating)
```

- Shows **HAVING clause** can also contain a **subquery**.
- Compare this with the query (**Q32**) where we considered only ratings with 2 sailors over 18!
 - **HAVING COUNT(*) >1**

Q37: Find those ratings for which the average age is the minimum over all ratings

- Aggregate operations **cannot be nested!**

```
SELECT S.rating
FROM Sailors S
WHERE AVG(S.age) =
      (SELECT MIN (AVG (S2.age)) FROM Sailors S2 GROUP BY S2.rating)
```

Cf. Oracle allows nested aggregation.

- Correct solution (in SQL/92 and Oracle):

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG (S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN (Temp.avgage)
                    FROM Temp)
```

```
WITH Temp AS ( -- Subquery factoring in Oracle
               SELECT S.rating, avg(S.sal) as avgsal
               FROM Sailors S
               GROUP BY S.rating )
SELECT Temp.rating, Temp.avgage
FROM Temp
WHERE Temp.avgsal = (SELECT MIN(Temp.avgage)
                    FROM Temp)
```

5.6 Null Values

- Field values in a tuple are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
 - SQL provides a special value *null* for such situations.
- The presence of *null* complicates **many issues**.
- Optional reading: [Nulls: nothing to worry about](#)

Null Values(2)

1. Special operators needed to check if value is/is not *null*.
 - IS NULL, IS NOT NULL
 - NVL(attr, const) function
2. Is *rating > 8* true or false when *rating IS NULL*? What is the result of AND, OR and NOT logical connectives involving NULL? We need a 3-valued logic (true, false and *unknown*; T, F, U)
 - If rating is null, rating > 8 is evaluated to “U”, not “T” or “F”
 - Null = Null → U
 - NOT (U) = U
 - T OR U = T; F OR U = U; F AND xx = F; U AND (T or U) = U

Null Values(3)

3. Meaning of constructs must be defined carefully.
- WHERE clause eliminates rows that don't evaluate to **true**. That is, **false and unknown** tuple is **discarded!!**
 - Duplicate semantics: duplicate if corresponding columns has nulls
 - +, -, *, / returns null if one of arguments is null
 - COUNT(*) counts null values; all other aggregate operations **simply ignore nulls!!**
 - ✓ AVG/SUM/MIN/MAX of nulls → null
 - ✓ Count(attr) also ignores null value

Null Values(4)

4. New operators (in particular, **outer joins**) possible/needed.

<i>sid</i>	<i>bid</i>	<i>day</i>
22	101	10/10/96
58	103	11/12/96

<i>sid</i>	<i>sname</i>	<i>rating</i>	<i>age</i>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

```
SELECT S.sid, R.bid
FROM Sailors S, Reserves R
WHERE S.sid = R.sid
```

<i>sid</i>	<i>bid</i>
22	101
31	null
58	103

INNER JOIN

```
SELECT S.sid, R.bid
FROM Sailors S, Reserves R
WHERE S.sid = R.sid(+)
```

<i>sid</i>	<i>bid</i>
22	101
31	null
58	103

OUTER JOIN

Null Values(5)

- ANSI join syntax

```
SELECT S.sid, R.bid  
FROM Sailors S, Reserves R  
WHERE S.sid = R.sid and S.age > 18 and R.date = '04/03/24'
```

↑
JOIN condition

↑
SELECTION condition

```
SELECT S.sid, R.bid  
FROM Sailors S NATURAL JOIN Reserves R  
WHERE S.age > 18 and R.date = '04/03/24'
```

- See [Oracle SQL Language Reference](#) for ANSI JOIN in Oracle

Null Values(6)

5. Disallowing NULL values

- sname CHAR(20) NOT NULL
- implicit NOT NULL in Primary key constraint!
- Column with UNIQUE constraint is **NULLable**; two or more tuples can have null values for the column

6. IN & NOT IN

- 1 in {1, null}: True, 2 in {1, null} : False
 - ✓ Thus, in == exists
- 1 not in {1, null}: False, 2 not in {1, null} : False
 - ✓ Thus, not in != not exists

5.7 Integrity Constraints (Review)

- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, **general constraints**.
 - *Domain constraints*: Field values must be of right type. Always enforced.
- In this section, we discuss the complex integrity constraints which utilize the full power of SQL queries

General Constraints

- Useful when more general ICs than keys are involved.
- Can use **queries** to express constraint.
- Constraints can be named.

```
CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10 )
```

Per-tuple Check

```
CREATE TABLE Reserves
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK ('Interlake' <>
        ( SELECT B.bname
          FROM Boats B
          WHERE B.bid=bid)))
```

Intra- or Inter-Table Check

5.7.3 Assertions: ICs over Multiple Relations

- Awkward and wrong!
- If **Sailors is empty**, the number of Boats tuples can be anything! – **why?**
- **Assertion** is the right solution; not associated with either table.

```
CREATE TABLE Sailors
```

```
( sid INTEGER,  
  sname CHAR(10),
```

```
  rating INTEGER,
```

```
  age REAL,
```

```
  PRIMARY KEY (sid),
```

```
  CHECK
```

```
    ( (SELECT COUNT (S.sid) FROM Sailors S)
```

```
      + (SELECT COUNT (B.bid) FROM Boats B) < 100 )
```

(# of boats + # of Sailors)
< 100

```
CREATE ASSERTION smallClub
```

```
  CHECK
```

```
    ( (SELECT COUNT (S.sid) FROM Sailors S)
```

```
      + (SELECT COUNT (B.bid) FROM Boats B) < 100 )
```

5.8 Triggers and Active Database

- Trigger: procedure that starts **automatically** if specified changes occur to the DBMS → **ACTIVE** Database!!
- Three parts: **ECA** Rule
 - **E**vent (activates the trigger)
 - **C**ondition (tests whether the triggers should run)
 - **A**ction (what happens if the trigger runs)

```
CREATE TRIGGER youngSailorUpdate
    AFTER INSERT ON SAILORS
    REFERENCEING NEW TABLE NewSailors
    FOR EACH STATEMENT
        INSERT INTO YoungSailors(sid, name, age, rating)
            SELECT sid, name, age, rating
            FROM NewSailors N
            WHERE N.age <= 18
```

Triggers

- Read Sec. 5.9.3
- See [here](#) for trigger example in Oracle 10g
 - Provide sophisticated auditing
 - Prevent invalid transactions
 - Enforce referential integrity (either those actions not supported by declarative integrity constraints or across nodes in a distributed database)
 - Enforce complex business rules
 - Enforce complex security authorizations
 - Provide transparent event logging
 -

Summary

- SQL was an important factor in the early acceptance of the relational model; more natural than earlier, procedural query languages, **more productive!!**.
- Relationally complete; in fact, significantly more expressive power than relational algebra.
- Even queries that can be expressed in RA can often be expressed more naturally in SQL.
- **Many alternative ways** to write a query; optimizer should look for most efficient evaluation plan.
 - In practice, users need to be aware of how queries are optimized and evaluated for best results.
- NULL for unknown field values brings many complications
- SQL allows specification of rich integrity constraints
- Triggers respond to changes in the database