

# Ch 9. Storing Data: Disks and Files

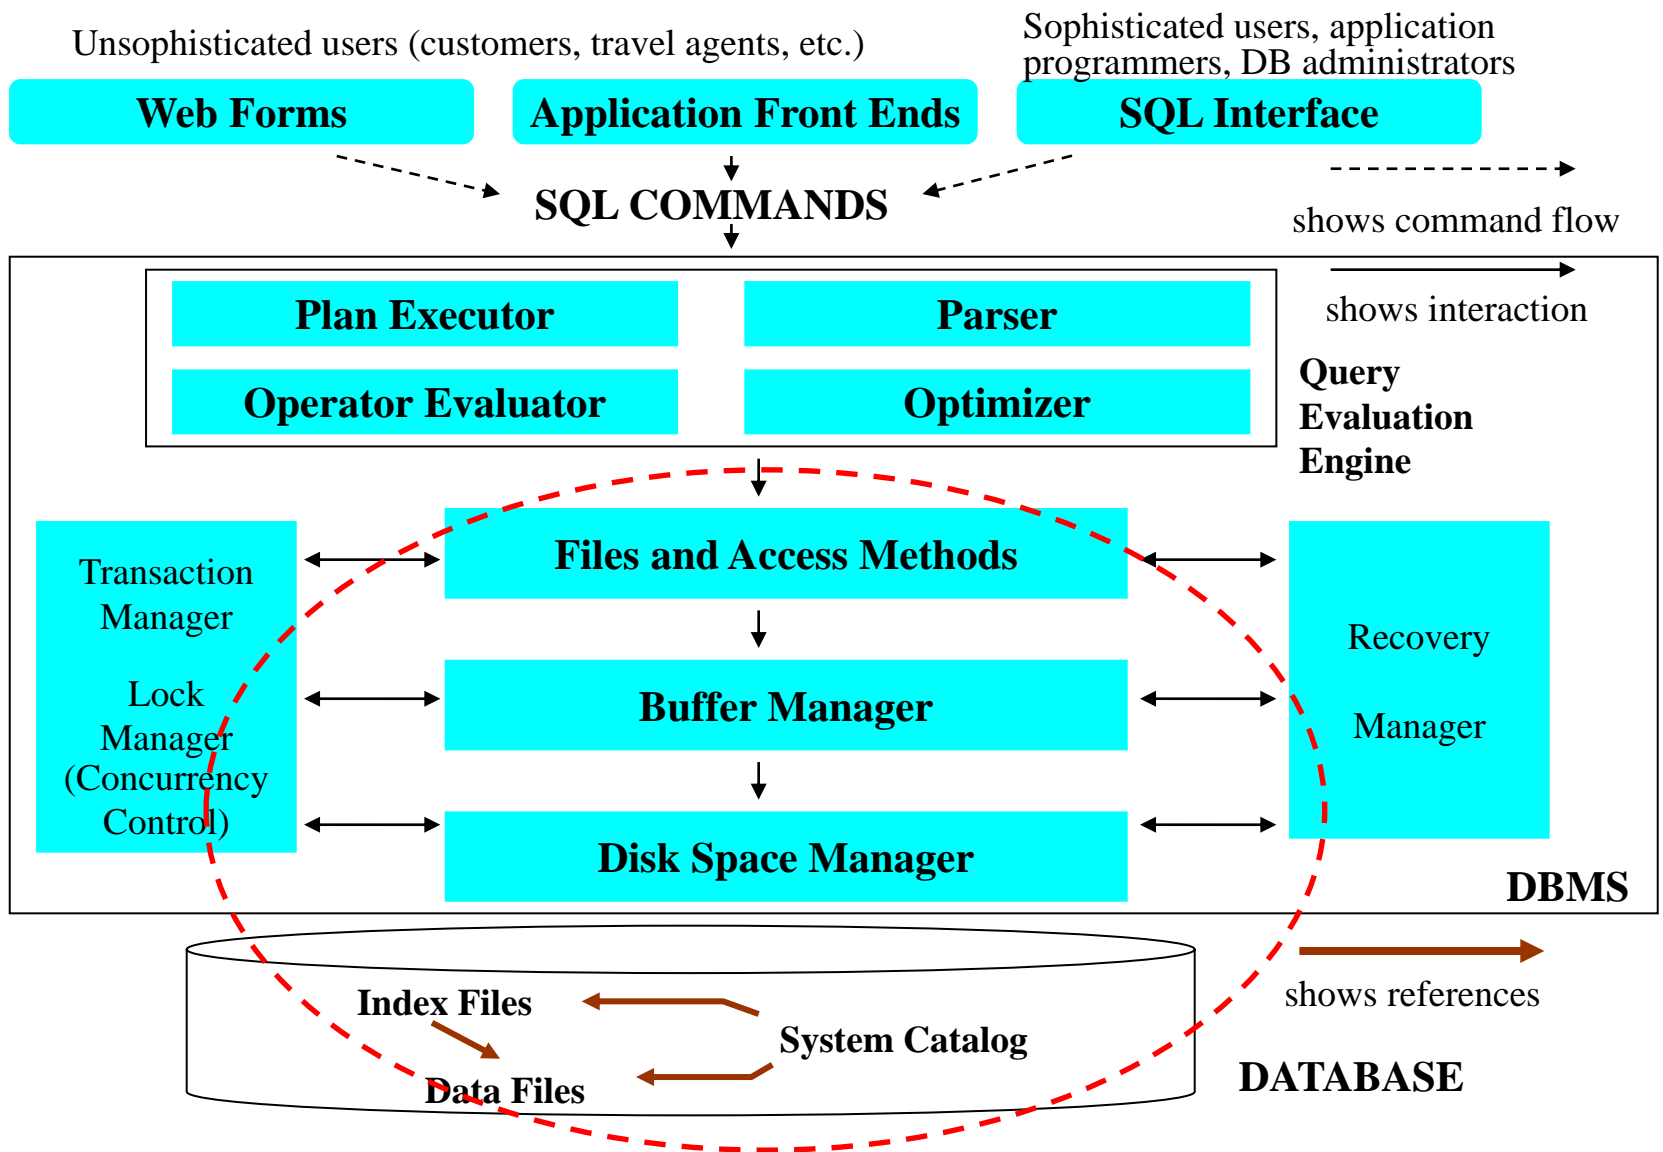
## - Heap File Structure -

Sang-Won Lee

<http://icc.skku.ac.kr/~swlee>

SKKU VLDB Lab.

( <http://vldb.skku.ac.kr/> )



**Figure 1.3 Anatomy of an RDBMS**

# Contents

9.0 Overview

9.1 Memory Hierarchy

9.2 RAID(Redundant Array of Independent Disk)

9.3 Disk Space Management

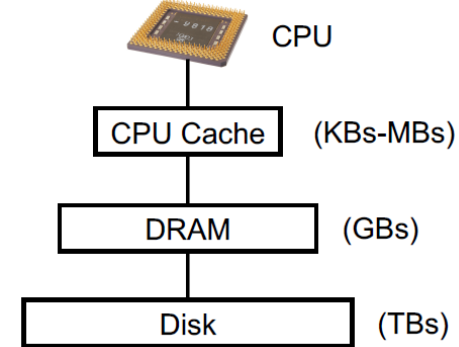
9.4 Buffer Manager

9.5 Files of Records

9.6 Page Format

9.7 Record Format

# 9.0 Disks and Files



- DBMS stores information on harddisks or flash SSDs.
  - Electronic (CPU, DRAM) vs. Mechanical (harddisk) vs. Electronic (SSD)
- This has major implications for DBMS design!
  - **READ**: transfer data from disk to main memory (RAM).
  - **WRITE**: transfer data from RAM to disk.
  - Both are expensive operations, relative to in-memory operations, so must be **planned carefully**!
    - ✓ DRAM: ~ 10 ns
    - ✓ Harddisk: ~ 10ms
    - ✓ SSD: **80us** ~ 10ms
- CS (and DBMS) is a discipline about numerous trade-offs.
  - Space vs. time; cost vs. performance

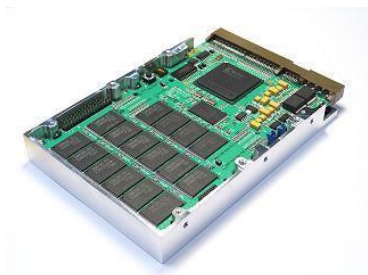
# Non-Volatile Secondary Storage: Flash SSD vs. Harddisk



60 years champion

**VS**

A new challenger!

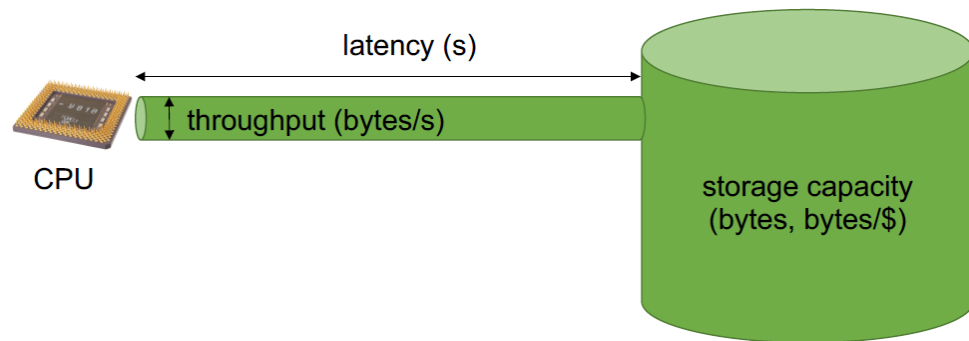
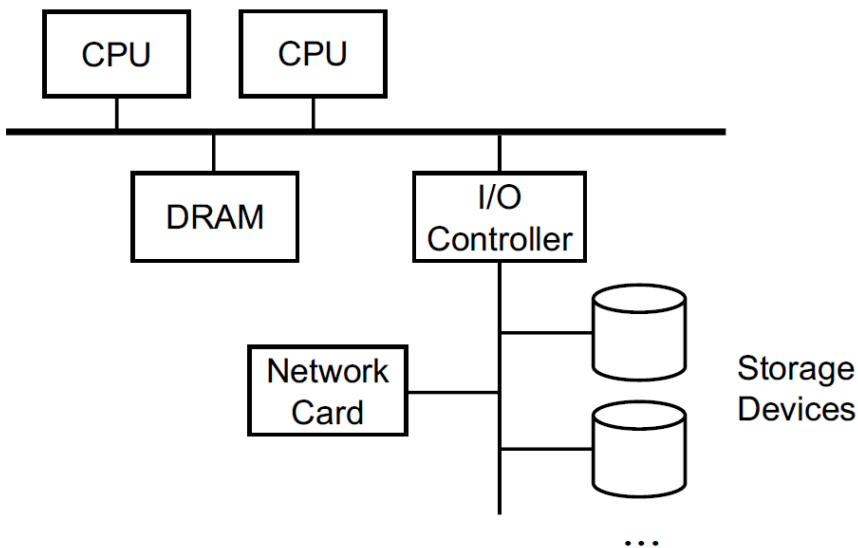


Identical  
Interface



Flash SSD	HDD
Electronic	Mechanical
Read/Write Asymmetric	Symmetric
No Overwrite	Overwrite

# Typical Server and Storage Performance Metrics



## Max throughput

large reads ( $\gg 1$ block):

- DRAM: 100GB/s
- NVMe SSD: 2GB/s
- HDD: 130MB/s

\$1,000 @ NewEgg:

- 0.25 TB of DRAM
- 9TB of NVMe SSD
- 50TB of HDD

Source: <http://web.stanford.edu/class/cs245/slides/03-System-Architecture-p2.pdf>

# Storage Performance Metrics

- **Capacity** (\$/GB) : Harddisk >> Flash SSD
- **Bandwidth** (MB/sec): Harddisk < Flash SSD
- **Latency** (IOPS): Harddisk << Flash SSD
  - e.g. Harddisk
    - ✓ Commodity hdd (7200rpm): 50\$ / 1TB / 100MB/s / 100 IOPS
    - ✓ Enterprise hdd(1.5Krpm): 250\$ / 72GB / 200MB/s / 500 IOPS
    - ✓ The price of harddisks is said to be **proportional to IOPS**, **not capacity**.

Storage Media	4KB Random Throughput (IOPS)		Sequential Bandwidth (MB/sec)		Capacity in GB	Price in \$ (\$/GB)
	Read	Write	Read	Write		
MLC SSD <sup>†</sup>	28,495	6,314	251.33	242.80	256	450 (1.78)
MLC SSD <sup>‡</sup>	35,601	2,547	258.70	80.81	80	180 (2.25)
SLC SSD <sup>§</sup>	38,427	5,057	259.2	195.25	32	440 (13.75)
Single disk <sup>¶</sup>	409	343	156	154	146.8	240 (1.63)
8-disk <sup>¶</sup> RAID-0	2,598	2,502	848	843	1,170	1,920 (1.63)

SSD: <sup>†</sup>Samsung 470 Series 256GB, <sup>‡</sup>Intel X25-M G2 80GB, <sup>§</sup>Intel X25-E 32GB

<sup>¶</sup>Disk: Seagate Cheetah 15K.6 146.8GB

TABLE I

PRICE AND PERFORMANCE CHARACTERISTICS OF FLASH MEMORY SSDS AND MAGNETIC DISK DRIVES

# Storage Device Metrics(2): HDD vs. Flash SSDs

- Other Metrics: Weight/shock resistance/heat & cooling, power(watt) , IOPS/watt, IOPS/\$ ....
  - Harddisk << Flash SSD

Table 1. Basic disk and solid-state disk (SSD) performance and cost.

Drive	Cost	Capacity	Power draw while reading data	Read IOPS	Throughput
2.5" disk	\$40	320 Gbytes	2.1 watts/0.75 watt	240	80 Mbytes per second (spec sheet)
Solid-state disk	\$220	80 Gbytes	0.1 watt	35,000	250 MBps

Table 2. Per-dollar and per-watt performance.

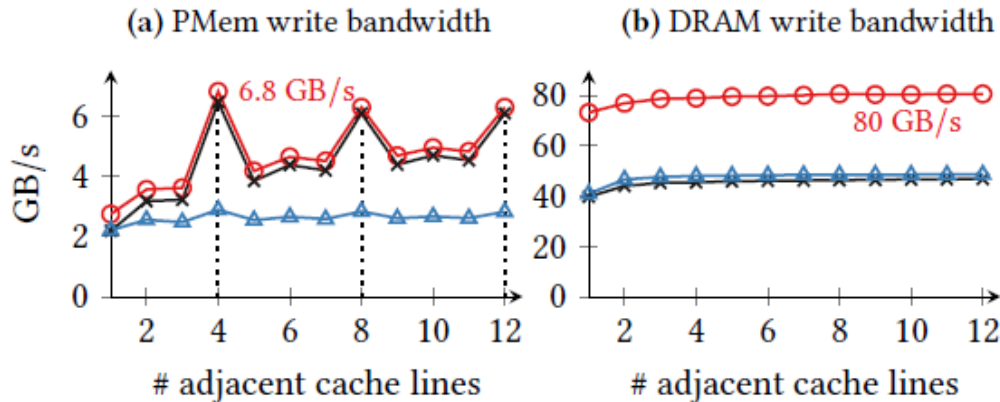
Drive	Gbytes per dollar	IOPS per dollar	IOPS per watt	Throughput per dollar	Throughput per watt
2.5" disk	4	3	104	1 Mbytes/second	40 MBps
Solid-state disk	0.36	159	220,000	1.13 MBps	2,500 MBps

*[Source: Rethinking Flash In the Data Center, IEEE Micro 2010]*



# Intel Optane DC Persistent Memory vs. Z-SSD

	DRAM	Optane DC PMM	SSD
Read Latency	73 ns	300 ns	230 $\mu$ s
Seq. Read BW	110 GB/s	36 GB/s	3.5 GB/s
Rand. Read BW	100 GB/s	10 GB/s	1.9 GB/s
Byte-addressable	Yes	Yes	No



Samsung SZ985 Z-NAND SSD	
Form Factor	HHHL
Interface	PCIe Gen3 x4
NAND	Z-NAND Technology
Port	Single
Data Transfer Rate (128KB data size)	
Sequential Read / Write (GB/s)	3.2 / 3.2
Data I/O Speed (4KB data Size, sustained)	
Random Read / Write (IOPs)	750K/ 170K
Latency (sustained random workload)	
Random Read	12 - 20 $\mu$ s
Random Write (Typical)	16 $\mu$ s
DWPD	30
Capacity	800GB

# Bandwidth Crisis in AI/ML era?

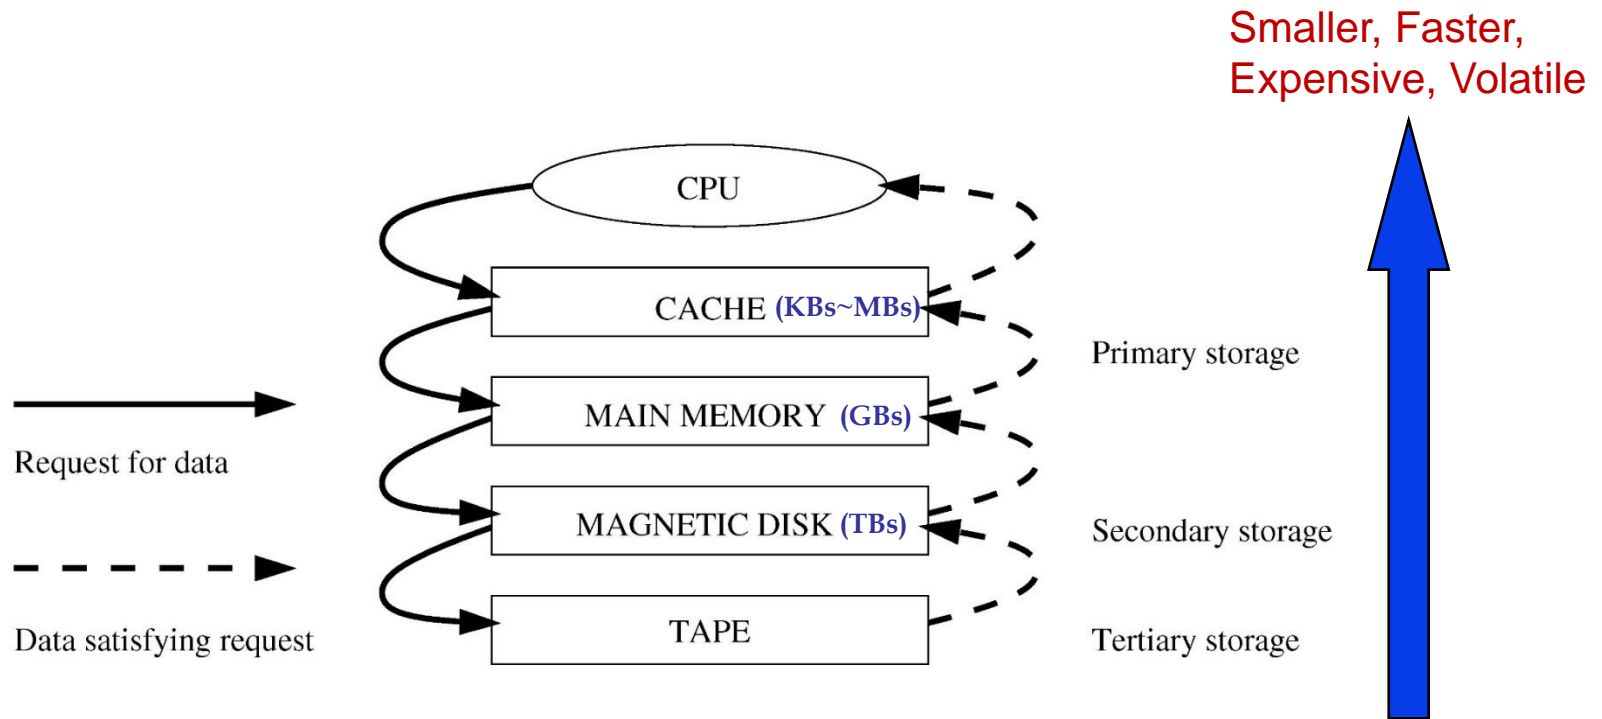
- Data-intensive ML algorithms. (source: [Compressed Linear Algebra for Declarative Large-Scale Machine Learning](#))
  - Many ML algorithms are **iterative, with repeated read-only data access**. These algorithms often rely on matrix-vector multiplications, which require one complete scan of the matrix with only **two floating point operations per matrix element**. This low operational intensity renders matrix-vector multiplication, even in-memory, I/O bound.<sup>18</sup> **Despite the adoption of flash-and NVM-based SSDs, disk bandwidth is usually 10x-100x slower than memory bandwidth, which is in turn 10x-40x slower than peak floating point performance.** Hence, it is crucial for performance to fit the matrix into available memory without sacrificing operations performance. This challenge applies to single-node in-memory computations, data-parallel frameworks with distributed caching like Spark,<sup>20</sup> and accelerators like GPUs with limited device memory. Even in the face of emerging memory and link technologies, the challenge persists due to increasing data sizes, different access costs in the memory hierarchy, and monetary costs.
- One real problem
  - Solution?: 1) cheaper but higher BW memory devices (Intel Optane DC PM), 2) data compression (above link), 3) making ML algorithms more computation intensive (?), 4) offloading ML algorithms near to storage, 5) use multiple low-spec CPU

# Storage Wars: File vs. Block vs. Object Storage

	File	Block	Object
Use cases	<ul style="list-style-type: none"><li>- File sharing</li><li>- Local archiving</li><li>- Data Protection</li></ul>	<ul style="list-style-type: none"><li>- DB</li><li>- Email server</li><li>- RAID</li><li>- VM</li></ul>	<ul style="list-style-type: none"><li>- Big data</li><li>- Web apps</li><li>- Backup archives</li></ul>

- Object storage is good for scalability for big data (<https://blog.storagecraft.com/object-storage-systems/>)
  - Scalability is where object-based storage does its most impressive work. Scaling out an object architecture is as simple as adding additional nodes to the storage cluster. Every server **has its physical limitations**. But thanks to location transparency and remarkable metadata flexibility, this type of storage can be scaled without the capacity limits that plague traditional systems.
- Amazon S3 (Simple Storage Service) and REST API
  - <https://docs.aws.amazon.com/AmazonS3/latest/API/Welcome.html>

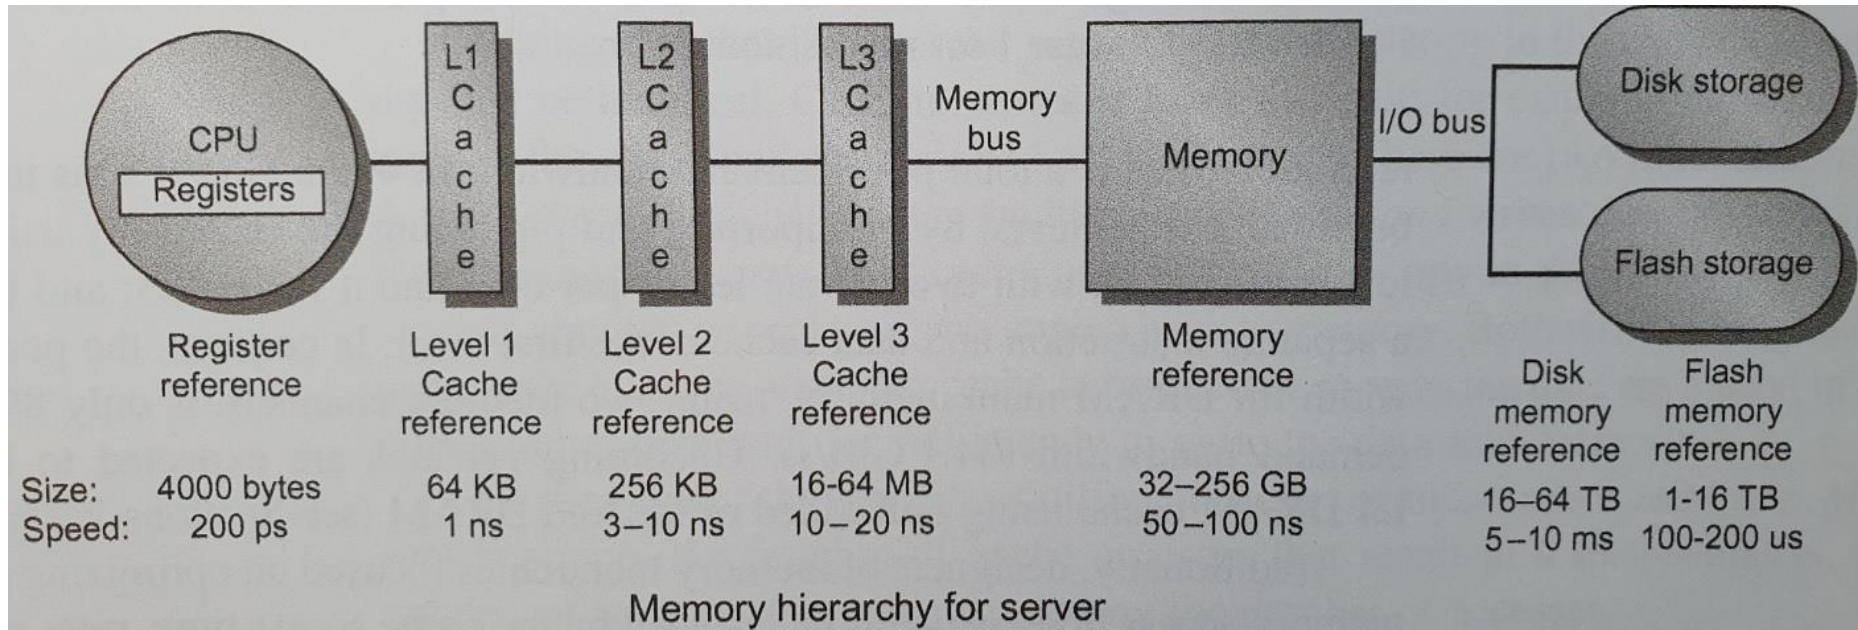
# 9.1 Memory Hierarchy



- Main memory (RAM) for currently used data.
- Disk for the main database (secondary storage).
- Tapes for archiving older versions of the data (tertiary storage)
- **WHY MEMORY HIERARCHY?**
- **What if ideal storage appear? Fast, cheap, large, NV..: PCM, MRAM, FeRAM?**

**Bigger, Slower,  
Cheaper, Non-Volatile**

# Memory Hierarchy for Server



Source: Figure 1.2 in "Computer Architecture: A Quantitative Approach (6<sup>th</sup> Ed.)"

# Disks

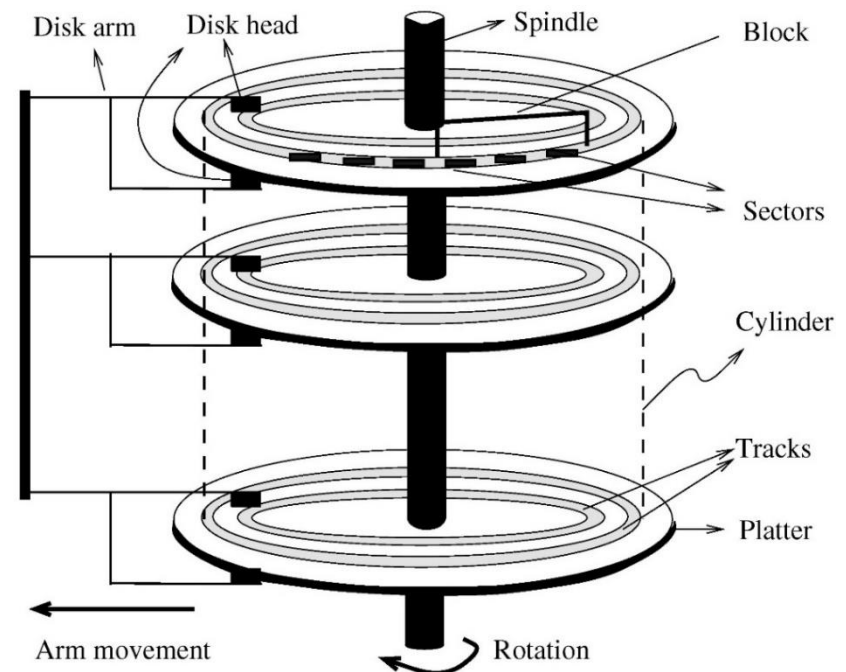


- Secondary storage device of choice. (non-volatile, durable)
- Main advantage over **tapes**: random access vs. sequential.
  - E.g. To find a record (with its address known) among 1 billion records:
- Data is stored and retrieved in disk blocks or pages unit.
- Unlike RAM, time to retrieve a disk page varies **depending upon location on disk**.
  - Thus, relative placement of pages on disk has big impact on DB performance!
    - ✓ e.g. adjacent allocation of the pages from the same tables.
  - We need to optimize both data placement and access
    - ✓ e.g. elevator disk scheduling algorithm

# Anatomy of a Disk

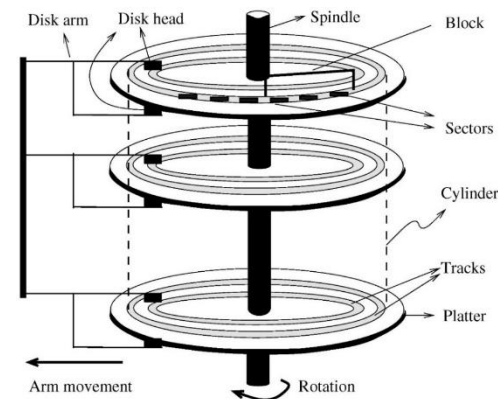


- The **platters** spin
  - e.g. 5400 / 7200 / 15K rpm
- **The arm assembly** is moved in or out to position a head on a desired track. Tracks under heads make a **cylinder**
  - **Mechanical storage** -> **low IOPS**
- Only **one head** reads/writes at any one time.
  - Parallelism degree: 1
- Block size is a **multiple** of sector size
- **Update-in-place**: poisoned apple
- **No atomic write**
- **Fsync** for ordering / durability





# Accessing a Disk Page



- Time to access (read/write) a disk block:
  1. **Seek time** (moving arms to position disk head on track)
  2. **Rotational delay** (waiting for block to rotate under head)
  3. **Transfer time** (actually moving data to/from disk surface)
- Mechanical devices: seek time and rotational delay dominate.
  - Seek time: about 1 to 20msec
  - Rotational delay: from 0 to 10msec
  - Transfer rate: about 1ms per 4KB page
- Key to lower I/O cost: **reduce seek/rotational delays!**
  - E.g. disk scheduling algorithm in OS, Linux 4 I/O schedulers



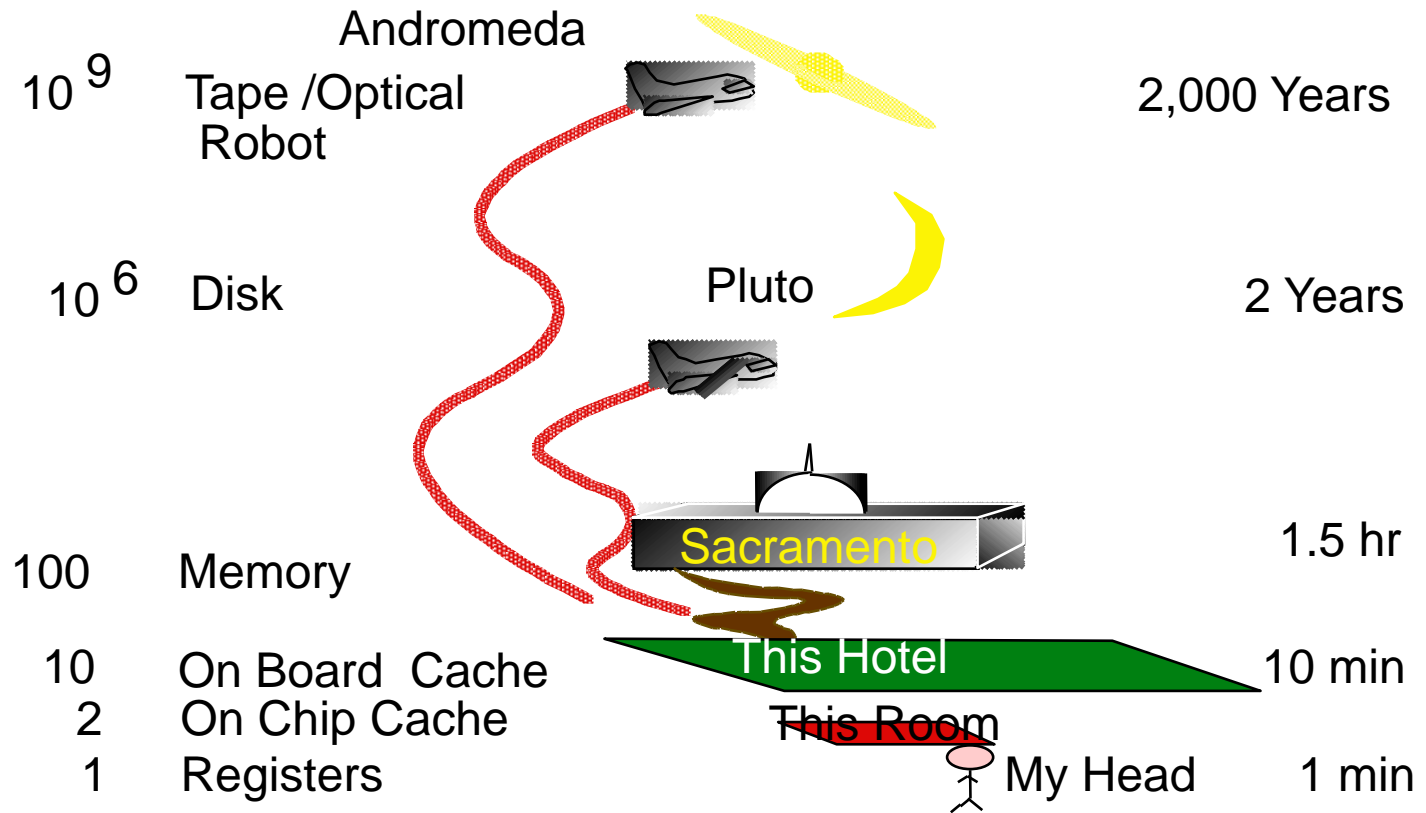
# Arranging Pages on Disk

- ‘Next’ block concept:
  - Blocks on same track, followed by
  - Blocks on same cylinder, followed by
  - Blocks on adjacent cylinder
- Blocks in a file should be arranged **sequentially** on disk (by ‘next’), to minimize seek and rotational delay.
  - E.g. extent-based allocation (Section 9.5)
- **Disk fragmentation problem** (see [https://en.wikipedia.org/wiki/File\\_system\\_fragmentation](https://en.wikipedia.org/wiki/File_system_fragmentation))
  - Is disk fragmentation still problematic in flash storage?
  - Not a big deal in read, but a big deal in write?

# Some Techniques to Hide IO Bottlenecks

- **Pre-fetching**: For a sequential scan, pre-fetching several pages at a time is a big win! Even cache / disk controller support prefetching
- **Caching**: modern disk controllers do their own **caching**.
- **IO overlapping**: CPU works while IO is performing
  - Double buffering, asynchronous IO
- **Multiple threads**
- **And, don't do IOs, avoid IOs**

# Jim Gray's Storage Latency Analogy: How Far Away is the Data?



# Latency Numbers Every Programmer Should Know

Typical one instruction .....	01 ns	
L1 cache reference .....	0.5 ns	
Branch mispredict .....	5 ns	
L2 cache reference .....	7 ns	
Mutex lock/unlock .....	25 ns	
Main memory reference .....	100 ns	
Compress 1K bytes with Zippy .....	3,000 ns	= 3 $\mu$ s
Send 2K bytes over 1 Gbps network .....	20,000 ns	= 20 $\mu$ s
SSD random read .....	150,000 ns	= 150 $\mu$ s
Read 1 MB sequentially from memory .....	250,000 ns	= 250 $\mu$ s
Round trip within same datacenter .....	500,000 ns	= 0.5 ms
Read 1 MB sequentially from SSD* .....	1,000,000 ns	= 1 ms
Disk seek .....	10,000,000 ns	= 10 ms
Read 1 MB sequentially from disk .....	20,000,000 ns	= 20 ms
Send packet CA->Netherlands->CA .....	150,000,000 ns	= 150 ms

X  $10^9$

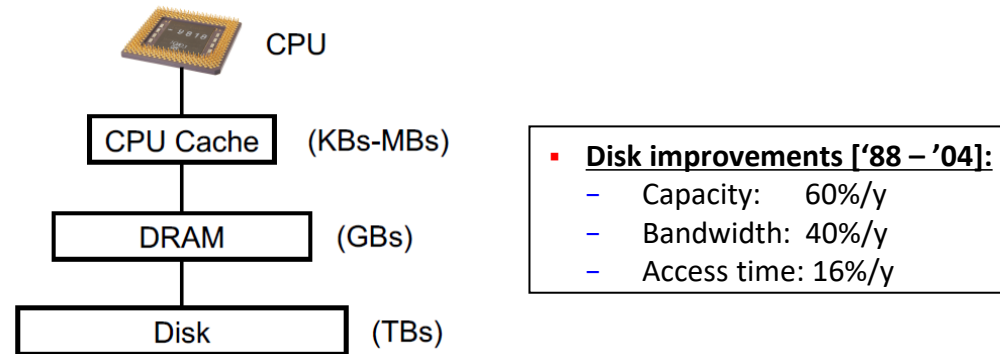
Assuming ~1GB/sec SSD

<https://gist.github.com/hellerbarde/2843375>

Data by [Jeff Dean](#); Originally by [Peter Norvig](#)

# Technology RATIOS Matter

[ Source: Jim Gray's PPT ]



- Technology ratio change: 1980s vs. 2020s
  - If everything changes in the same way, then nothing really changes.
  - If some things get much cheaper/faster than others, then that is **real change**.
    - Some things are not changing much (e.g., cost of people, speed of light) while other things are changing a LOT (e.g., Moore's law, disk capacity)
  - Harddisk: "Latency lags behind bandwidth" and "bandwidth does behind capacity"
- Flash memory/NVRAMs and its role in the memory hierarchy?
  - Disruptive technology ratio change → new disruptive solution?

# Evolution of DRAM, HDD, and SSD (1987 – 2017)

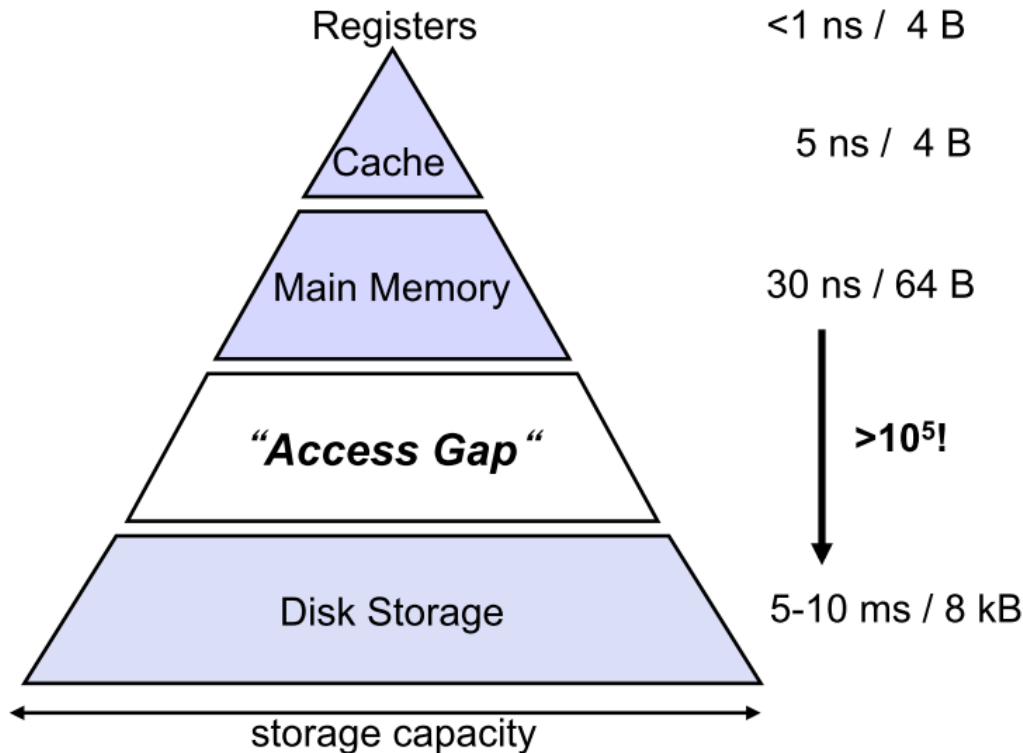
- Raja et. al., The Five-minute Rule Thirty Years Later and its Impact on Storage Hierarchy, ADMS '17
- **the Storage Hierarchy**

Metric	DRAM				HDD				SATA Flash SSD	
	1987	1997	2007	2017	1987	1997	2007	2017	2007	2017
Unit price(\$)	5k	15k	48	80	30k	2k	80	49	1k	560
Unit capacity	1MB	1GB	1GB	16GB	180MB	9GB	250GB	2TB	32GB	800GB
\$/MB	5k	14.6	0.05	0.005	83.33	0.22	0.0003	0.00002	0.03	0.0007
Random IOPS	-	-	-	-	15	64	83	200	6.2k	67k (r)/20k (w)
Sequential b/w (MB/s)	-	-	-	-	1	10	300	200	66	500 (r)/460 (w)

Table 1: The evolution of DRAM, HDD, and Flash SSD properties

# Latency Gap in Memory Hierarchy

## Typical access latency & granularity



We need 'gap filler'!  
: Flash memory SSD!

*[Source: Uwe Röhm's Slide]*

- Latency lags behind bandwidth [David Patterson, CACM Oct. 2004]
  - Bandwidth problem can be cured with money, but latency problem is harder

# Why Not Store It All in Main Memory?

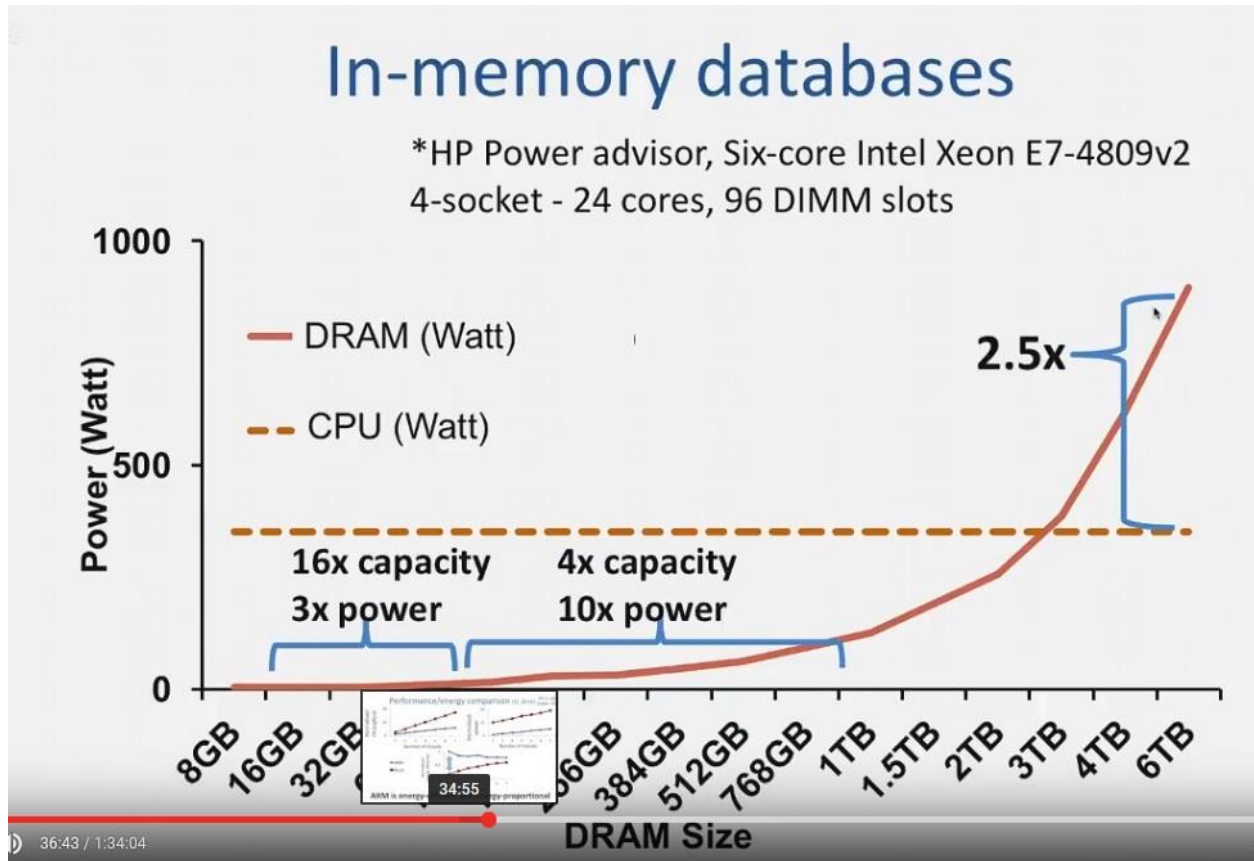
- **Cost!**: 20\$ /1GB DRAM vs. 50\$ / 150 GB of disk (EIDI/ATA) vs. 100\$/30GB (SCSI).
  - High-end databases today are in the 10-100 TB range.
  - Approx. **60% of the cost** of a production system is in the **disks**.
- Some specialized systems (e.g. Main Memory(MM) DBMS) store entire database in main memory.
  - Vendors claim 10x speed up vs. traditional DBMS in main memory.
  - Sap Hana, MS Hekaton, [Oracle In-memory](#), Altibase ..
- Main memory is **volatile**: data should be saved between runs.
  - **Disk write** is **inevitable**: 1) log write for recovery and 2) periodic checkpoint



# MM-DBMS

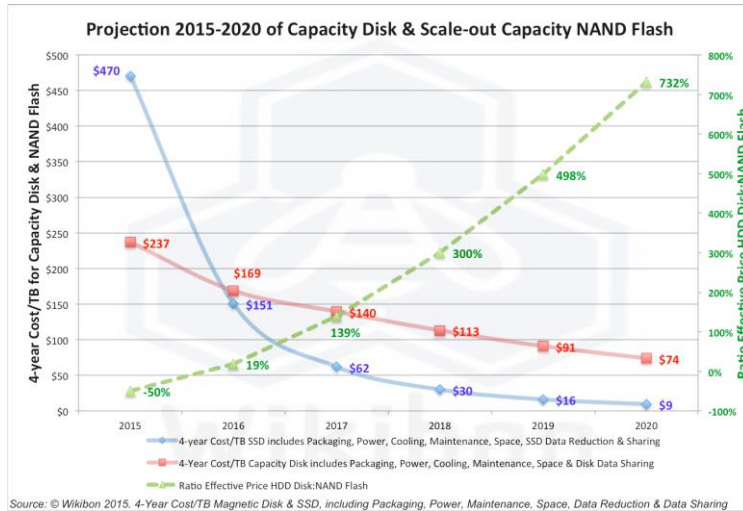
- Why MMDBMS has been recently popular since mid-2000s?
  - Sap Hana, MS Hekaton, [Oracle In-memory](#), Altibase, ....
  - The price of DRAM had ever dropped for the last two decades
    - ✓ \$/IOPS @ DISK >> \$/GB @ DRAM
  - The overhead of disk-based DBMS is not negligible
  - Applications with extreme performance requirements?

# Power Consumption Issue in Big Memory



- Why exponential?
- 1KWh = 15 ~ 50 cents, 1 year = 1,752\$

# HDD vs. SSD [Patterson 2016]



## Future Memory Hierarchy Deeper

- Storage hierarchy gets more and more complex:
  - L1 cache
  - L2 cache
  - L3 cache
  - Fast DRAM (on interposer with CPU)
  - 3D XPoint based storage
  - SSD
  - (HDD)
- Need to design software to take advantage of this hierarchy



## SSDs vs. HDDs

- SSDs will soon become cheaper than HDDs
- Transition from HDDs to SSDs will accelerate
  - Already most instances in Amazon Web Service have SSDs
- Going forward we can assume SSD-only clusters

"Tape is dead, Disk is tape, Flash is disk."  
Jim Gray, 2007

# Evolution of secondary storages

- Source: Oracle Magazine, July/August 2014
- [The Life of a Data Byte](#) (CACM, Dec. 2020)
  - A short history about modern storage medias

56

Time Capsule

Flashbacks: Culture. Industry. Oracle. Oracle Magazine.

BY RICH SCHWERIN



1951

Tape Drive

The Remington Rand UNISERVO was the primary I/O device on the UNIVAC I computer, and stored up to 224 KB on a 1,200-foot-long metal tape.

1956



Hard Disk

The refrigerator-sized IBM 350 disk drive held 3.75 MB and leased for US\$3,200 a month. Inflation adjusted, that's nearly US\$28,000 today—or US\$7,400 per MB.


1970s



Floppy Disks

Data storage in the '70s and into the '80s? Floppy. From 8-inch to 5¼-inch to 3½-inch, these disks of thin, flexible magnetic storage medium were state of the art. Just ask your mom.

1999



SD Cards

Initially 64 MB, the Secure Digital (SD) memory card storage from SanDisk, Matsushita, and Toshiba has been getting smaller in size and larger in capacity ever since. (Today's MicroSD holds 128 GB.)


2000



USB Flash Drive

The first ThumbDrive from Trek Technology plugged into any USB port and offered a whopping 8 MB storage capacity. And within just a few years, thumb drives were making fashion statements. Sushi, anyone?

2011



Storage from A to ZFS

Organizations are optimizing storage with tiered Sun flash, disk, and tape solutions from Oracle and enabling unified storage with the Oracle ZFS Storage Appliance.

2013



Extreme Memory

In a single rack, Oracle's Exadata Database Machine X4 supports 88 TB of user data in flash—a capacity sufficient to hold the majority of online transaction processing databases in flash memory.

YOUR TURN

FROM 8-INCH FLOPPIES TO 88 TB IN FLASH, tell us about your first storage, your ultimate storage, and where you think storage will be in five years. Visit Facebook/OracleMagazine and let us know. [bit.ly/orclmagfb](http://bit.ly/orclmagfb)

JULY/AUGUST 2014 [ORACLE.COM/ORACLEMAGAZINE](http://ORACLE.COM/ORACLEMAGAZINE)

# Implications for DBMS Design

- The access characteristics of storage devices (e.g. hard disks and flash SSDs) necessitate that database systems have the ability to control *where*, *how* and *when* data is physically accessed.
- **Disk Space Management:** ‘Spatial control’
  - **Where** on the secondary storage is the data stored?
- **Buffer Management:** ‘Temporal control’
  - **When** is data physically read from or written to disk?
- **Query Optimization and Execution:** ‘Access pattern control’
  - **How** is data accessed? Sequentially or Random Access?

# Mega Changes in Computer Architectures and Implications on Database Technology

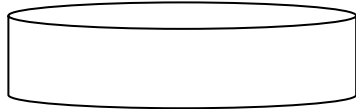
- CPU: Single-core → Multi-core (End of Moore's Law?)
- DRAM: small and expensive → large and become cheaper
- Storage: HDD → Flash SSD (→ NVRAM ?)
- Data center/Cloud: disaggregation, object storage



- MMDBMS
- New concurrency control: 2PL → OCC
- Buffer replacement algorithm
- Cloud-native DBMS (Amazon Aurora, Snowflake)
- And, many others

## 9.2 RAID

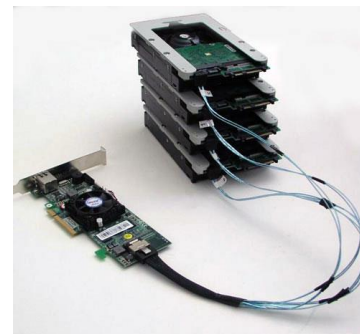
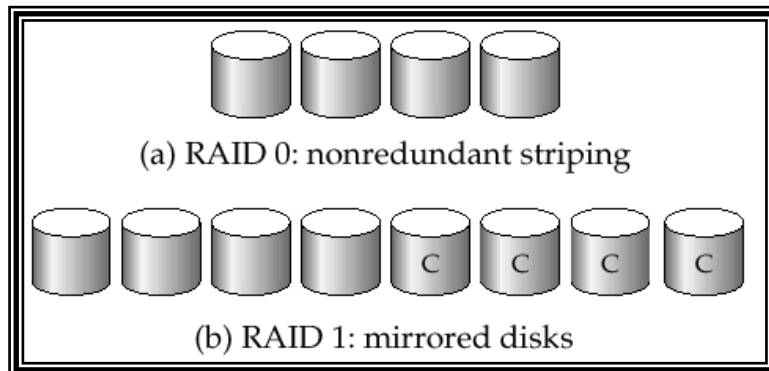
- **SLED** (Single Large Expensive Disk) approach till 1980s



vs.



- Redundant Arrays of Independent(or Inexpensive) Disks
  - Disk array: arrangement of several disks that gives abstraction of a single, large disk.
- Goals: Increase **performance** and **reliability**.

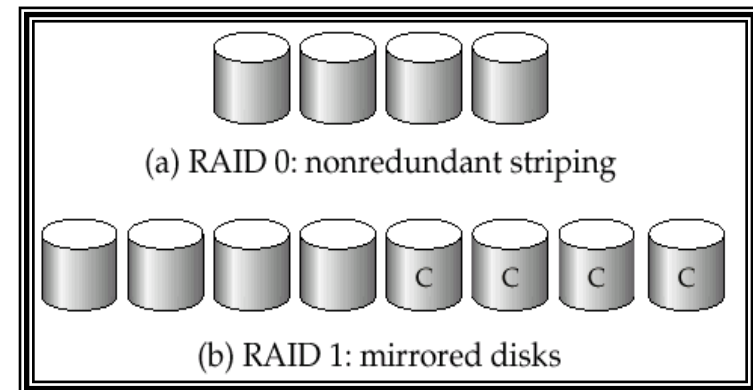


Cf. Tesla Battery and Rocket Tech.



# RAID

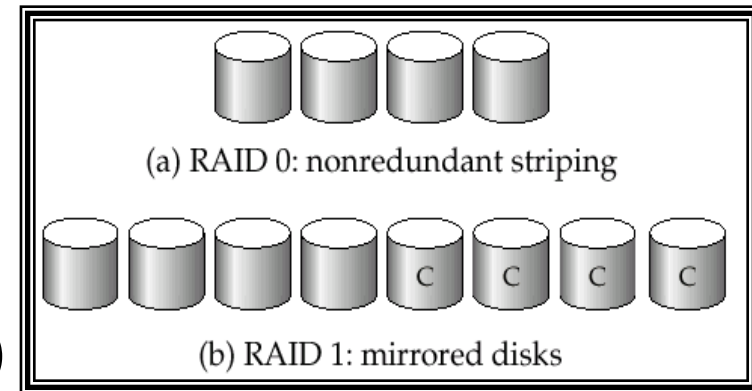
- Two main techniques:
  - **Data striping**: Data is partitioned; size of a partition is called the **striping unit**. Partitions are distributed over several disks.
    - ✓ For large data, **larger bandwidth** (i.e. transfer rate)
    - ✓ For small random data, **higher IOPS**
  - **Mirroring for redundancy**: More disks => more failures. Redundant information allows reconstruction of data if a disk fails.
- Benefits of RAID
  - Bandwidth for sequential IOs
  - IOPS for random IOs
  - Reliability by redundancy
- Another beauty in computer science
  - Simple and powerful!!





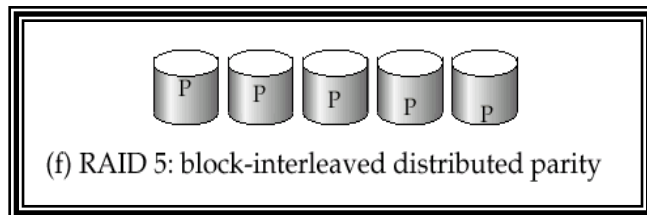
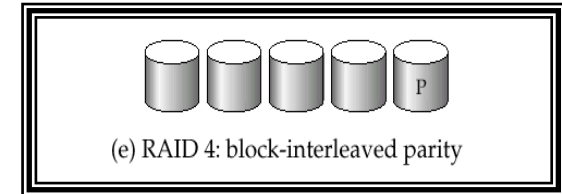
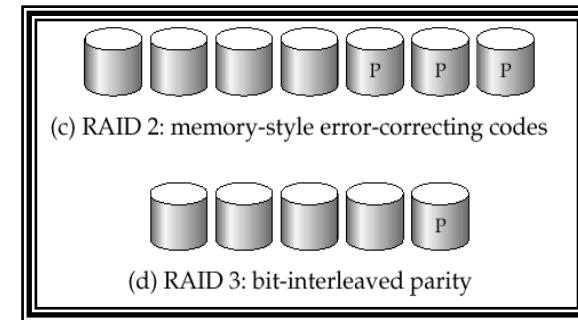
# RAID Levels

- Level 0: No redundancy
- Level 1: Mirrored (two identical copies)
  - Each disk has a mirror image (check disk)
  - Parallel reads, a write involves two disks.
  - Maximum transfer rate = transfer rate of one disk
- Level 0+1: Striping and Mirroring
  - Parallel reads, a write involves two disks.
  - Maximum transfer rate = aggregate bandwidth



# RAID Levels (Contd.)

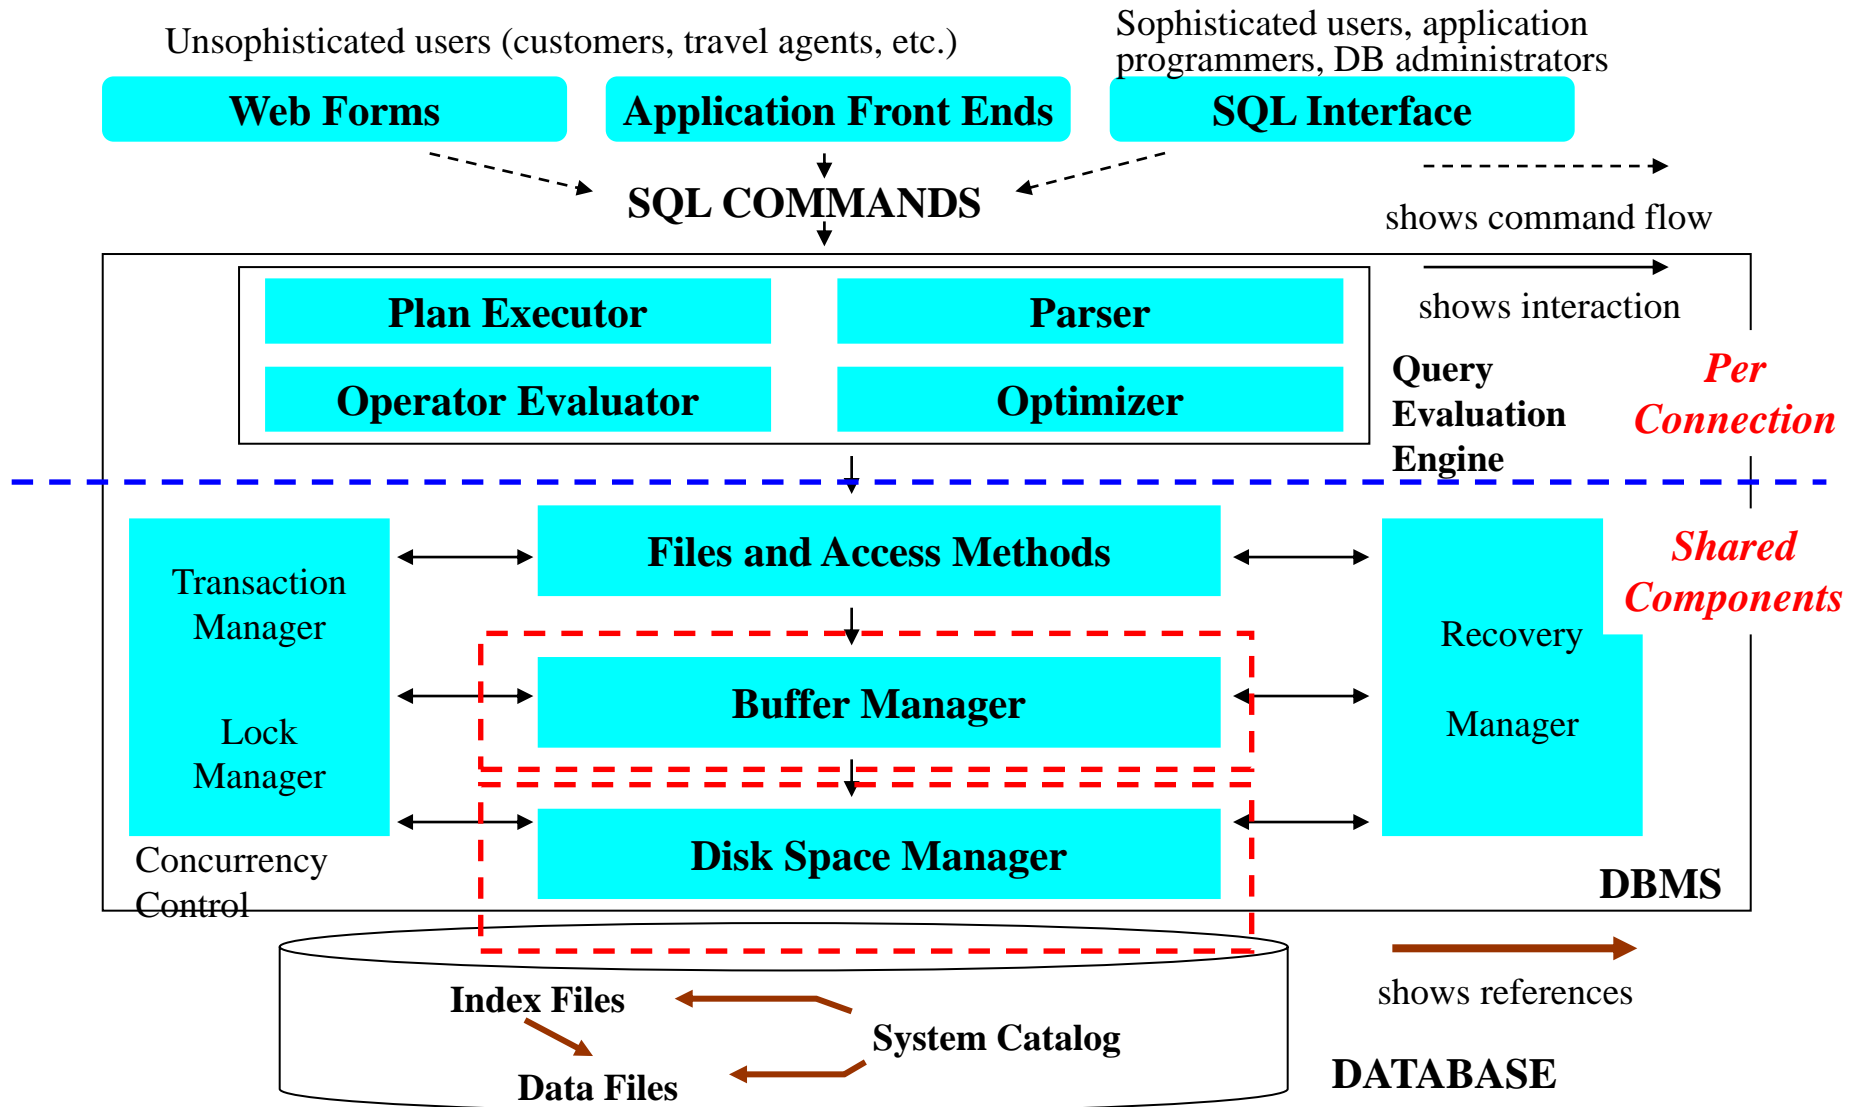
- Level 3: Bit-Interleaved Parity
  - Striping Unit: One bit. One check disk.
  - Each read and write request involves all disks; disk array can process one request at a time.
- Level 4: Block-Interleaved Parity
  - Striping Unit: One disk block. One check disk.
  - Parallel reads possible for small requests, large requests can utilize full bandwidth
  - Writes involve modified block and check disk
- Level 5: Block-Interleaved Distributed Parity
  - Similar to RAID Level 4, but parity blocks are distributed over all disks



P0	0	1	2	3
4	P1	5	6	7
8	9	P2	10	11
12	13	14	P3	15
16	17	18	19	P4

## 9.3 Disk Space Management

- Lowest layer of DBMS software manages space on disk.
- Higher levels call upon this layer to:
  - allocate/de-allocate a page
  - read/write a page
- Request for a sequence of pages must be satisfied by allocating the pages sequentially on disk
- Higher levels don't need to know how this is done, or how free space is managed.



**Figure 1.3 Architecture of a DBMS**

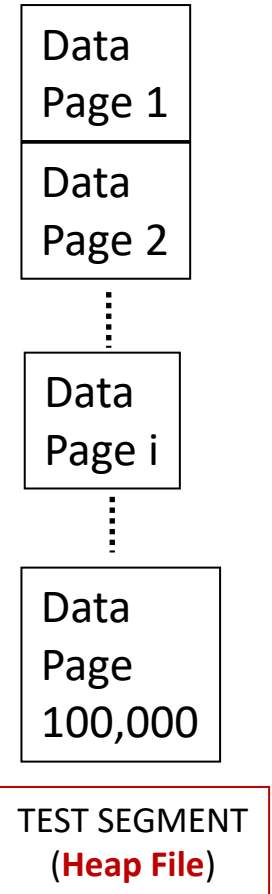
## 9.4 Buffer Management in DBMS

# Table, Insertions, Heap Files

```
CREATE TABLE TEST (a int, b int, c varchar2(650));

/* Insert 1M tuples into TEST table (approximately 664 bytes per
tuple) */
BEGIN
  FOR i IN 1..1000000 LOOP
    INSERT INTO TEST (a, b, c) values (i, i, rpads('X', 650, 'X'));
  END LOOP;
END;

/*
Page = 8KB
10 tuples / page
100,000 pages in total
TEST table = 800MB
*/
```



# OLAP vs. OLTP

- On-Line Analytical vs. Transactional Processing

```
SQL> SELECT SUM(b) FROM TEST;
```

```
SUM(B)
```

```
-----
```

```
5.0000E+11
```

```
Execution Plan
```

```
-----
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
----	-----------	------	------	-------	-------------	------

```
-----
```

0	SELECT STATEMENT		1	5	22053 (1)	00:04:25
---	------------------	--	---	---	-----------	----------

1	SORT AGGREGATE		1	5		
---	----------------	--	---	---	--	--

2	TABLE ACCESS FULL	TEST	996K	4865K	22053 (1)	00:04:25
---	-------------------	------	------	-------	-----------	----------

```
-----
```

```
Statistics
```

```
-----
```

```
179 recursive calls
```

```
0 db block gets
```

```
100152 consistent gets
```

```
100112 physical reads
```

```
.....
```

```
1 rows processed
```

Data  
Page 1

Data  
Page 2

⋮

Data  
Page i

⋮

Data  
Page  
100,000

TEST SEGMENT  
(Heap File)

# OLAP vs. OLTP

- On-Line Transactional Processing: Point or Range Query

```
SQL> SELECT B FROM TEST WHERE A = 500000;
```

B

500000

Execution Plan

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	5	22053 (1)	00:04:25
1	SORT AGGREGATE		1	5		
2	TABLE ACCESS FULL	TEST	996K	4865K	22053 (1)	00:04:25

Statistics

179 recursive calls  
0 db block gets  
100152 consistent gets  
100112 physical reads

.....

1 rows processed

No Index on A column  
Thus, Full Table Scan

Data  
Page 1

Data  
Page 2

...

Data  
Page i

...

Data  
Page  
100,000

TEST SEGMENT  
(Heap File)



# OLAP vs. OLTP

**CREATE INDEX TEST\_A ON TEST(A);**

```
SELECT B          /* point query */
```

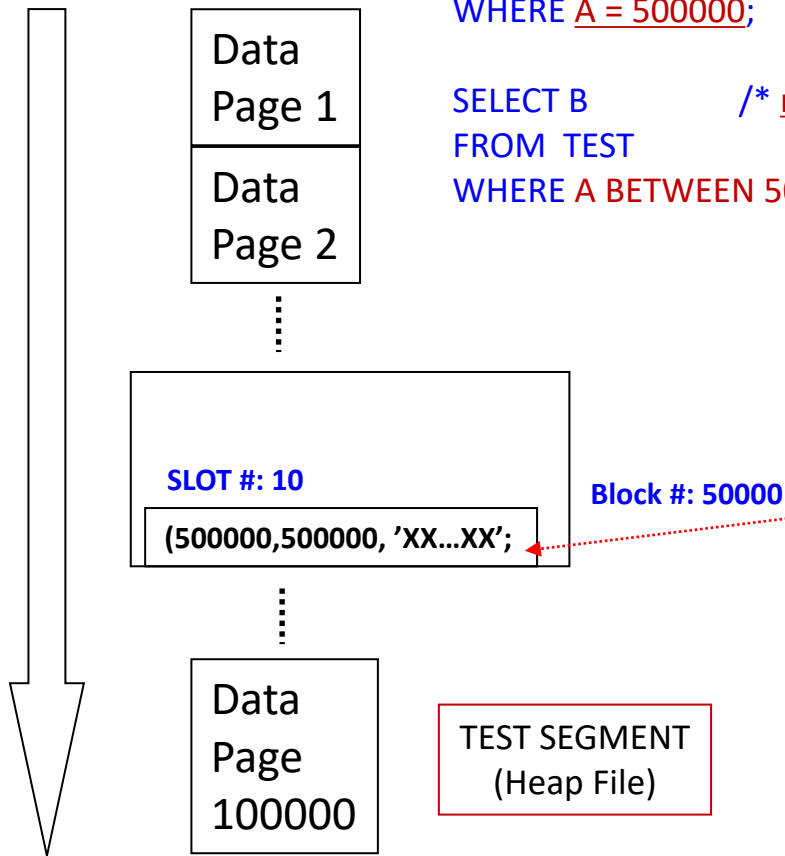
```
FROM TEST
```

```
WHERE A = 500000;
```

```
SELECT B          /* range query */
```

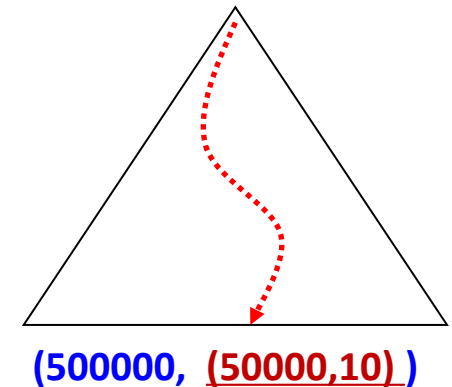
```
FROM TEST
```

```
WHERE A BETWEEN 50001 and 50101;
```



B-tree Index on TEST(A)

**SEARCH KEY: 500000**



Cost:

- Full Table Scan: 100,000 Block Accesses
- Index: (3~4) + Data Block Access
  - Point query: 1
  - Range queries: depending on range

# OLAP vs. OLTP

- Index-based Table Access

```
SQL> SELECT B FROM TEST WHERE A = 500000;
```

```
B
```

```
500000
```

```
Execution Plan
```

```
-----  
| Id | Operation                | Name  | Rows | Bytes | Cost (%CPU)| Time     |  
-----  
| 0 | SELECT STATEMENT          |       |      |      | 4 (0)| 00:00:01 |  
| 1 | TABLE ACCESS BY INDEX ROWID| TEST  | 1    | 10    | 4 (0)| 00:00:01 |  
|* 2 | INDEX RANGE SCAN          | TEST_A | 1    |      | 3 (0)| 00:00:01 |  
-----
```

```
Statistics
```

```
.....
```

```
5 consistent gets
```

```
4 physical reads
```

```
.....
```

```
1 rows processed
```

Column A is now Indexed,  
Thus, Index-Scan!

# OLTP: TPC-A/B/C Benchmark

```
ACCOUNT (ACCOUNT_NUMBER, CUSTOMER_NUMBER, ACCOUNT_BALANCE, HISTORY)
CUSTOMER (CUSTOMER_NUMBER, CUSTOMER_NAME, ADDRESS, ..... )
HISTORY (TIME, TELLER, CODE, ACCOUNT_NUMBER, CHANGE, PREV_HISTORY)
CASH_DRAWER (TELLER_NUMBER, BALANCE)
BRANCH_BALANCE (BRANCH, BALANCE)
TELLER (TELLER_NUMBER, TELLER_NAME, ..... )
```

*From Gray's Presentation*

```
exec sql begin declare section;
long Aid, Bid, Tid, delta, Abalance;
exec sql end declare section;
DCApplication()
{  read input msg;
```

```
    exec sql begin work;
```

```
    exec sql update accounts set Abalance = Abalance + :delta where Aid = :Aid;
```

```
    exec sql select Abalance into :Abalance from accounts where Aid = :Aid;
```

```
    exec sql update tellers set Tbalance = Tbalance + :delta where Tid = :Tid;
```

```
    exec sql update branches set Bbalance = Bbalance + :delta where Bid = :Bid;
```

```
    exec sql insert into history(Tid, Bid, Aid, delta, time) values (:Tid, :Bid, :Aid, :delta, CURRENT);
```

```
    send output msg;
```

```
    exec sql commit work; }
```

A transaction

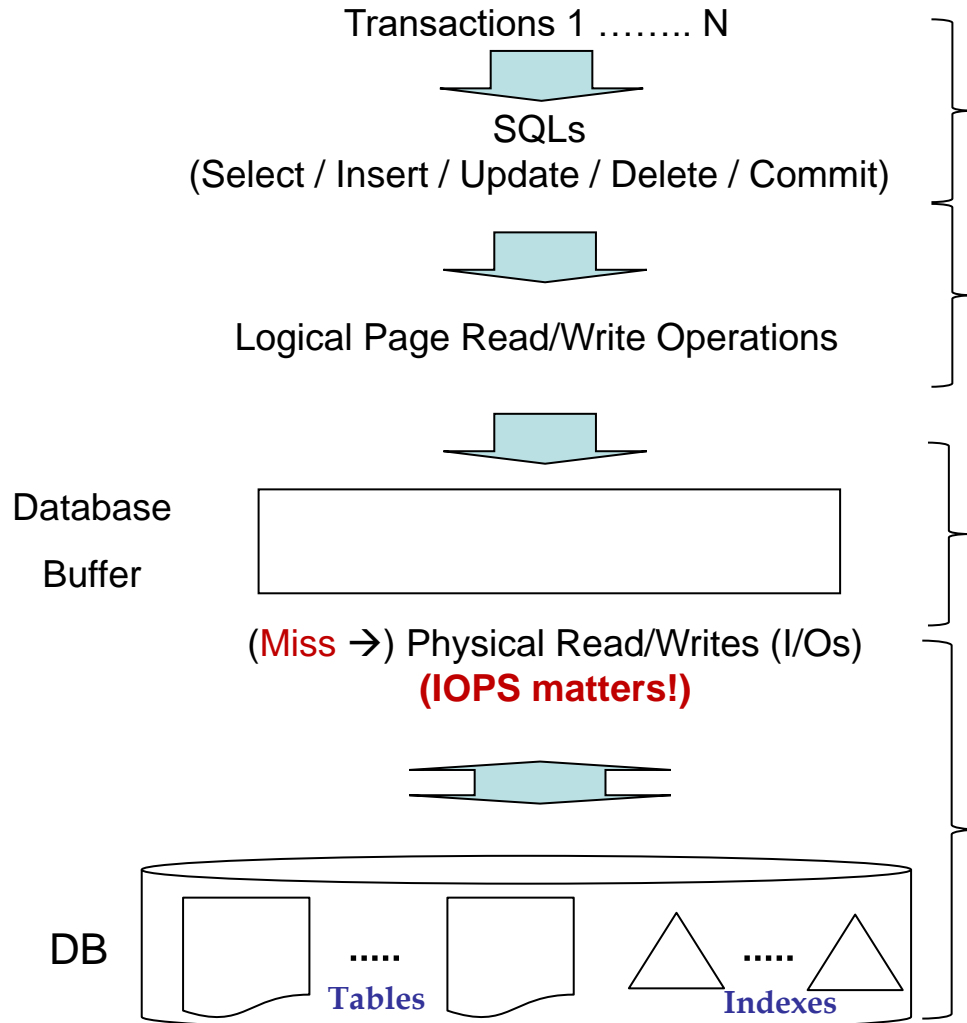
= A sequence of SQL statements

= A sequence of Reads and Writes

**NOTE:**

**Most tables in OLTP are indexed!! Thus, index-based access!!**

# Database IO Architecture (OLTP vs. OLAP)

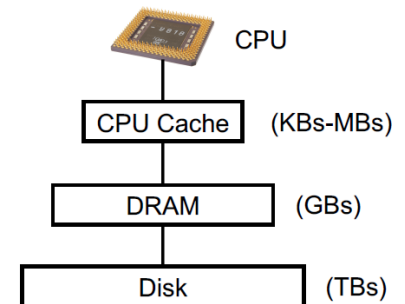
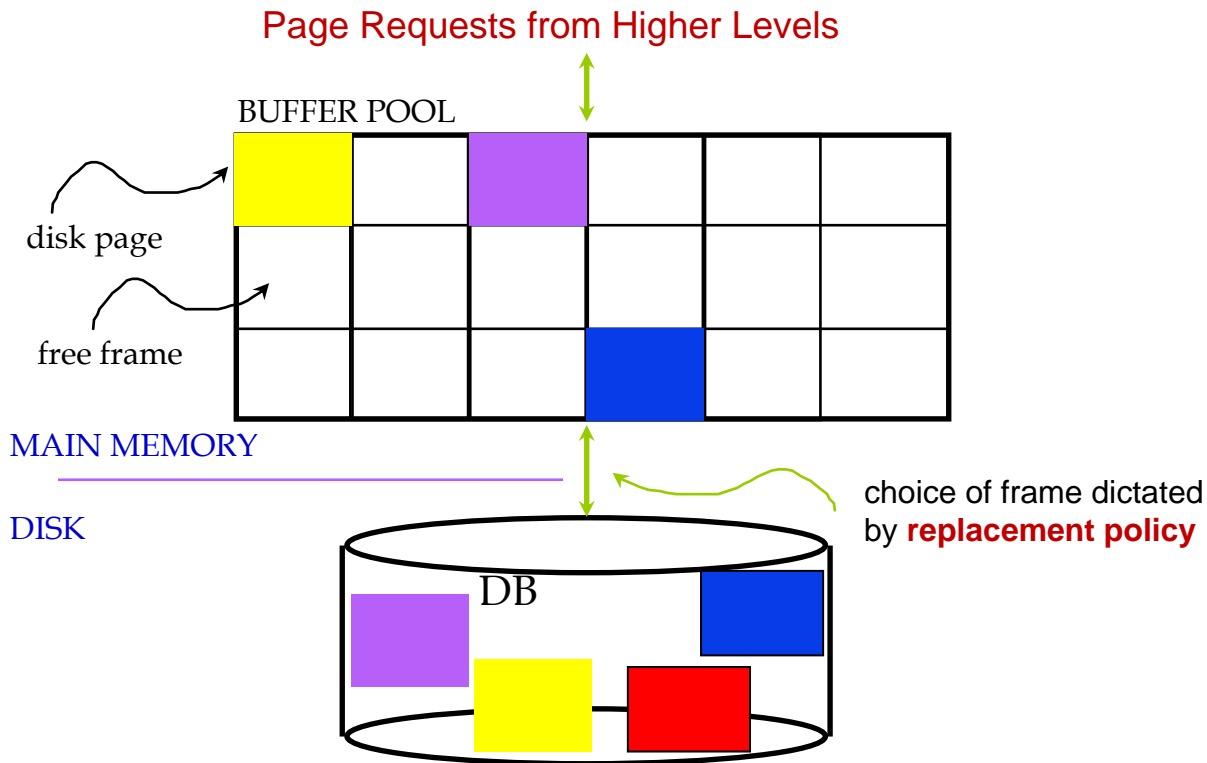


- A transaction consists of SQL statements executed in sequence
- Each SQL (Select/Insert/Delete/Update)
  - Select reads tuples from page(s) while Insert/Delete/Update changes records in page(s)
  - Thus, access one or more pages
- When page(s) are in buffer (i.e., HIT): DRAM operation
- Otherwise (i.e., MISS), IOs are issued
  - read page(s) from the storage
  - In case of dirty victim, write page to the storage
  - Also, checkpoint writes (chap 18)

IO patterns  
: Random vs. Sequential  
: Hot vs. cold page writes

# 9.4 Buffer Management in a DBMS

- Data must be in **RAM** for DBMS to operate on it!



Typically want to cache **frequently accessed data** at a high level of the storage hierarchy to improve performance

Why Small Buffer Cache Works?  
Access Skew and Temporal Locality Exist in the access pattern

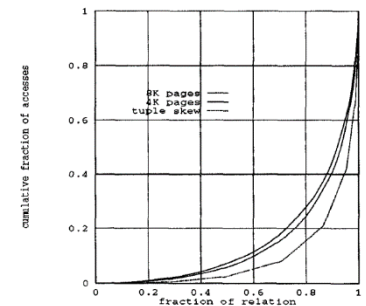


Figure 5: Stock Relation CDF

# When a Page is Requested ...

- Buffer pool information table: <frame#, pageid, pin\_cnt, dirty>
  - In big systems, it is not trivial to just check whether a page is in pool
- If requested page is not in pool:
  - Choose a frame for replacement
  - If frame is dirty, write it to disk
  - Read requested page into chosen frame
- Pin the page and return its address.
- If requests can be predicted (e.g., sequential scans) pages can be pre-fetched several pages at a time!

# More on Buffer Management

- Requestor of page must unpin it, and indicate whether page has been **modified** (using **dirty bit**)
- Page in pool may be requested many times,
  - a **pin\_count** is used.
  - a page is a candidate for replacement iff  $\text{pin\_count} = 0$ .
- CC & recovery may entail additional I/O when a frame is chosen for replacement. (e.g. Write-Ahead Log protocol)

# Buffer Manager Pseudo Code

```
void* BufferManager.getPage( pid int ) {  
    void * frame = search_buffer-for_page(pid);  
    if ( !frame ) {  
        // we get here if the requested page is not in the buffer  
        frame = get_next_empty_frame();  
        if ( !frame ) // and we ran out of space...  
            frame = replacement_policy.choose();  
        if ( frame.dirty ) { // only if we allow to 'steal' a frame  
            write_page(frame);  
            frame.dirty = false;  
        }  
        load(pid, frame);  
    }  
    frame.pin ++;  
    return frame;  
}
```

[Source: Uwe Röhm's Slide]



# Buffer Replacement Policy

- Hit vs. miss
- Hit ratio = # of hits / ( # of page requests to buffer cache)
  - One miss incurs one (or two) physical IO. Hit saves IO.
  - Rule of thumb: at least 80 ~ 90%
- **Problem:** for the given (future) references, which victim should be chosen for highest hit ratio (i.e. least # of IOs)?
  - Numerous policies
  - Does one policy win over the others?
  - One policy does not fit all reference patterns!

## Buffer Replacement Policy (2)

- Frame is chosen for replacement by a **replacement policy**:
  - Random, FIFO, LRU, MRU, LFU, Clock etc.
  - Replacement policy can have big impact on # of I/O's; depends on the [access pattern](#)
- For a given workload, one replacement policy, A, achieves 90% hit ratio and the other, B, does 91%.
  - How much improvement? **1%** or **10%**?
  - We need to interpret its impact in terms of miss ratio, not hit ratio

# Buffer Replacement Policy (3)

- Least Recently Used (LRU)
  - For each page in buffer pool, keep track of time last *unpinned*
  - Replace the frame that has the oldest (earliest) time
  - Very common policy: intuitive and simple
  - Why does it work?
    - ✓ “Principle of (temporal) locality” (of references) ([https://en.wikipedia.org/wiki/Locality\\_of\\_reference](https://en.wikipedia.org/wiki/Locality_of_reference))
    - ✓ Why temporal locality in database?: e.g. hot items, insertion to heap files
  - The correct implementation is not trivial
    - ✓ Especially in large scale systems: e.g. time stamp
- Variants
  - Linked list of buffer frames, LRU-K, 2Q (Linux cache), midpoint-insertion and touch count algorithm (Oracle, MySQL/InnoDB), Clock (MS-SQL), ARC, [LIRS](#) ...
  - Implication of big memory: “random” > “LRU”??

All About Oracle's Touch Count Algorithm  
https://www.oracle.com/technetwork/middleware/arc/arc-114647-01.pdf

# Buffer Replacement Policies (4): Taxonomy

- Recency vs. Frequency

replacement decision based on		Age		
		no	since last usage	since arrival
References	none	RANDOM		FIFO
	last reference		LRU MRU CLOCK	
	all references	LFU	GCLOCK	

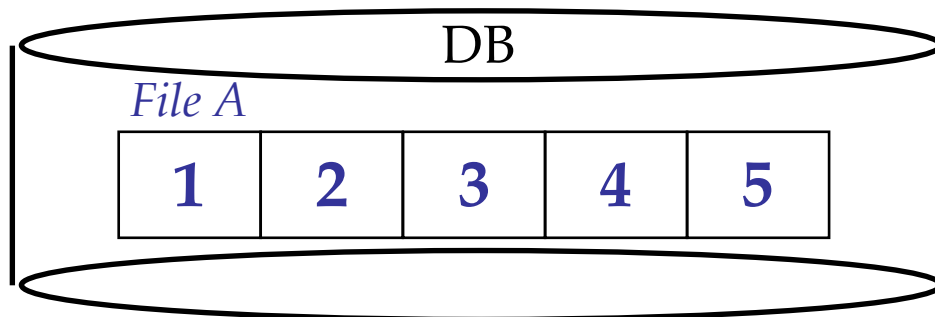
*[Source: Uwe Röhm's Slide]*

# Buffer Replacement Policy (5):

## LRU is **NOT** Scan-Resistant

- Problem of LRU - **sequential flooding**
  - caused by LRU + repeated sequential scans.
  - **# buffer frames < # pages in file** means each page request causes an I/O. MRU much better in this situation (but not in all situations, of course).

BUFFER POOL SIZE: 4 Blocks

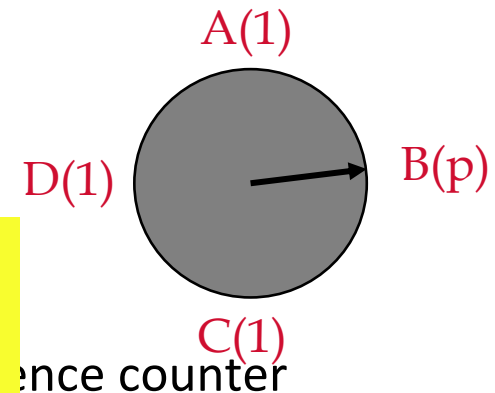


- Assume repeated sequential scans of file A
- What happens when reading 5<sup>th</sup> blocks? and when reading 1<sup>st</sup> block again? ....

# Buffer Replacement Policy (6): “Clock” Algorithm

- An **approximation** of LRU
- Arrange frames into a cycle, store one **reference bit** *per frame*
  - Can think of this as the **2nd chance** bit
- When pin count reduces to 0, turn on reference bit
- When replacement necessary

```
do for each page in cycle {  
    if (pincount == 0 && ref bit is on)  
        turn off ref bit;  
    else if (pincount == 0 && ref bit is off)  
        choose this page for replacement;  
} until a page is chosen;
```



# Buffer Replacement: Hit Ratio vs. Buffer Size

Why Small Buffer Cache Works?  
Access Skew and Temporal Locality  
Exist in the access pattern

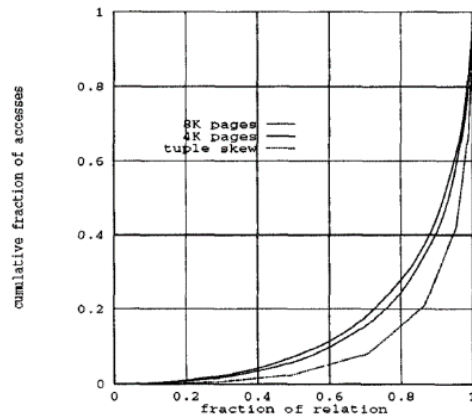
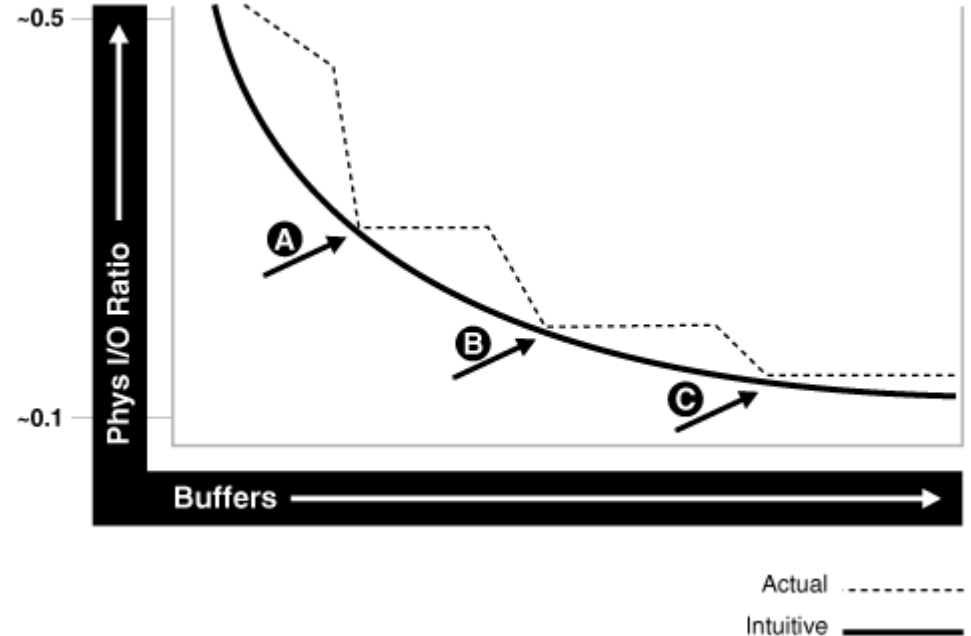


Figure 5: Stock Relation CDF



# OLTP Through the Looking Glass [sigmod 08]

## New Order Transaction

1. Select(whouse-id) from Warehouse
2. Select(dist-id, whouse-id) from District
3. Update(dist-id, whouse-id) in District
4. Select(customer-id, dist-id, whouse-id) from Customer
5. Insert into Order
6. Insert into New-Order
7. For each item (10 items):
  - (a) Select(item-id) from Item
  - (b) Select(item-id, whouse-id) from Stock
  - (c) Update(item-id, whouse-id) in Stock
  - (d) Insert into Order-Line
8. Commit

## New Order

```

begin
for loop(10)
.....Btree lookup(I), pin
Btree lookup(D), pin
Btree lookup (W), pin
Btree lookup (C), pin
update rec (D)
for loop (10)
.....Btree lookup(S), pin
.....update rec (S)
.....create rec (O-L)
.....insert Btree (O-L)
create rec (O)
insert Btree (O)
create rec (N-O)
insert Btree (N-O)
insert Btree 2ndary(N-O)
commit
    
```

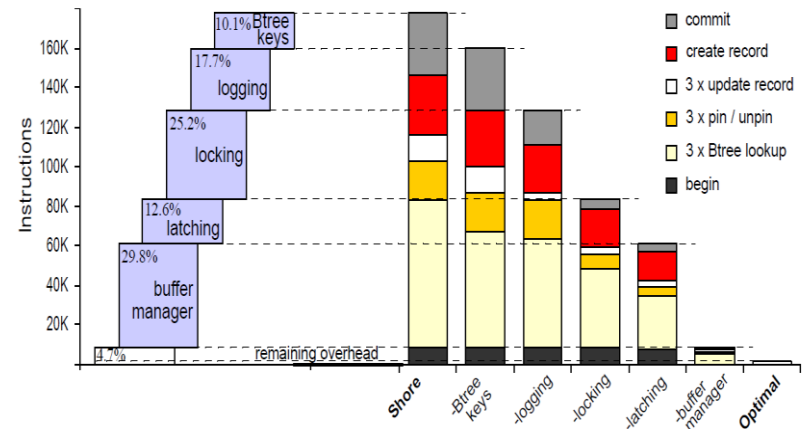


Figure 5. Detailed instruction count breakdown for Payment transaction.

Relation	New Order	Payment	Order Status	Delivery	Stock Level	Average
warehouse	U(1)	U(1)				0.87
district	U(1)	U(1)			P(1)	0.93
customer	NU(1)	NU(2.2)	NU(2.2)	P(10)		1.524
stock	NU(10)				P(200)	12.4
item	NU(10)					4.4
order	A(1)		P(1)	P(10)		0.53
new-order	A(1)			P(10)		0.49
order-line	A(10)		P(10)	P(100)	P(200)	13.3
history		A(1)				0.43



# Belady's MIN Algorithm

- Theoretical optimal buffer replacement algorithm
  - “The *most* efficient caching algorithm would be to always **discard the information that will not be needed for the longest time in the future**. ... Since it is generally impossible to predict how far in the future information will be needed, this is generally not implementable in practice. The practical minimum can be calculated only after experimentation, and one can compare the effectiveness of the actually chosen cache algorithm.” ([https://en.wikipedia.org/wiki/Cache\\_algorithms#B.C3.A9I.C3.A1dy.27s\\_Algorithm](https://en.wikipedia.org/wiki/Cache_algorithms#B.C3.A9I.C3.A1dy.27s_Algorithm))
- Offline algorithm (vs. online algorithm)
  - “All practical solutions are **attempts to approximate the optimal Belady's MIN policy**” (from “Principles of Operating Systems: Design and Applications”, Brian L. Stuart)
- Belady's anomaly
  - If the number of page frames is increased, would always the hit ratio will be higher or at least same?

# Belady's Anomaly

- FIFO Algorithm: 3 vs. 4 page frames

All pages frames initially empty

	0	1	2	3	0	1	4	0	1	2	3	4
Youngest page		0	1	2	3	0	1	4	4	4	2	3
			0	1	2	3	0	1	1	1	4	2
Oldest page				0	1	2	3	0	0	0	1	4
	P	P	P	P	P	P	P				P	P

9 Page faults

(a)

		0	1	2	3	0	1	4	0	1	2	3	4
Youngest page		0	1	2	3	3	3	4	0	1	2	3	4
			0	1	2	2	2	3	4	0	1	2	3
Oldest page				0	1	1	1	2	3	4	0	1	2
					0	0	0	1	2	3	4	0	1
		P	P	P	P			P	P	P	P	P	P

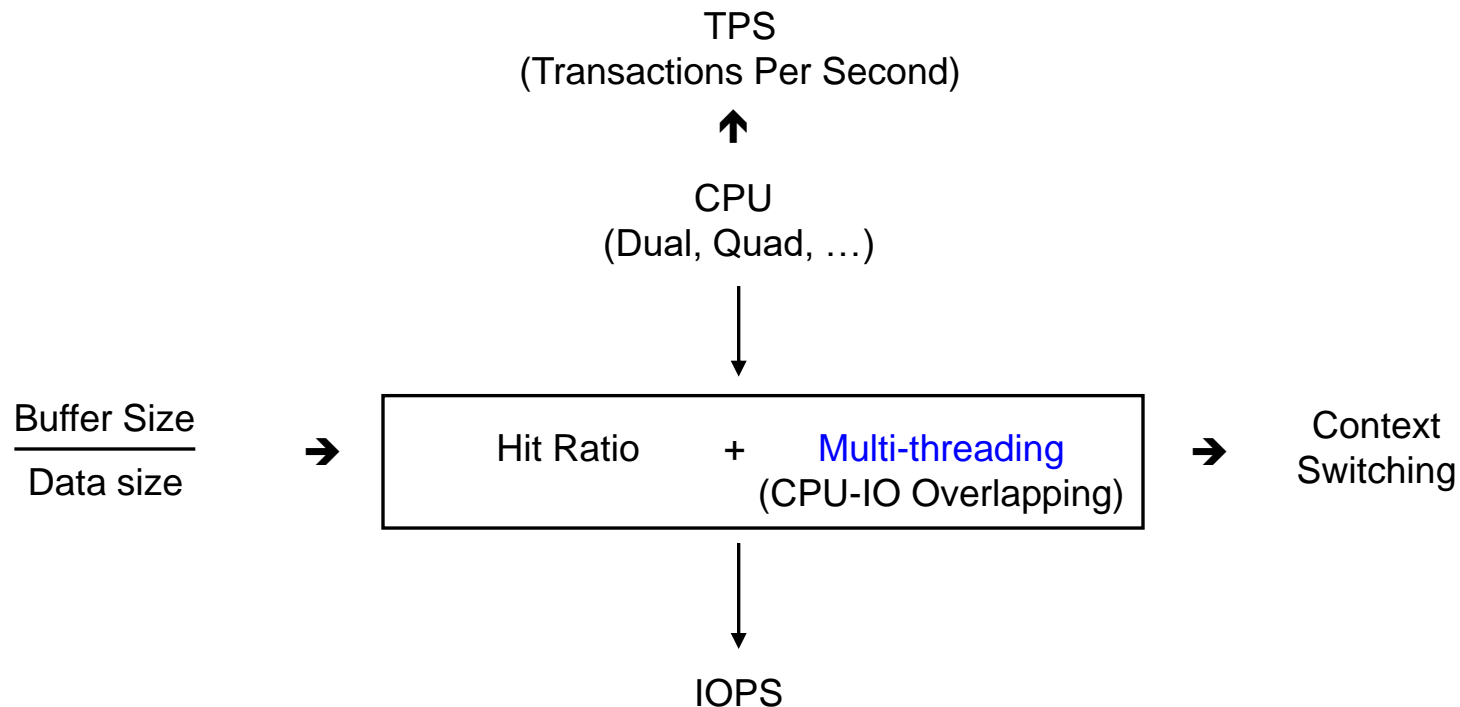
10 Page faults

(b)

Source: <https://www.cs.ucsb.edu/~chris>

# CPU-IO Overlapping

- 3 States: CPU Bound, IO Bound, Balanced




**For perfect CPU-IO overlapping,  
IOPS Matters!**

# IOPS Crisis in OLTP

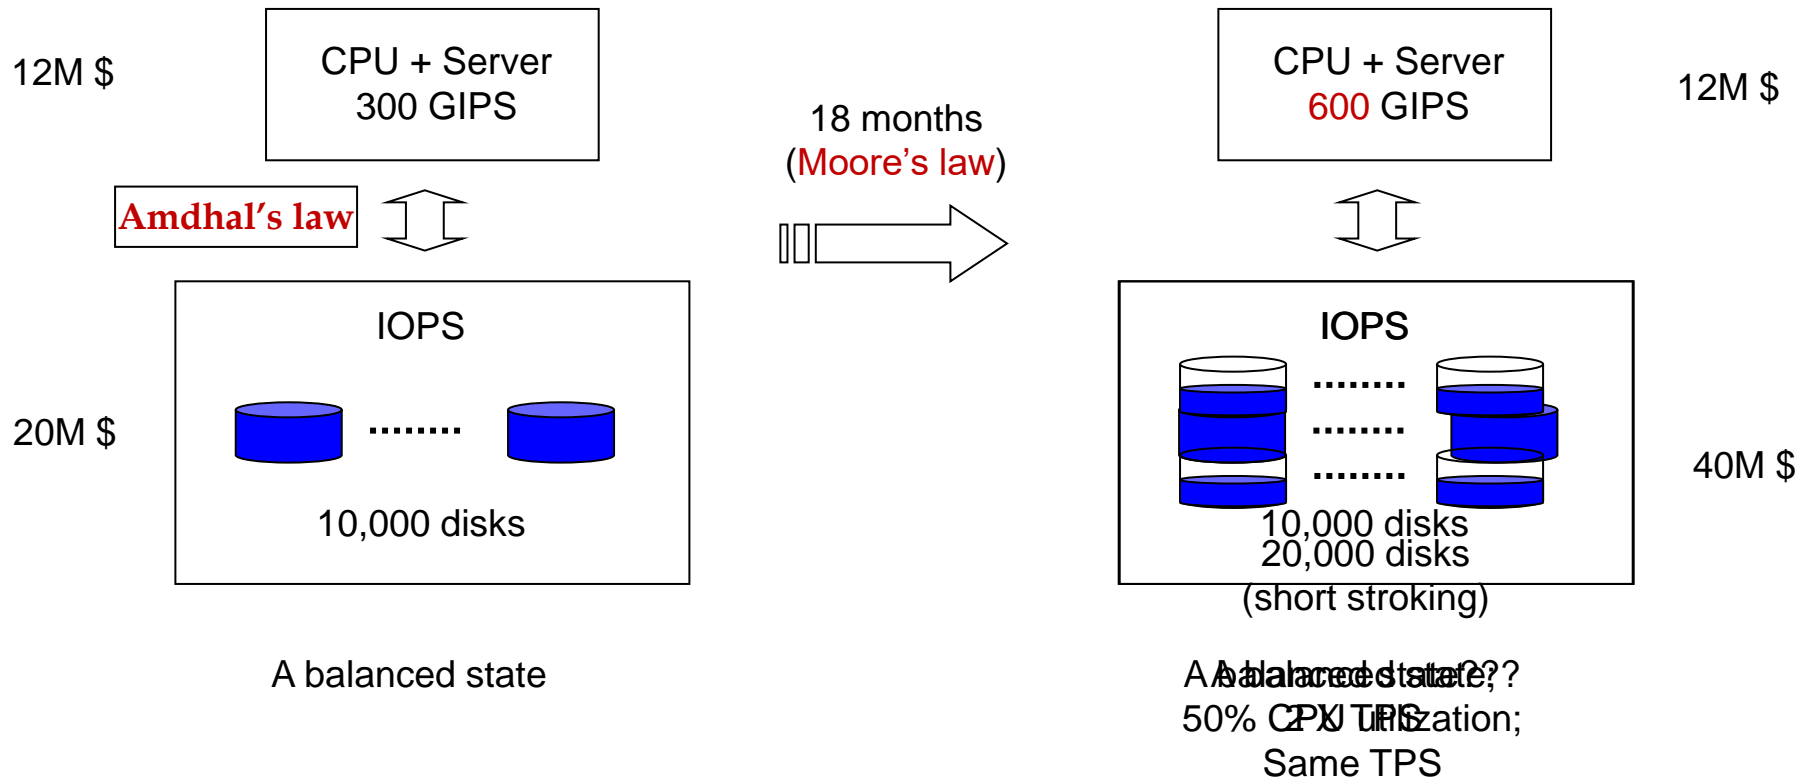
## IBM for TPC-C (2008 Dec.)

- 6M tpmC
- Total cost: **35M \$**
  - Server HW: 12M \$
  - Server SW: 2M \$
  - Storage: **20M \$**
  - Client HW/SW: 1M \$
- They are buying **IOPS**, **not** capacity

IBM		IBM Power 595 Model 9119-FHA		TPC-C Rev. 5.9	
				Report Date: June 10, 2008	
Total System Cost	TPC-C Throughput	Price/Performance		Availability Date	
\$17,111,788 USD	6,085,166	\$2.81 USD		December 10, 2008	
Database Server Processor Chip/Core/Thread	Database Manager	Operating System	Other Software	No. Users	
32/64/128 POWER6 5.0GHz	DB2 9.5	AIX 5L V5.3	Microsoft Visual C++ Microsoft COM+	5,184,000	
<div><div><p><b>128 Clients</b> IBM System x3550 Dual-core 2.0GHz Intel® Xeon™ 4MB L2 Cache 1GB Memory 73.4GB SAS Drive Integrated dual-port 10/100/1000 Ethernet</p><p><b>IBM Power 595 Server</b> 32 Processor Chips with 64 5.0GHz POWER6™ Cores 32MB L3 Cache per chip 4096GB Memory 8 146GB Internal SCSI Drives 68 4Gb dual-port Fibre Channel Adapters 10 10/100/1000 Ethernet Adapters</p><p><b>Storage</b> 68 IBM System Storage DS4800 784 IBM System Storage DS4000 EXP810 10,992 73.4GB 15K RPM 4Gbps Drives</p></div></div>					
System Components		Server		Each of the 128 Clients	
	Quantity	Description	Quantity	Description	
Processors Chips /Cores/Threads	32/64/128	5.0GHz POWER6	1/2/2	2.0GHz 4MB L2 Xeon Processor	
Memory	64	64GB	2	512 MB	
Disk Controllers	1 68 68	SAS Controller 4Gb FC Adapters IBM System Storage DS4800	1	SAS Controller	
Disk Drives	8 10,992	146.8GB 15K RPM SCSI 73.4GB 15K RPM 4Gb FC	1	73.4GB 15K RPM SAS	
Total Storage		746,467GB		67.86GB	
Terminals	1	System Console	1	System Console	

# IOPS Crisis in OLTP(2)

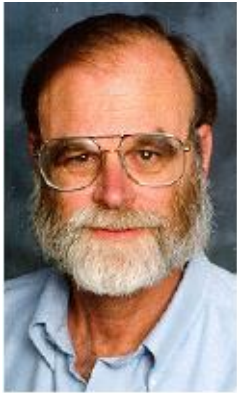
- For **balanced systems**, OLTP systems pay \$\$\$ on disks for high IOPS
  - vs. IO-bound, CPU-bound



# Flash Disk Opportunity for Server Applications

## ([MS-TR 2007](#))

What If FLASH Disks Delivered Thousands of IO/s and Were “Big”?



Jim Gray

My tests and those of several others suggest that FLASH disks can deliver about 3K random 8KB reads/second and with some re-engineering about 1,100 random 8KB writes per second. Indeed, it appears that a single FLASH chip could deliver nearly that performance and there are many chips inside the “box” – so the actual limit could be 4x or more. But, even the current performance would be VERY attractive for many enterprise applications. For example, in the TPC-C benchmark, has approximately equal reads and writes. Using the graphs above, and doing a weighted average of the 4-deep 8 KB random read rate (2,804 IOps), and 4-deep 8 KB sequential write rate (1233 IOps) gives *harmonic average* of 1713 (1-deep gives 1,624 IOps). TPC-C systems are configured with ~50 disks per cpu. For example the most recent [Dell TPC-C system](#) has ninety 15Krpm 36GB SCSI disks costing 45k\$ (with 10k\$ extra for maintenance that gets “discounted”). Those disks are 68% of the system cost. They deliver about 18,000 IO/s. That is comparable to the requests/second of ten FLASH disks. So we could replace those 90 disks with ten NSSD if the data would fit on 320GB (it does not). That would save a lot of money and a lot of power (1.3Kw of power and 1.3Kw of cooling).

The current flash disks are built with 16 Gb NAND FLASH. When, in 2012, they are built with a 1 terabit part, the device will have 2TB of capacity and will indeed be able to store the TPC-C database. So we could replace a 44k\$ disk array with a few (say 10) 400\$ flash disks (maybe).

If one looks at the system diagram of the Samsung NSSD there are many opportunities for innovation. It suggests interesting RAID options for fault tolerance (combining the MSR-TR-2006-176 ideas with non-volatile storage map and a block-buffer, and with writing raid-5 stripes of data across the chip array) adding a battery, adding logic for copy-on-write snapshots, and so on. These devices enable whole new approaches to file systems. They are potential gap fillers between disks and RAM and they are interesting “hot data” storage devices in their own right.

“My tests and those of several others suggest that FLASH disks can deliver about 3K random 8KB reads/second and with some re-engineering about 1,100 random 8KB writes per second. Indeed, it appears that a single FLASH chip could deliver nearly that performance and there are many chips inside the “box” – so the actual limit could be 4x or more. But, even the current performance would be VERY attractive for many enterprise applications. For example, in the TPC-C benchmark, has approximately equal reads and writes. Using the graphs above, and doing a weighted average of the 4-deep 8 KB random read rate (2,804 IOps), and 4-deep 8 KB sequential write rate (1233 IOps) gives *harmonic average* of 1713 (1-deep gives 1,624 IOps). TPC-C systems are configured with ~50 disks per cpu. For example the most recent [Dell TPC-C system](#) has ninety 15Krpm 36GB SCSI disks costing 45k\$ (with 10k\$ extra for maintenance that gets “discounted”). Those disks are 68% of the system cost. They deliver about 18,000 IO/s. That is comparable to the requests/second of ten FLASH disks. So we could replace those 90 disks with ten NSSD if the data would fit on 320GB (it does not). That would save a lot of money and a lot of power (1.3Kw of power and 1.3Kw of cooling).” (excerpts from “Flash disk opportunity for server applications”)


# Our Message in SIGMOD 2009

## One FlashSSD can beat Ten Harddisks in OLTP - Performance, Price, Capacity, Power – (in 2008)

- In 2015, one FlashSSD can beat several 10x or more HDDs in OLTP
- “My tests and those of several others suggest that FLASH disks can deliver about 3K random 8KB reads/second and with some re-engineering about 1,100 random 8KB writes per second. Indeed, it appears that a single FLASH chip could deliver nearly that performance and there are many chips inside the “box” – so the actual limit could be 4x or more. But, even the current performance would be VERY attractive for many enterprise applications. For example, in the TPC-C benchmark, has approximately equal reads and writes. Using the graphs above, and doing a weighted average of the 4-deep 8 KB random read rate (2,804 IOps), and 4-deep 8 KB sequential write rate (1233 IOps) gives *harmonic average* of 1713 (1-deep gives 1,624 IOps). TPC-C systems are configured with ~50 disks per cpu. For example the most recent [Dell TPC-C system](#) has ninety 15Krpm 36GB SCSI disks costing 45k\$ (with 10k\$ extra for maintenance that gets “discounted”). Those disks are 68% of the system cost. They deliver about 18,000 IO/s. That is comparable to the requests/second of ten FLASH disks. So we could replace those 90 disks with ten NSSD if the data would fit on 320GB (it does not). That would save a lot of money and a lot of power (1.3Kw of power and 1.3Kw of cooling).” (excerpts from Flash disk opportunity for server-applications (Jim Gray)

# Flash-based TPC-C @ 2013 September

- Oracle + Sun Flash Storage
  - 8.5M tpmC
- Total cost: **4.7M \$**
  - Server HW: .6M \$
  - Server SW: 1.9M \$
  - Storage: **1.8M \$**
    - ✓ 216 400GB Flash Module: **1.1M \$**
    - ✓ 86 3TB 7.2K HDD : **0.07M**
  - Client HW/SW: 0.1M \$
  - Others: 0.1M\$
- Implications
  - More vertical stacks (by SW vendor )
  - Harddisk vendors (e.g. Seagate)

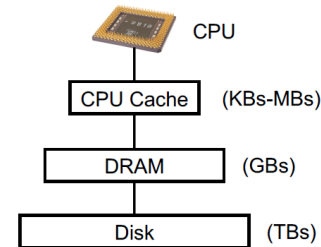
ORACLE		SPARC T5-8 Server		TPC-C 5.11.0 TPC-Pricing 1.7.0
Total System Cost		TPC-C Throughput	Price/Performance	Report Date March 26, 2013
\$4,663,073USD		8,552,523 tpmC	\$0.55USD/tpmC	Availability Date <b>September 25, 2013</b>
Database Server Processors/Cores/Threads	Database Manager	Operating System	Other Software	Number of Users
SPARC T5 3.6GHz 8 / 128 / 1,024	Oracle Database 11g Release 2 Enterprise Edition With Oracle Partitioning	Oracle Solaris 11.1	Oracle Tuxedo CFSR Oracle Web Tier 1	6,800,000
<div> <div> <b>Clients</b>                      16 Sun Server X3-2                      2 Intel® Xeon®                      E5-2690 2.93GHz                      64GB Memory                      2 600GB SAS disk                 </div> <div> <b>Database System</b>   <b>SPARC T5-8 Server</b>                      8 SPARC T5 3.6GHz                      4TB Memory                      2 600GB 10K RPM SAS                      12 8Gb/s FC HBA, 2 port                      10GbE SFP+                 </div> <div> <b>Storage</b>  <b>54 DATA COMSTAR</b>                      Sun Server X3-2L w/                      2 Intel® Xeon® E5-2609 2.4GHz                      16GB DDR3                      46 w/ 1 3TB 7.2K RPM SAS                      8 w/ 2 3TB 7.2K RPM SAS                      4 F40 PCI-E 400GB SSD    <b>2 REDO COMSTAR</b>                      Sun Server X3-2L w/                      2 Intel® Xeon® E5-2609 2.4GHz                      16GB DDR3                      12 3TB 7.2K RPM SAS                 </div> </div>				
System Component	Each Server Node		Each Client	
Processors/Cores/Threads and cache	8/128/1024	SPARC T5 3.6GHz 8MB L3 Cache	2/16/32	Intel® Xeon® E5-2690 12MB Smart Cache
Memory		4TB		64GB
Disk Controllers	12	8Gb/s FC HBA 2 Port	1	8 port Internal SAS
OS Disks (each system)	2	600GB 10K RPM SAS	2	600GB 10K RPM SAS
External Storage	216	Flash Accelerator F40 PCI-E 400GB eMLC		
	86	3TB 7.2K RPM SAS		
Total Storage		344.4TB		



# MMDBMS vs. All-Flash DBMS: Personal Thoughts

- Why MMDBMS has been recently popular?
  - Sap Hana, MS Hekaton, [Oracle In-memory](#), Altibase, ....
  - The price of DRAM had ever dropped for the last two decades
    - ✓  $\$/\text{IOPS} @ \text{DISK} \gg \$/\text{GB} @ \text{DRAM}$
  - The overhead of disk-based DBMS is not negligible
  - Applications with extreme performance requirements?
- Flash storage
  - Lowered  $\$/\text{IOPS}$
  - $\$/\text{IOPS} @ \text{SSD} \ll \$/\text{GB} @ \text{DRAM}$
- MMDBMS vs. **All-Flash** DBMS (with some optimizations)
  - Winner? Time will tell: read “[Umbra: A Disk-Based System with In-Memory Performance@CIDR2020](#)”
  - But, in 2017, the average DRAM DDR4 price has increased by 2.3.
    - ✓ “Reducing DRAM footprint with NVM in Facebook” (Eurosys 2018)

# Buffer Management vs. 5 Min Rule & Pareto's Law



## Five minute rule

- “The 5 Minute Rule for Trading Memory Accesses for Disc Accesses”, SIGMOD record 1985;
- See also “The Five-minute Rule Thirty Years Later and its Impact on the Storage Hierarchy”, ADMS, 2017.
- “How much high-tier storage size should we have?”
  - ✓ Can determine based on data’s access frequency and devices’ cost
- @1987: If a data object is accessed in every **5 minutes or more frequently**, it should be memory-resident, not in harddisk @ 1987
- @2017: **3 hour rule** with harddisk; this justifies MMDBMS
- @2017: **10 second rule** with low latency flash SSD

## Pareto’s law: 80/20 rules

- E.g. A Modeling Study of the TPC-C Benchmark, SIGMOD 1993
  - ✓ **84 % of accesses go to 20 % of the tuples.**

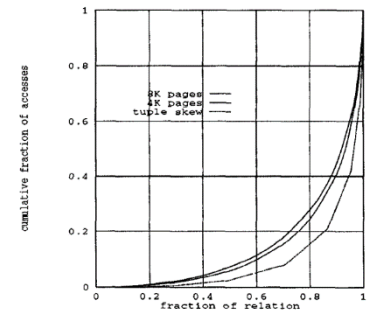
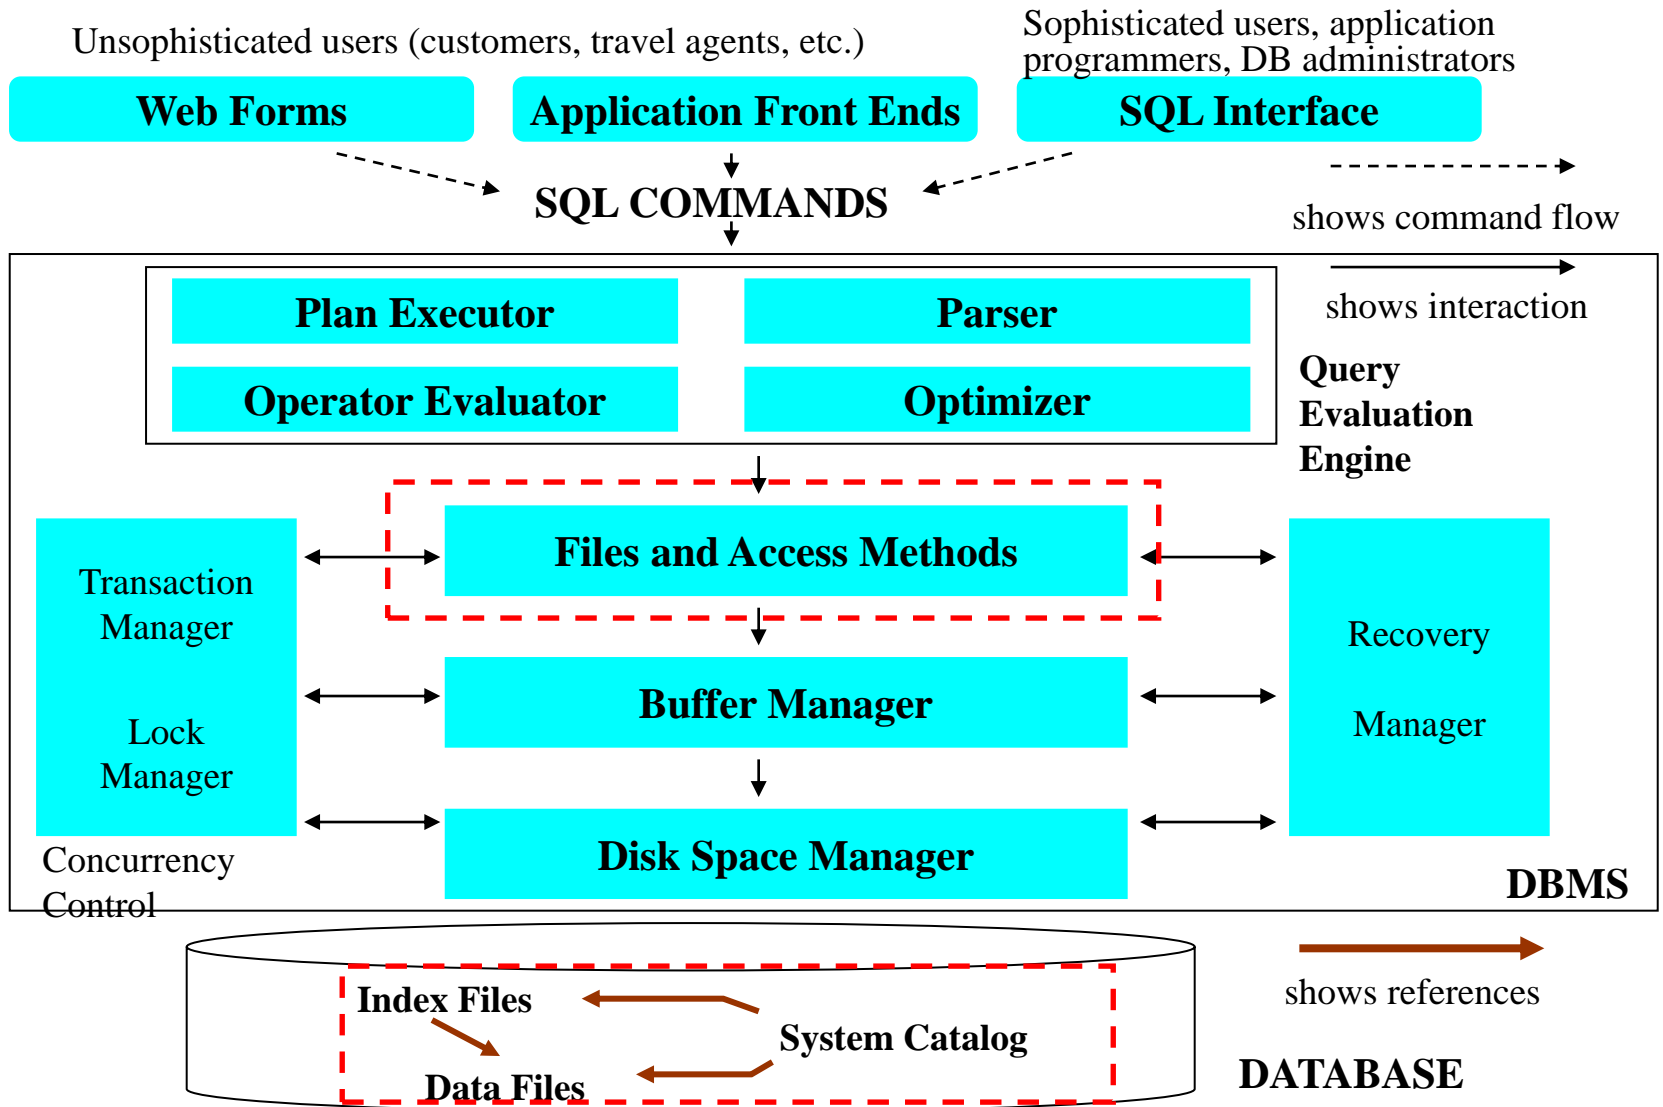


Figure 5: Stock Relation CDF

# DBMS vs. OS File System

OS does disk space & buffer mgmt: **why not** let OS manage these tasks? (See appendix slides for more issues)

- Differences in OS support: portability issues
- Some limitations, e.g., files can't span disks.
  - note, this is changing --- OS File systems are getting smarter (i.e., more like databases!)
- Buffer management in DBMS requires ability to:
  - **pin** a page in buffer pool, **force** a page to disk (important for implementing CC & recovery),
  - adjust replacement policy, and pre-fetch pages based on **access patterns** in typical DB operations.



**Figure 1.3 Architecture of a DBMS**

# How to Physically Store a Logical Table in Disk?

- Logical database level
  - DB: a set of relations, Relation: a set of records(or tuples), Tuple: a sequence of attributes
  - e.g. 

```
CREATE TABLE EMP (  
    empno NUMBER PRIMARY KEY,  
    ename VARCHAR(30),  
    ...  
);
```
- Physical database level
  - How to represent SQL data types, a (variable-length) tuple with several attributes, a collection of tuples in a page? A set of pages in a relation / index?
  - How to handle insert, update, delete?

## 9.5 Files of Records (or Tuples/Rows)

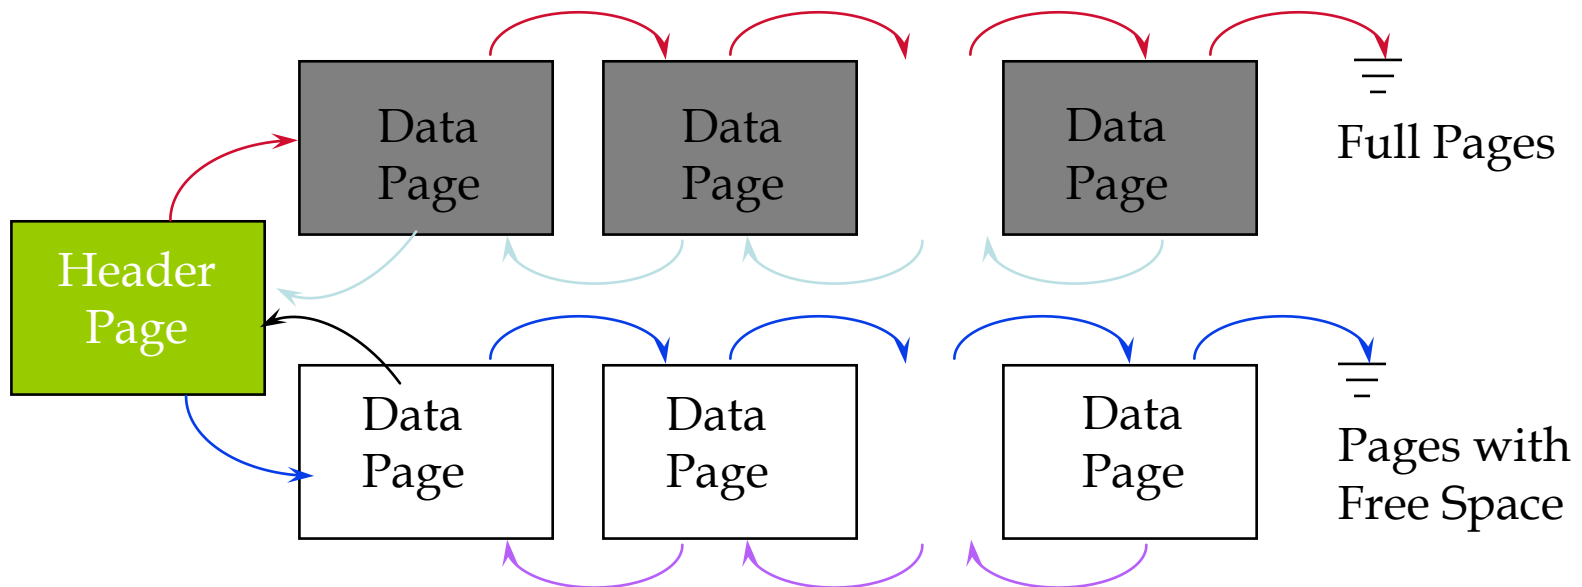
- **Blocks interface for I/O** between buffer manager and disk space manager, but...
- Higher levels of DBMS operate on **records**, and **files of records**.
- **FILE** (per table): a collection of pages, each containing a collection of records (which in general belong to same table). Must support the following **operations**:
  1. Insert/delete/modify record
  2. Fetch a particular record (specified using **record id**)
  3. Scan all records (possibly with **some conditions** on the records to be retrieved)

# Unordered (**Heap**) Files

- Simplest file structure contains records **in no particular order**.
  - As a file grows and shrinks, disk pages are allocated and de-allocated.
- To support record level operations, we must:
  - keep track of the **pages** in a file
  - keep track of **free space** on pages
  - keep track of the **records** on a page
- There are many alternatives for keeping track of this

# Heap File: Implemented as a List

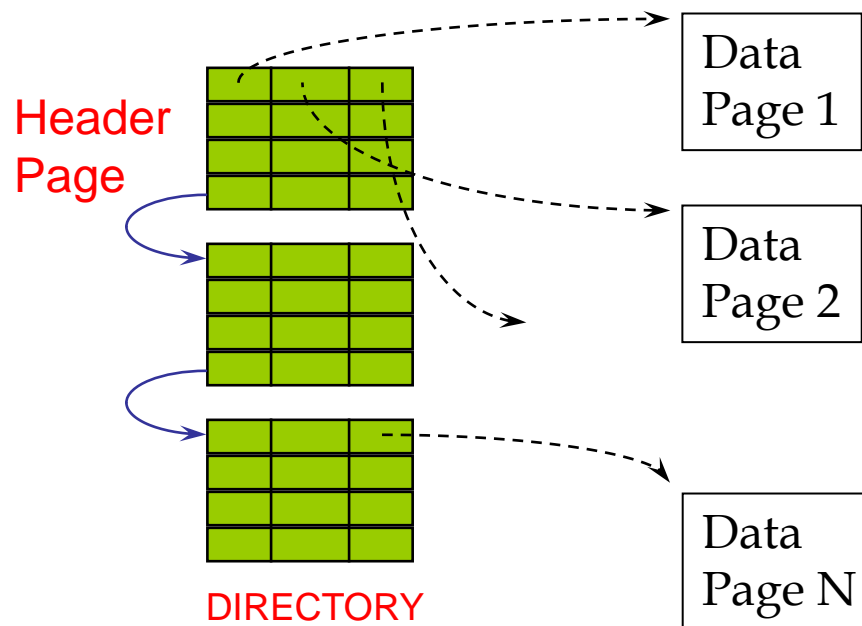
- The header page id and heap file name must be stored someplace: [database catalog](#)
- Each page contains 2 'pointers' plus data.
- **A critical problem:** inefficient to find a page for record insertion.





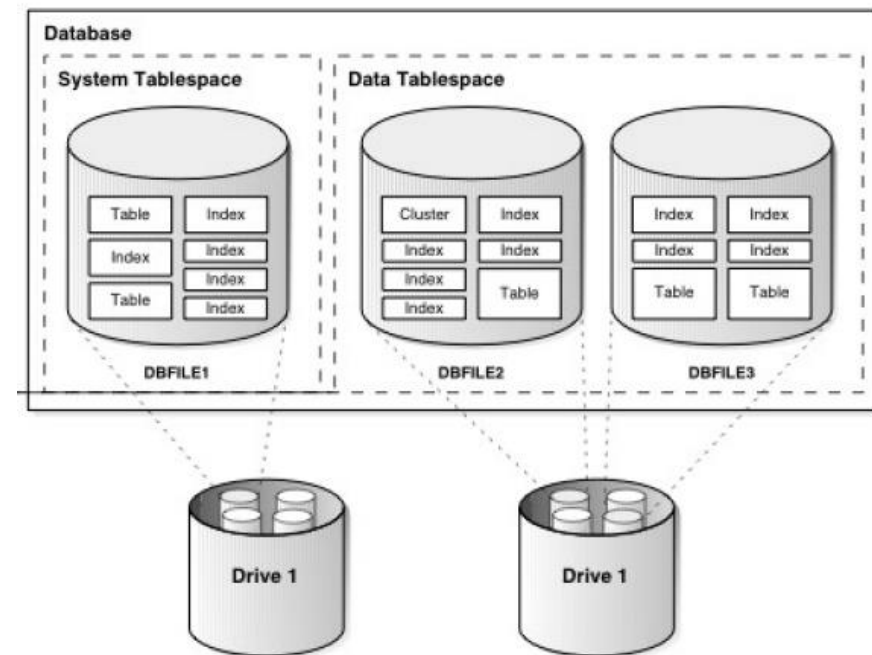
# Heap File: Using a Page Directory

- The entry for a page can include # of free bytes on the page.
- The directory is a collection of pages; linked list implementation is just one alternative.
  - Much smaller than linked list of all HF pages!

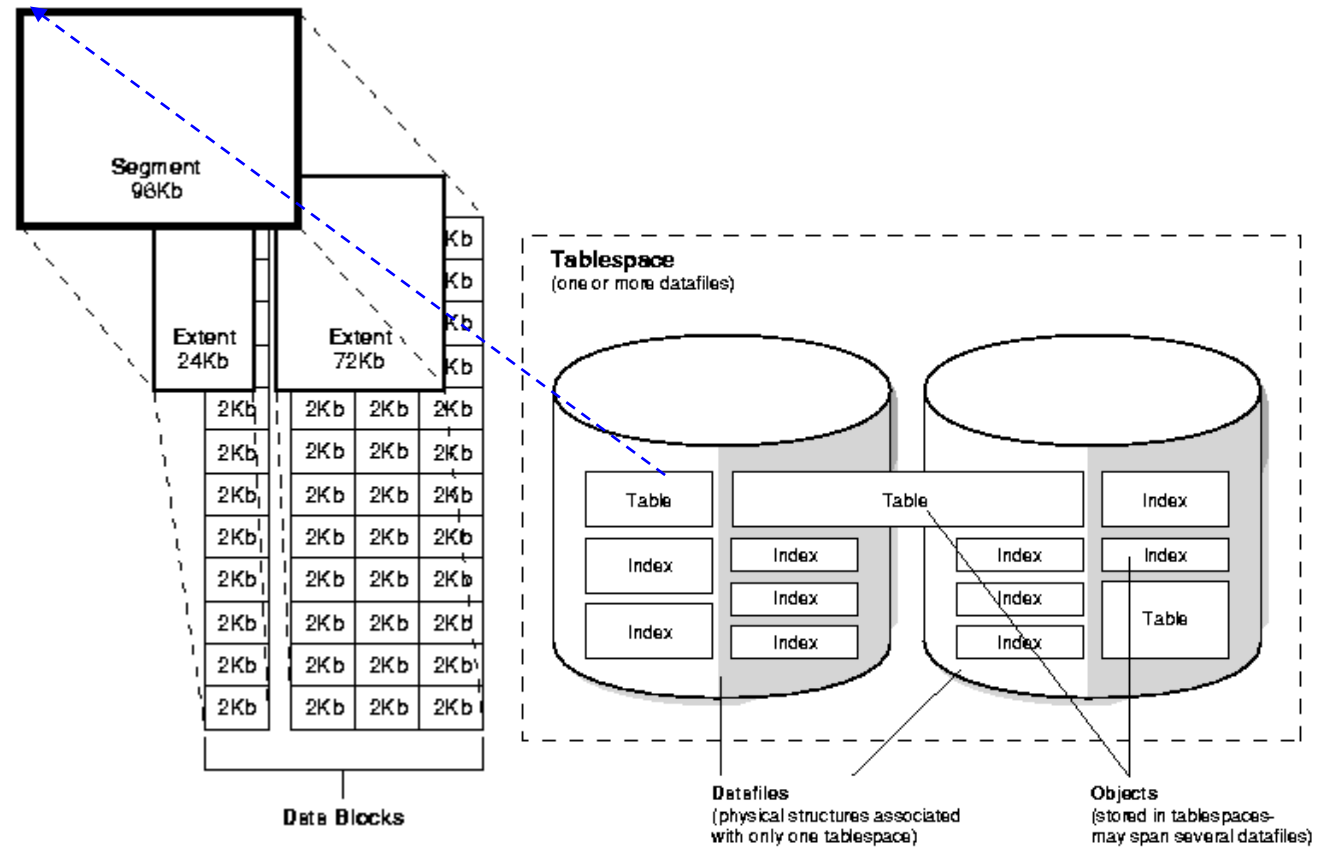


# Oracle: Tablespace, Segments, Extents, and Blocks

- **Tablespace** as logical DBMS 'file'
  - Consist of several physical files in the file system; may span several disks
- Separate data **segments** for
  - each table, index,...
- Data segments and hence Tablespaces can grow (by **extents**))
- An **extent** is the unit of disk space allocation: a sequence of disk blocks
- Rows/tuples are stored on disk blocks (or pages)



# Oracle: Segments, Extents, and Blocks (2)



# Oracle: Segments, Extents, and Blocks (3)

```
COLUMN table_name FORMAT a10  
COLUMN tablespace_name FORMAT a10
```

```
select table_name, tablespace_name, blocks, pct_free, avg_row_len, avg_space  
from user_tables  
where table_name = 'TEST'; /* TEST table in Ch9-Script.sql */
```

TABLE_NAME	TABLESPACE_NAME	BLOCKS	PCT_FREE	AVG_ROW_LEN	AVG_SPACE
TEST	USERS	1000	10	664	1410

# Oracle: Segments, Extents, and Blocks (4)

```
conn scott/tiger as sysdba
```

```
COLUMN segment_name FORMAT a10
```

```
COLUMN segment_type FORMAT a10
```

```
SELECT segment_name, segment_type, header_file, header_block,  
       blocks, extents, max_extents
```

```
FROM dba_segments
```

```
WHERE segment_name = 'TEST';
```

SEGMENT_NAME	SEGMENT_TYPE	HEADER_FILE	HEADER_BLOCK	BLOCKS	EXTENTS	MAX_EXTENTS
TEST	TABLE	4	363	1024	23	2147483645

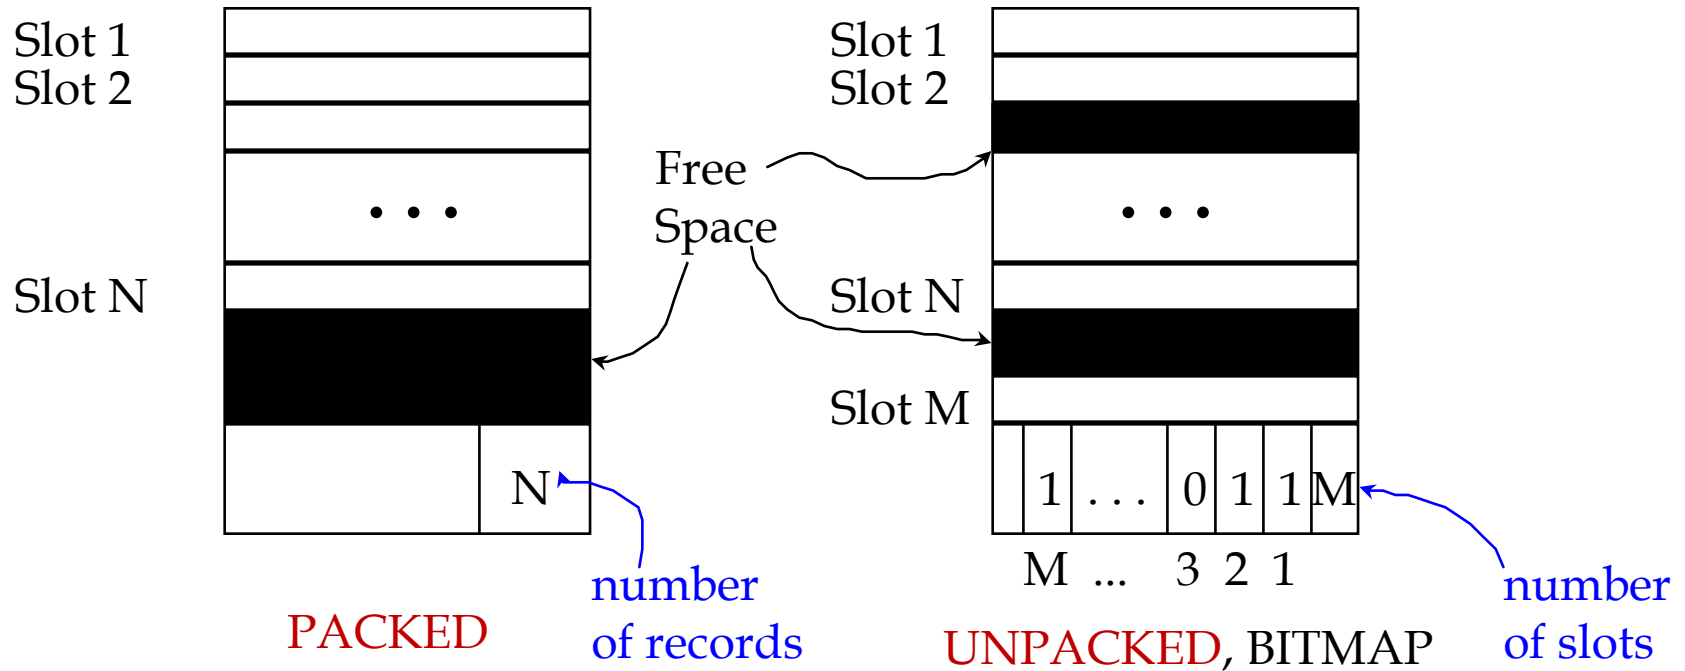
# Oracle: Segments, Extents, and Blocks (5)

```
COLUMN segment_name FORMAT a10  
COLUMN tablespace_name FORMAT a10
```

```
SELECT segment_name, tablespace_name, extent_id, file_id, block_id, blocks  
FROM    dba_extents  
WHERE   segment_name = 'TEST';
```

SEGMENT_NAME	TABLESPACE	EXTENT_ID	FILE_ID	BLOCK_ID	BLOCKS
TEST	USERS	0	4	361	8
TEST	USERS	1	4	369	8
TEST	USERS	2	4	377	8
.....					
TEST	USERS	14	4	473	8
TEST	USERS	15	4	489	8
TEST	USERS	16	4	521	128
TEST	USERS	17	4	649	128
.....					
TEST	USERS	21	4	1161	128
TEST	USERS	22	4	1289	128

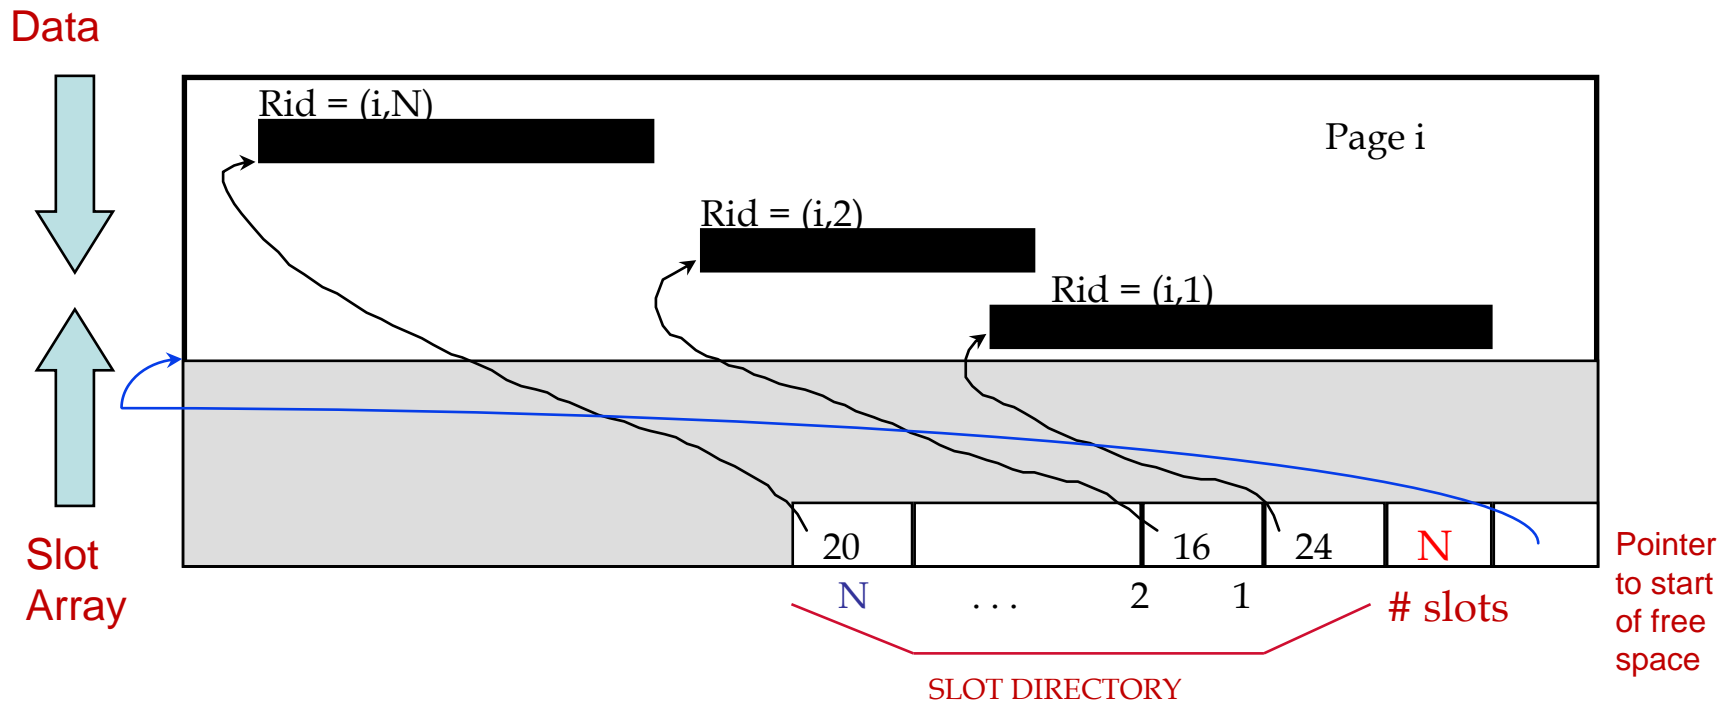
## 9.6 Page Formats: Fixed Length Records



- record id = <page id, slot #> /\* can fetch a record using its record\_id \*/
  - In first alternative (that is, left figure), moving records for free space management changes rid; may not be **acceptable**
    - ✓ why? external references

# Page Formats: Variable Length Records

## - *Tuple-Oriented, Slotted Page Structure* -



- Can move records on page without changing rid; so, attractive for fixed-length records too.
  - each slot entry = **<record offset, length>**
- Page is **full** when data space and slot array **meet**.



# Block Dump in Oracle

```
conn scott/tiger as sysdba
```

```
select header_file, header_block, bytes, blocks  
from dba_segments  
where segment_name = 'EMP';
```

HEADER_FILE	HEADER_BLOCK	BYTES	BLOCKS
-----	-----	-----	-----
4	27	65536	8

```
alter system dump datafile 4 block min 27 block max 30;  
--- Check the trace file in admin/udump/xxx.trc file.
```

# Oracle Formatted Block Dumps

**data\_block\_dump**, data header at 0x3041074

=====

tsiz: 0x1f88

hsiz: 0x2e

pbl: 0x03041074

bdba: 0x0040c652

76543210

flag=-----

ntab=1

nrow=14

frre=-1

fsbo=0x2e

fseo=0x1d51

avsp=0x1d23

tosp=0x1d23

0xe:pti[0]                    nrow=14    offs=0

0x12:pri[0]                  offs=0x1f62

0x14:pri[1]                  offs=0x1f37

0x16:pri[2]                  offs=0x1f0c

0x18:pri[3]                  offs=0x1ee3

. . . . .

0x22:pri[8]                  offs=0x1e16

0x24:pri[9]                  offs=0x1deb

0x26:pri[10]                offs=0x1dc5

0x28:pri[11]                offs=0x1d9f

0x2a:pri[12]                offs=0x1d78

0x2c:pri[13]                offs=0x1d51

**block\_row\_dump:**

tab 0, row 0, @0x1f62

**t1: 38** fb: --H-FL-- lb: 0x0 cc: 8

col 0: [ 3] c2 4a 46

col 1: [ 5] 53 4d 49 54 48

col 2: [ 5] 43 4c 45 52 4b

col 3: [ 3] c2 50 03

col 4: [ 7] 77 b4 0c 11 01 01 01

col 5: [ 2] c2 09

col 6: \*NULL\*

col 7: [ 2] c1 15

. . . . .

tab 0, row 13, @0x1d51

**t1: 39** fb: --H-FL-- lb: 0x0 cc: 8

col 0: [ 3] c2 50 23

col 1: [ 6] 4d 49 4c 4c 45 52

col 2: [ 5] 43 4c 45 52 4b

col 3: [ 3] c2 4e 53

col 4: [ 7] 77 b6 01 17 01 01 01

col 5: [ 2] c2 0e

col 6: \*NULL\*

col 7: [ 2] c1 0b

end\_of\_block\_dump

End dump data blocks tsn: 0 file#: 1

minblk 50769 maxblk 50770

# Page Size

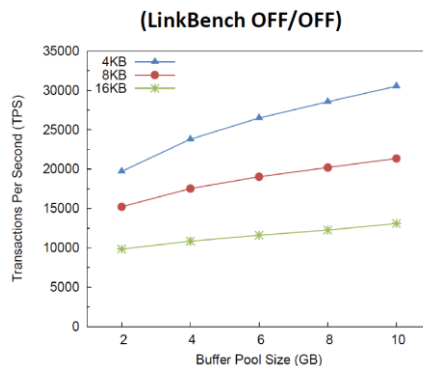
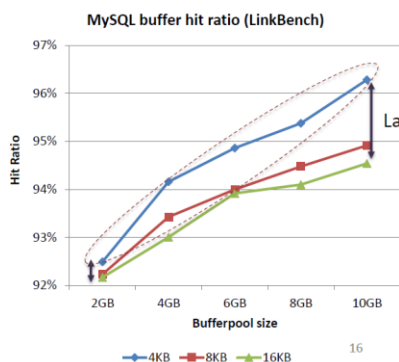
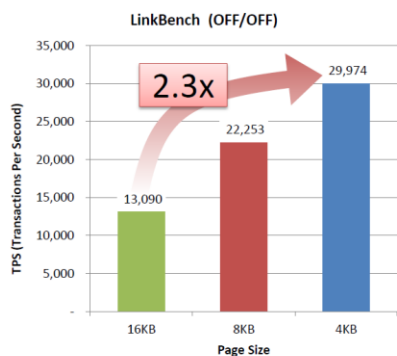
- Oracle
  - “Oracle recommends smaller Oracle Database block sizes (2 KB or 4 KB) for online transaction processing or mixed workload environments and larger block sizes (8 KB, 16 KB, or 32 KB) for decision support system workload environments.” (see chapter “IO Config. And Design” in “Database Performance Tuning Guide” book )

*Table 17-3 Block Size Advantages and Disadvantages*

Block Size	Advantages	Disadvantages
Smaller	<p>Good for small rows with lots of random access.</p> <p>Reduces block contention.</p>	<p>Has relatively large space overhead due to metadata (that is, block header).</p> <p>Not recommended for large rows. There might only be a few rows stored for each block, or worse, row chaining if a single row does not fit into a block,</p>
Larger	<p>Has lower overhead, so there is more room to store data.</p> <p>Permits reading several rows into the buffer cache with a single I/O (depending on row size and block size).</p> <p>Good for sequential access or very large rows (such as LOB data).</p>	<p>Wastes space in the buffer cache, if you are doing random access to small rows and have a large block size. For example, with an 8 KB block size and 50 byte row size, you waste 7,950 bytes in the buffer cache when doing random access.</p> <p>Not good for index blocks used in an OLTP environment, because they increase block contention on the index leaf blocks.</p>

# Page Size (2): e.g. InnoDB Page Tuning on SSD: 16KB -> 4KB

- Why better with smaller 4KB on SSD?
  - Higher hit ratio; Higher IOPS
  - Better throughput inside SSD
    - ✓ Better parallelism and less interference among requests
  - Less latch contention for hot pages (Contention Split @ CIDR '21)
  - Less space amplification (Bohyun's TPC-C exp.)



Random IOPS	Page Size		
	16KB	8KB	4KB
Read-only (128 threads)	29,870	57,847	89,083
Write-only (1-fsync)	196	206	225
Write-only (256-fsync)	4,563	7,978	12,647
Write-only (128 no-barrier)	13,446	25,546	49,009

(a) *DuraSSD*

Random IOPS	Page Size		
	16KB	8KB	4KB
Read-only (128 threads)	516	528	538
Write-only (128 threads)	428	439	444

(b) Harddisk (Seagate Cheetah 15K.6 146.8GB)

Table 2: Effect of page size on IOPS

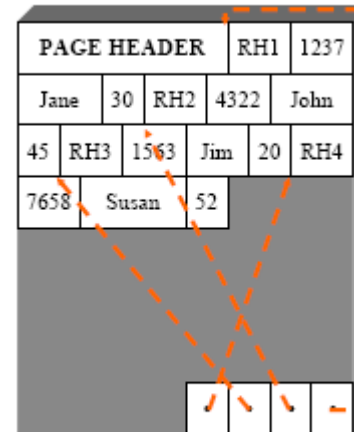
Source: Durable Write Cache in Flash Memory SSD for Relational and NoSQL Databases (SIGMOD 2014)

# “Classical” Data Layout on Disk Page: Row-Store

## □ NSM (n-ary Storage Model, or Slotted Pages)

**R**

RID	SSN	Name	Age
1	1237	Jane	30
2	4322	John	45
3	1563	Jim	20
4	7658	Susan	52
5	2534	Leon	43
6	8791	Dan	37



Records are stored sequentially  
Attributes of a record are stored together

# Other Layouts: C-Store and Pax

- C-Store(DSM) and Pax(Hybrid)

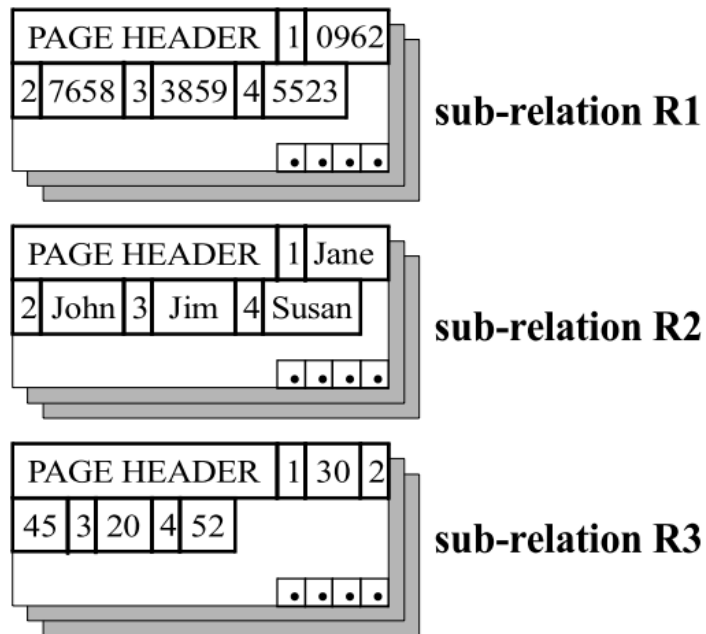


FIGURE 2: The Decomposition Storage Model (DSM).

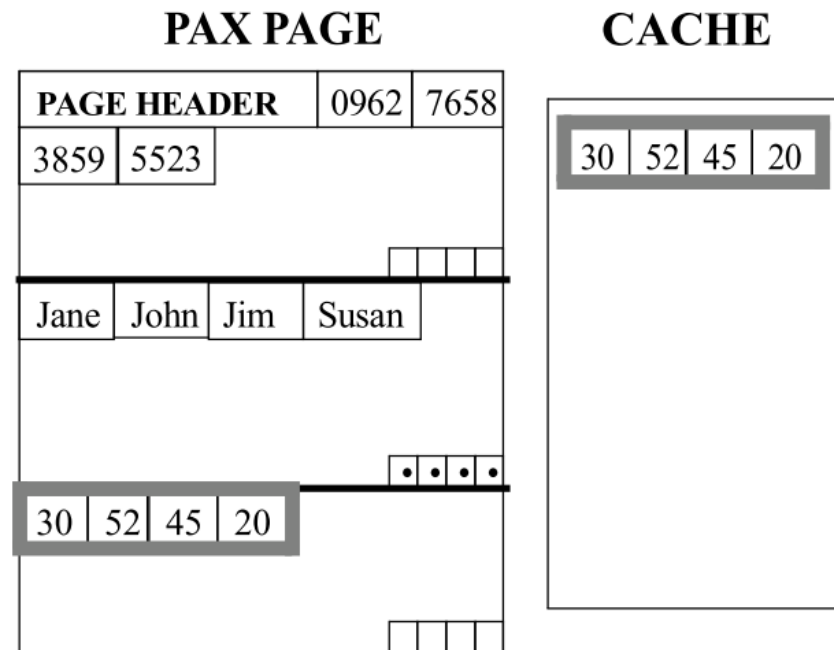
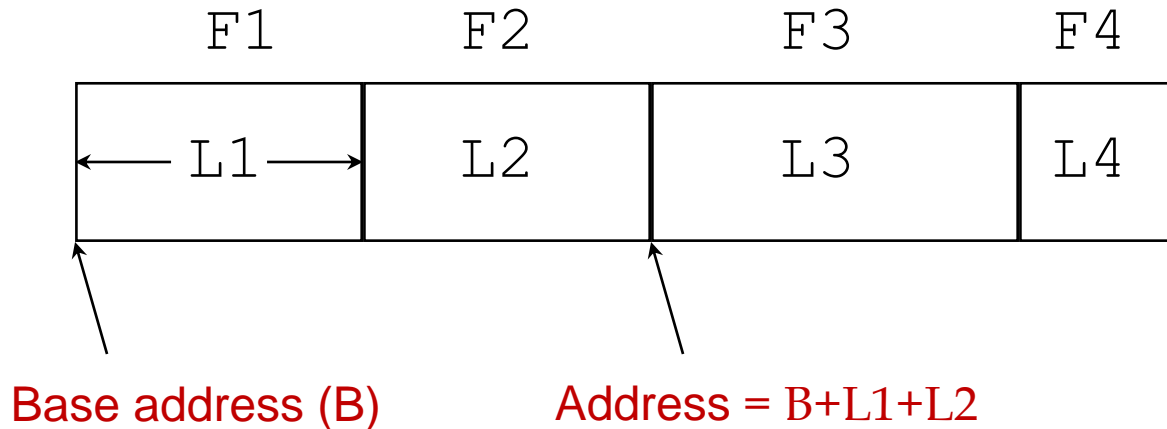


FIGURE 3: The cache behavior of PAX.

[Source: Weaving Relations for Cache Performance, Ailamaki et. al., VLDB 2001]

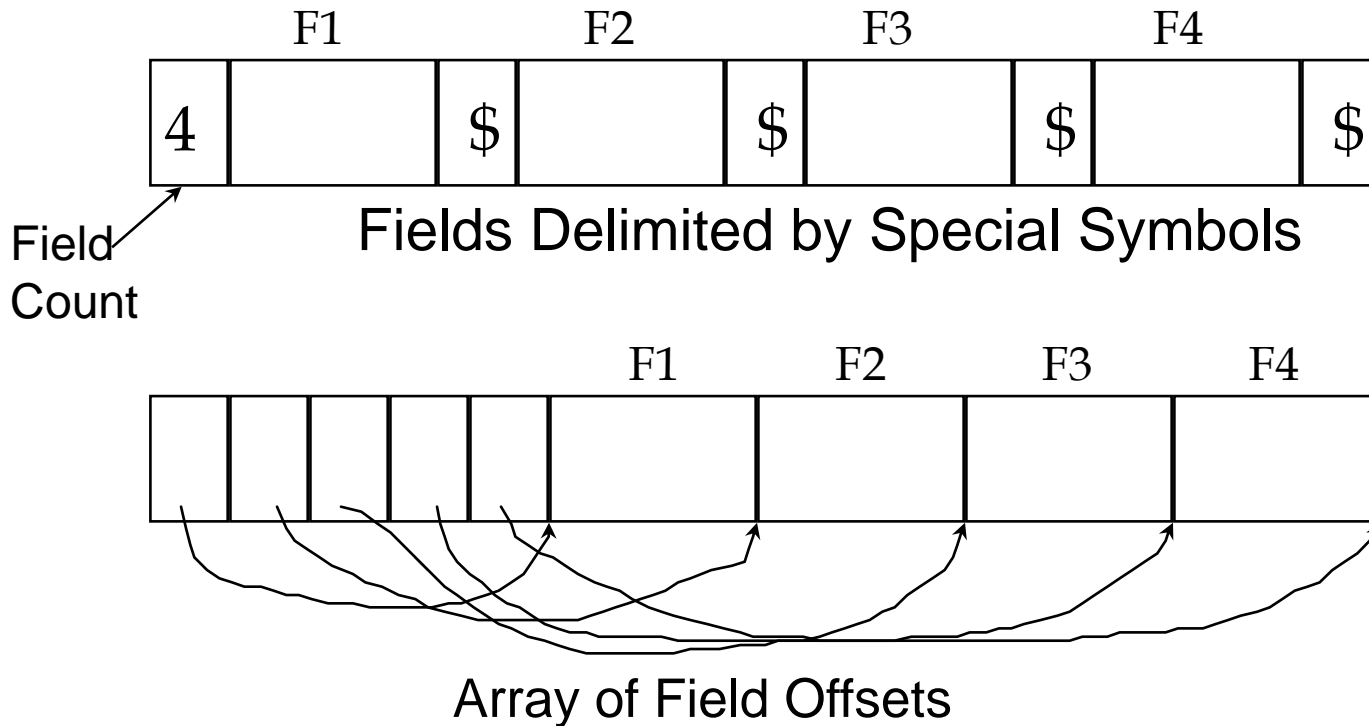
## 9.7 Record Formats: Fixed Length

- Information about field types same for all records in a file; stored in system catalogs.



# Record Formats: Variable Length

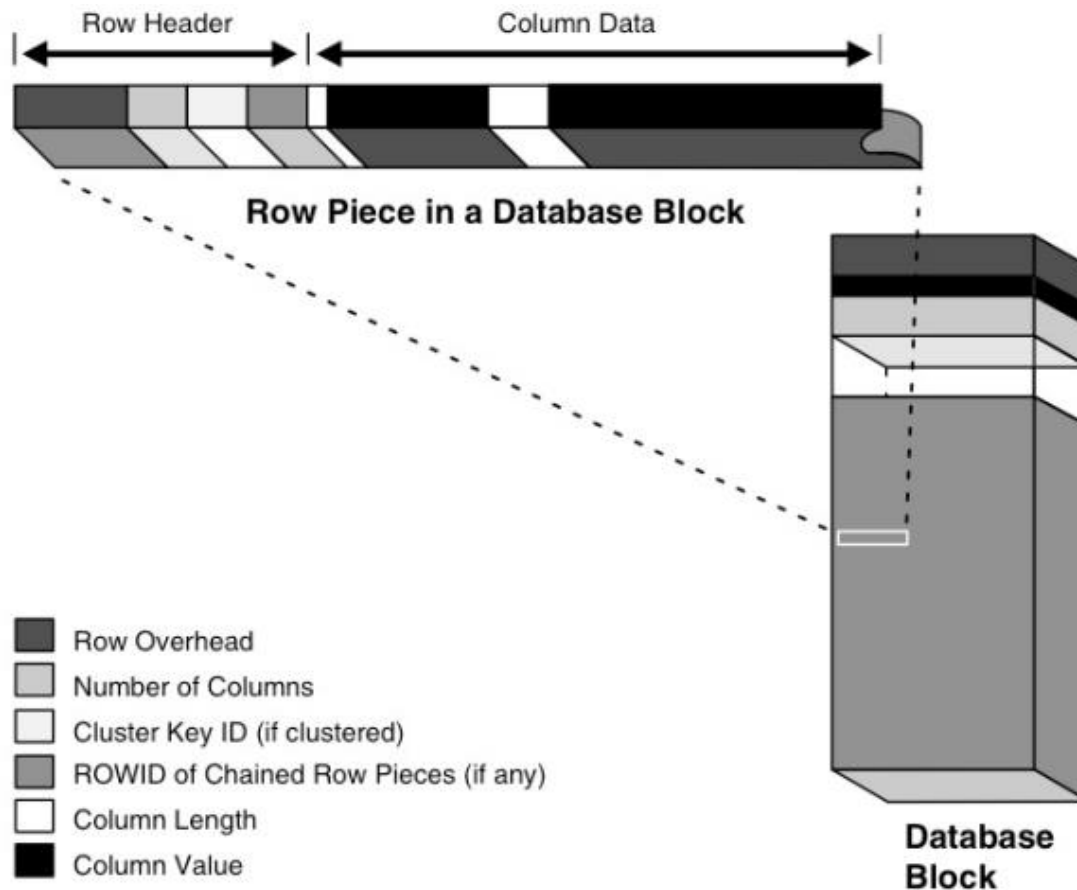
- Two alternative formats (# fields is fixed):



- Second offers direct access to i'th field, efficient storage of [nulls](#);
  - small directory overhead.



# Row Layout in Oracle



"Nulls are stored in the database if they fall between columns with data values. In these cases they require 1 byte to store the length of the column (zero).

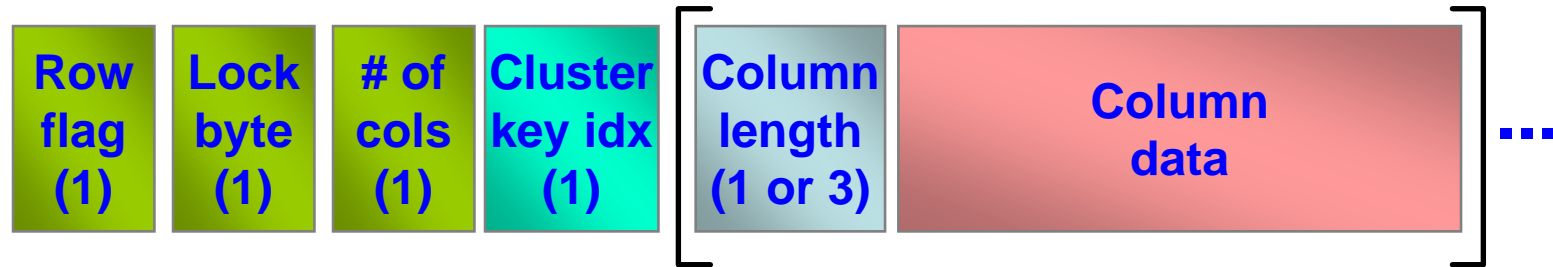
...

To conserve space, a null in a column only stores the column length (zero).

Oracle does not store data for the null column. Also, for trailing null columns, Oracle does not even store the column length."

[Oracle 10g Database Concepts, Chap. 5.4]

# Row Layout in Oracle (2)



Row overhead

tab 0, row 0, @0x1f62

tl: 38 fb: --H-FL-- lb: 0x0 cc: 8

col 0: [ 3] c2 4a 46

col 1: [ 5] 53 4d 49 54 48

col 2: [ 5] 43 4c 45 52 4b

col 3: [ 3] c2 50 03

col 4: [ 7] 77 b4 0c 11 01 01 01

col 5: [ 2] c2 09

col 6: \*NULL\*

col 7: [ 2] c1 15

0xFE  
(1)

Length field  
(2)

0xFF  
(NULL)

**fb: row migration**

-normal case: H FL

-Migration case

- original row: H

- migration row: FL

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
7369	SMITH	CLERK	7902	80/12/17	800		20

# Column Value in Oracle

===== SQL DUMP function =====

// dump() shows how a data value is stored in disk

```
select dump(a)
from test
where a = 1000;
```

DUMP(A)

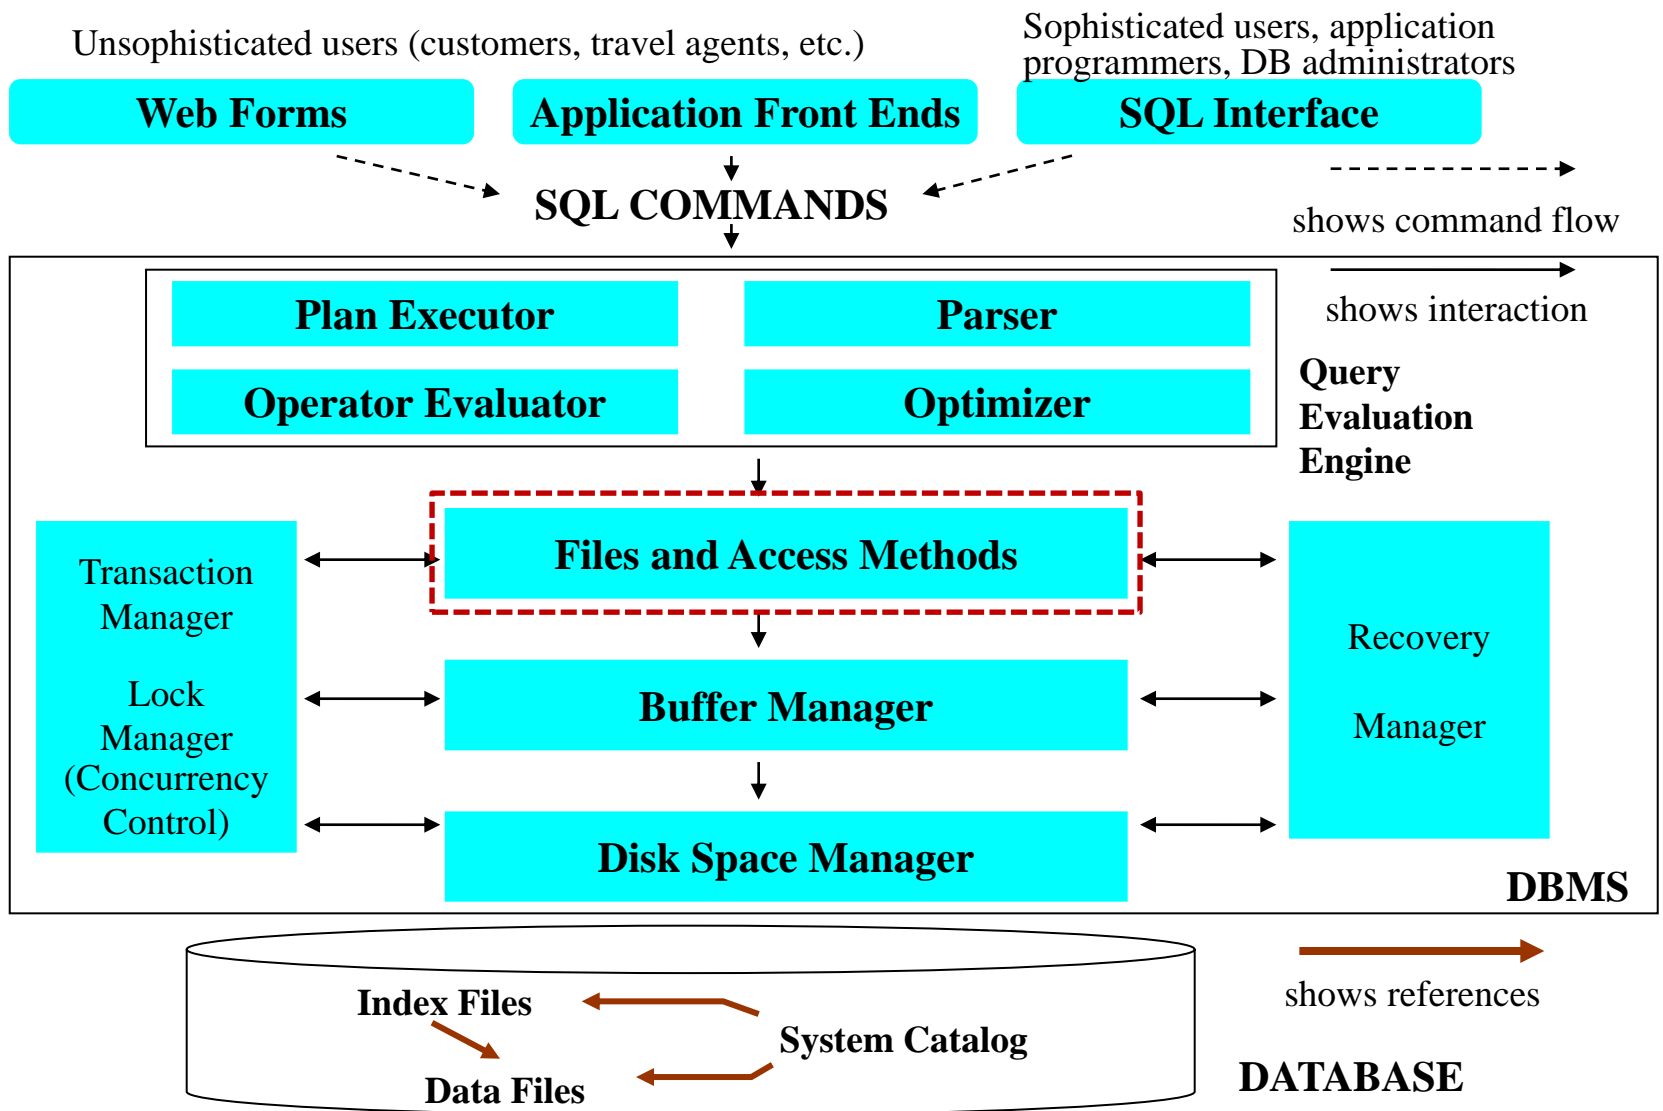
-----  
Typ=2 Len=2: 194,11

← Integer value 1000: type 2: int, column length 2 byte  
← data value 1000 is represented as (194, 11)

# System Catalogs(or Data Dictionary)

- For each relation:
  - name, file name, file structure (e.g., Heap file)
  - attribute name and type, for each attribute
  - index name, for each index
  - integrity constraints
- For each index:
  - structure (e.g., B+ tree) and search key fields
- For each view:
  - view name and definition
- + statistics, authorization, buffer pool size, etc.
- See [here](https://docs.oracle.com/cd/B13789_01/server.101/b10755/toc.htm) for Oracle's Catalog ([https://docs.oracle.com/cd/B13789\\_01/server.101/b10755/toc.htm](https://docs.oracle.com/cd/B13789_01/server.101/b10755/toc.htm))

**Catalogs are themselves stored as relations!!**



**Figure 1.3 Anatomy of an RDBMS**

# Full Table Scan: An Access Method

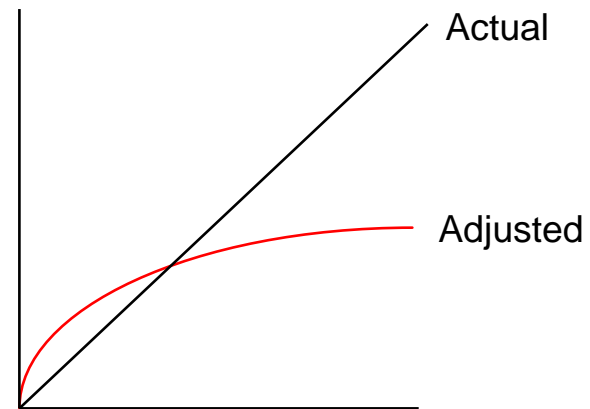
```
SELECT SUM(a)  
FROM test;
```

- To process the above query, we should scan all the blocks of the 'TEST' table
  - **FULL TABLE SCAN**: the only access method for the query
  - Access method: data structure and algorithms for organizing and accessing data (e.g. heap file, index)
- # of blocks = 1000
- How many disk I/Os do you guess are necessary?
  - **Cost model!!**

# Cost Model of Full Table Scan

- Full table scan
  - Cost = # of Blocks / Adjusted *db\_file\_multiblock\_read\_count* + 1 (in Oracle 9.2)
  - *db\_file\_multiblock\_read\_count*
  - Why adjust? *memory caching!*

Actual	Adjusted
4	4.175
8	6.589
16	10.398
32	16.409
64	25.895
128	40.865



(How the CBO works by Lewis, Jonathan, 2003  
[http://www.nocoug.org/download/2003-08/how\\_cbo\\_works.ppt](http://www.nocoug.org/download/2003-08/how_cbo_works.ppt))

# Estimated Cost of the Sample Query

- $\text{Cost} = 1000 / 10.4 + 1 = 98$
- Query optimizer estimates that the given query will **incurs** 98 disk I/Os
- What about the real disk I/O #?
- NOTE: **estimated** vs. **real** cost



# OLAP vs. OLTP

- On-Line { Analytical vs. Transactional} Processing

```
SQL> SELECT SUM(b) FROM TEST;  
SUM(B)
```

```
5.0000E+11
```

```
Execution Plan
```

```
-----  
| Id | Operation          | Name | Rows | Bytes | Cost (%CPU)| Time     |  
-----  
| 0 | SELECT STATEMENT    |      |    1 |    5 | 22053  (1)| 00:04:25 |  
| 1 | SORT AGGREGATE      |      |    1 |    5 |           |          |  
| 2 | TABLE ACCESS FULL | TEST | 996K | 4865K | 22053  (1)| 00:04:25 |  
-----
```

```
Statistics
```

```
179 recursive calls
```

```
0 db block gets
```

```
100152 consistent gets
```

```
100112 physical reads
```

```
.....
```

```
1 rows processed
```

Data  
Page 1

Data  
Page 2

⋮

Data  
Page i

⋮

Data  
Page  
100,000

TEST SEGMENT  
(Heap File)

# OLAP vs. OLTP

- Point or Range Query

```
SQL> SELECT B FROM TEST WHERE A = 500000;  
B
```

500000

Execution Plan

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	5	22053 (1)	00:04:25
1	SORT AGGREGATE		1	5		
2	TABLE ACCESS FULL	TEST	996K	4865K	22053 (1)	00:04:25

Statistics

179 recursive calls

0 db block gets

100152 consistent gets

100112 physical reads

.....

1 rows processed

Data  
Page 1

Data  
Page 2

⋮

Data  
Page i

⋮

Data  
Page  
100,000

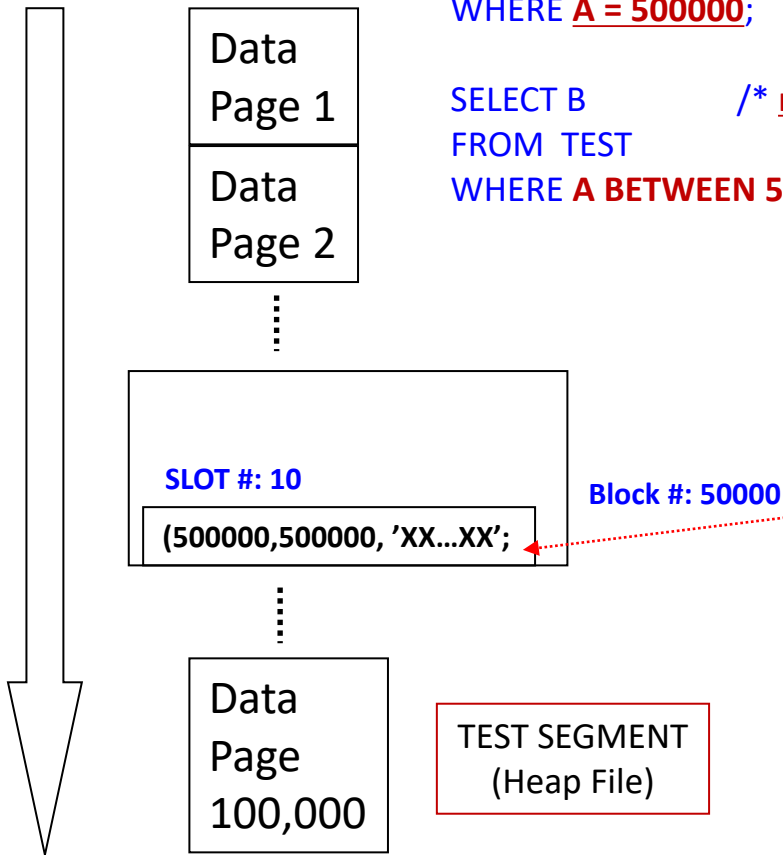
TEST SEGMENT  
(Heap File)

# OLAP vs. OLTP

```
CREATE INDEX TEST_A ON TEST(A);
```

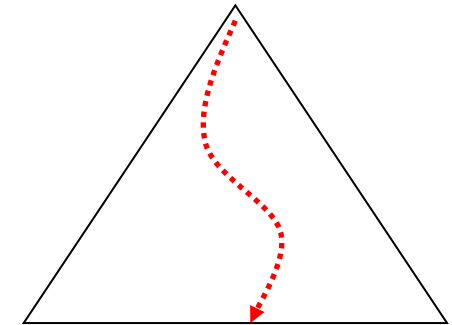
```
SELECT B          /* point query */  
FROM TEST  
WHERE A = 500000;
```

```
SELECT B          /* range query */  
FROM TEST  
WHERE A BETWEEN 50001 and 50101;
```



B-tree Index on TEST(A)  
(chapter 10.)

SEARCH KEY: 500000



(500000, (50000, 10))

## Cost:

- Full Table Scan: 100,000 Block Accesses
- Index: (3~4) + Data Block Access
  - Point query: 1
  - Range queries: depending on range

# OLAP vs. OLTP

- Index-based Table Access

```
SQL> SELECT B FROM TEST WHERE A = 500000;
```

```
B
```

```
-----  
500000
```

```
Execution Plan
```

```
-----  
| Id | Operation                | Name   | Rows | Bytes | Cost (%CPU)| Time     |  
-----  
| 0 | SELECT STATEMENT          |        |    1 |    10 |    4 (0)| 00:00:01 |  
| 1 | TABLE ACCESS BY INDEX ROWID| TEST   |    1 |    10 |    4 (0)| 00:00:01 |  
|* 2 | INDEX RANGE SCAN          | TEST_A |    1 |      |    3 (0)| 00:00:01 |  
-----
```

```
Statistics
```

```
-----  
.....  
      5 consistent gets  
      4 physical reads  
.....  
      1 rows processed
```

# Summary

- **Disks** provide cheap, non-volatile storage.
  - Random access, but cost depends on location of page on disk; important to arrange data sequentially to minimize *seek* and *rotation* delays.
- **Buffer manager** brings pages into RAM.
  - Page stays in RAM until released by requestor.
  - Written to disk when frame chosen for replacement (which is sometime after requestor releases the page).
  - Choice of frame to replace based on *replacement policy*.
  - Tries to *pre-fetch* several pages at a time.
- **DBMS vs. OS File Support**
  - **DBMS needs features not found in many OS's**, e.g., forcing a page to disk, controlling the order of page writes to disk, files spanning disks, ability to control pre-fetching and page replacement policy based on predictable access patterns, etc.

## Summary (Contd.)

- **Variable** length record format with field offset directory offers support for direct access to i'th field and null values.
- **Slotted page format** supports variable length records and allows records to move on page
- **File layer** keeps track of pages in a file, and supports abstraction of a collection of records.
  - Pages with free space identified using linked list or directory structure (similar to how pages in file are kept track of).
- **Catalog relations** store information about the various database objects including relations, indexes, and views.

*“If you want truly to understand a system,  
try to change it”*

*Kurt Lewin (A Psychologist)*