

Ch 18. Crash Recovery

Sang-Won Lee

<http://icc.skku.ac.kr/~swlee>

SKKU VLDB Lab.

(<http://vldb.skku.ac.kr/>)

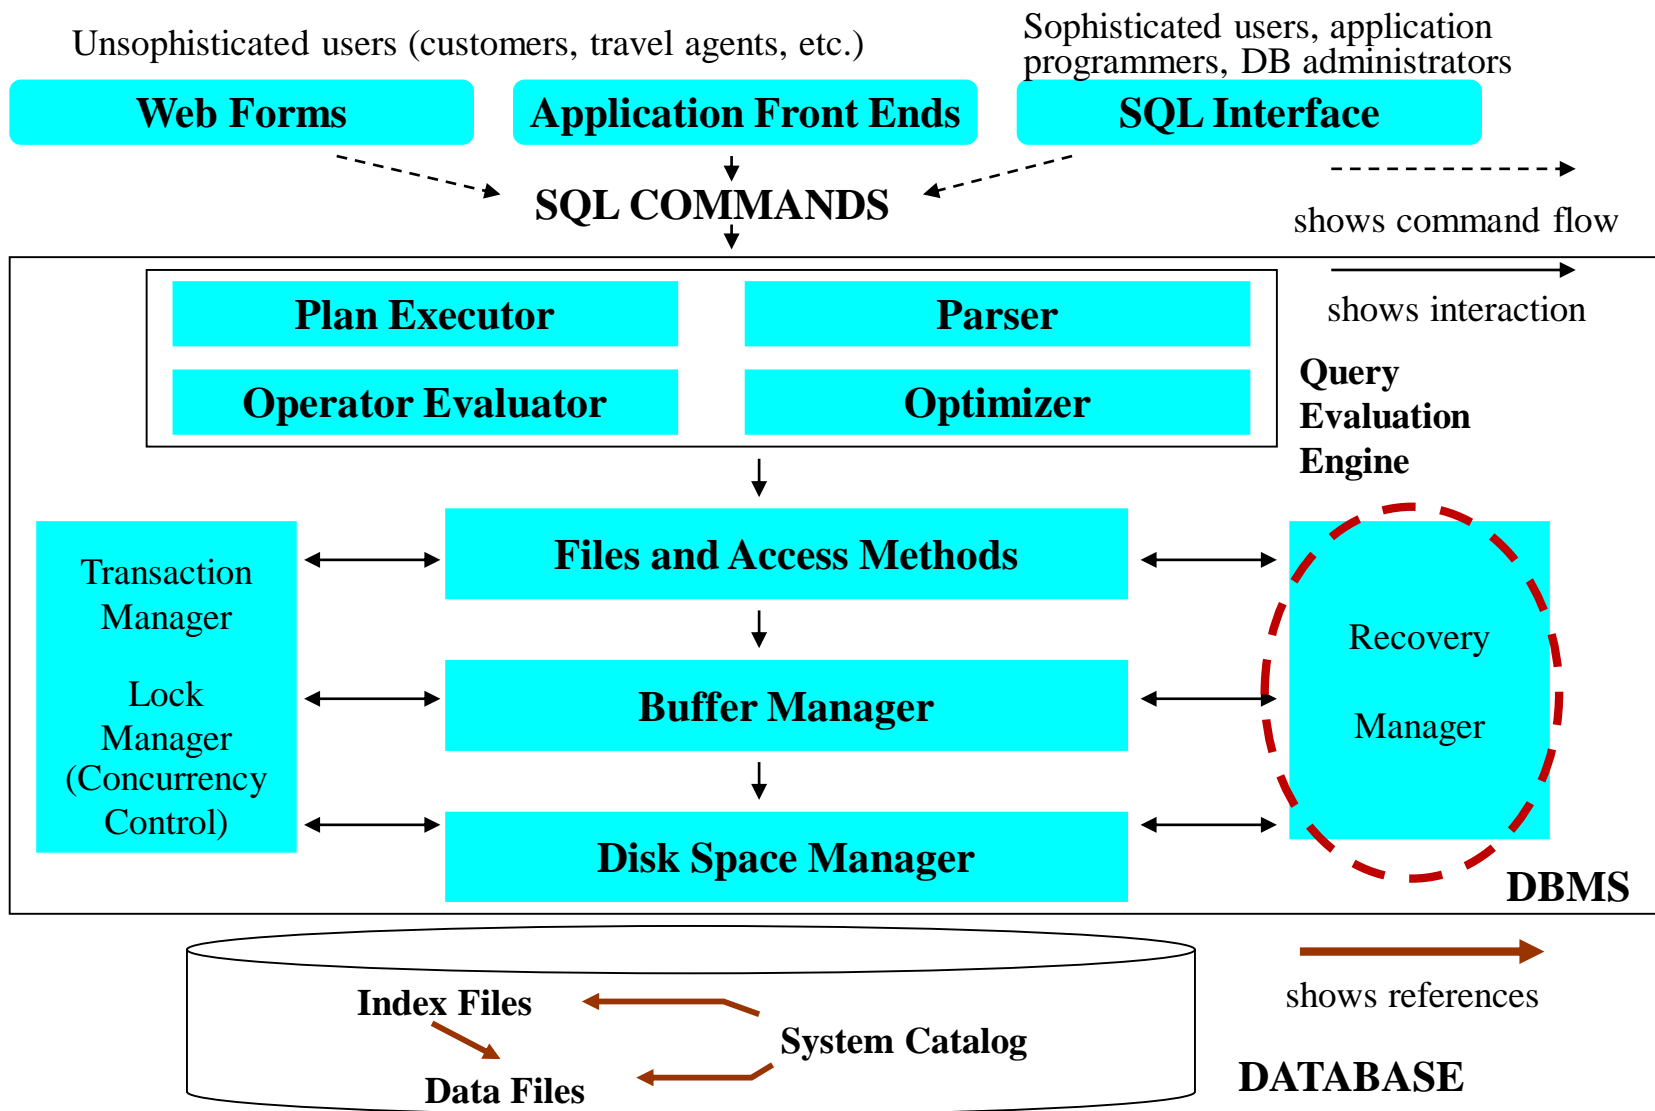


Figure 1.3 Anatomy of an RDBMS


The ACID Properties

- **Atomicity**: All actions in the Xact happen, or none happen.
- **Consistency**: If each Xact is consistent, and the DB starts consistent, it ends up consistent.
- **Isolation**: Execution of one Xact is isolated from that of other Xacts.
- **Durability**: If a Xact commits, its effects persist.

The **Recovery Manager** guarantees Atomicity & Durability.

Example: Transactional Atomicity and Durability

- Money transfer
 - BeginTX; $A = A - 100$; $B = B + 100$; Commit (or Rollback or Failure)
 - BeginTX; Read A; Update $A = A - 100$; Read B; Update $B = B + 100$; ..



Various Failures and Recovery

```
create table account (id number, balance number);  
insert into account values (1, 1000);  
insert into account values (2, 2000);  
commit;
```

```
Begin_tx;  
update account set balance = balance - 100 where id = 1;  
update account set balance = balance + 100 where id = 2;  
commit; -- or rollback;
```

Buffer
Cache

Pi [1, 1000]

Pj [2, 2000]

Storage

Pi [1, 1000]

Pj [2, 2000]

- NOTE that system can fail **at any point** of TX execution
- **Various DB statuses** upon failure time
 - ✓ What if rollback after Pi is updated? -> Undo
 - ✓ Both Pi and Pj are not written back to the storage before system crash → Redo
 - ✓ Updated Pi is already written to the storage before commit and then system crashes; → Undo
 - ✓ Updated pages at DRAM after commit -> Redo
 - ✓ Non-atomic page write
 - ✓
- Database recovery manager has to be able to cope with all the above scenarios.

Failures and Recovery


- Transaction failure
 - Rollback or killed by deadlock → undo recovery
- System failure
 - OS, DBMS bug, Power shortage → redo/undo recovery
 - Atomic write of page(s) is not guaranteed
- Media failure – not the main topic in this slide
 - Disk head crash → archive recovery
- “A very reliable (and available) system with unreliable component?”
 - [[John Von Neumann](#): stored-program, CPU + RAM + Disk]
 - THE GIANT in database field: [Jim Gray](#)
- Replication (though we do not cover) - <https://www.khan.co.kr/culture/culture-general/article/202108090600001>



Motivation

- **A**tomicity:
 - Transactions may abort (“Rollback”).
- **D**urability:
 - What if DBMS stops running? (Causes?)

- Desired behavior after system restarts:
 - T1, T2 & T3 should be durable. - **Durability**
 - T4 & T5 should be aborted (effects not seen). - **Atomicity**




18.0 Basics



Disk Writes

Updates@Transactions
: SQL Insert/Delete/Update




- Three assumptions:
 1. Page: unit of update propagation from host DRAM to storage
 2. In-place update vs. out-of-place update (or Copy-on-write: CoW)
 3. Atomic propagation: single page, multiple pages
 - ✓ How? e.g. shadow page
- Technical issues:
 - Buffer replacement
 - Commit, abort, crash, recovery, redo/undo

Transaction – Commit and Force


Updates@Transactions
: SQL Insert/Delete/Update

BUFFER POOL



MAIN MEMORY (Volatile)

DISK (Non- Volatile)



Transaction T


- Begin
- Read page P_i , P_j , P_k (from disk)
- Update P_i , P_k
- Commit (Write P_i , P_k to disk)

→ Achieve “All” and Durability

→ Commit policy: “Force at Commit”

Transaction – Rollbacks and **No-Steal**

Updates@Transactions
: SQL Insert/Delete/Update



Transaction T

- Begin
- Read page P_i (from disk)
- Update P_i
- *[Never write P_i to disk before commit]*
- Rollback

➔ Achieve “Nothing”

➔ Buffer management: “No Steal”

Recovery Modes vs. Buffer Policies

- **Force** every write to disk?
 - Poor response time.
 - But, provides durability.
- **Steal** buffer frames from uncommitted Xacts?
 - If not, poor throughput.
 - If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial	
No Force		Desired

No-Steal and Force – Economical Solution?

- The economics of “volatile RAM and non-volatile disk”
 - No-steal policy is impossible because of limited RAM size
 - Force policy is inefficient since many random writes to disk reduce the throughput
- Therefore, all major DBMSs choose “steal + no-force” policy
 - Steal policy as buffer management
 - No-force policy for commit




More on Steal and Force

- **STEAL** (why enforcing Atomicity is hard)
 - *To steal frame F*: Current page in F (say P) is written to disk; some Xact holds lock on P.
 - ✓ What if the Xact with the lock on P aborts?
 - ✓ Must remember the old value of P at steal time (to support **UNDO**ing the write to page P). (modern file systems lacks UNDO)
- **NO FORCE** (why enforcing Durability is hard)
 - What if system crashes before a modified page is written to disk?
 - Write as little as possible, in a convenient place at commit time, to support **REDO**ing modifications.



Transaction – Commit and No-Force

Updates@Transactions
: SQL Insert/Delete/Update




Transaction T

- Begin
- Read page P_i (from disk)
- Update P_i
- Commit
 - *[Do not force P_i to disk]*

➔ In case of system failure, how to achieve “all” and durability?

➔ “Redo log”

Transaction – Rollbacks and Steal



Transaction T

- Begin
 - Read page P_i (from disk)
 - Update P_i
 - [prior to commit, P_i can *reach disk* due to buffer replacement]
 -
 - Rollback (or system failure)
- ➔ In case of system failure, how to achieve “nothing”?
- ➔ “Undo log”

Recovery Modes

- **Force** every write to disk?
 - Poor response time.
 - But provides durability.

Force	Trivial	
No Force		Desired
	No Steal	Steal

- **Steal** buffer-pool frames from uncommitted Xacts?
 - If not, poor throughput.
 - If so, how can we ensure atomicity?

SQLite: Case Study

- Single user, mobile device: very **simple** approach for recovery
- SQLite demo
- Two journal modes for recovery
 - Rollback (RBJ) and WAL journal: a kind of **side file**
 - ✓ Simple and seemingly inefficient: redundant write or journaling
 - **But**, enable 1) atomic page write, 2) atomic propagation of N pages, and 3) simple UNDO;
 - And, in case of WAL, transaction-consistent checkpoint is supported
- Commit policy: **FORCE**, thus NO REDO
- Buffer policy : **STEAL**, but simple UNDO by RBJ and WAL



SQLite Case Study: Concurrency Control

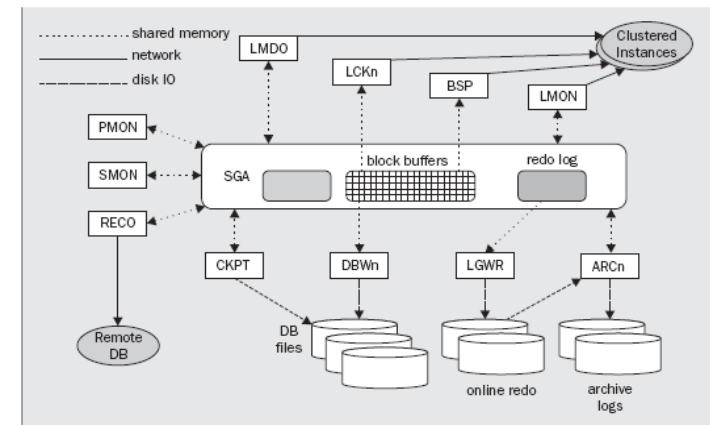
- SQLite takes coarse-grained concurrency control
 - **File-level** vs. Oracle's tuple-level fine-grained CC
- Coarse-Grained CC allows to take the force commit
 - In case of tuple-level CC like Oracle, one page can be updated concurrently by T1 (committing) and T2(in-progress)
 - THEN, even with RBJ / WAL, there is no way to undo T2's update when system crash is encountered before T2 commits.
 - WHY??

18.2 Log



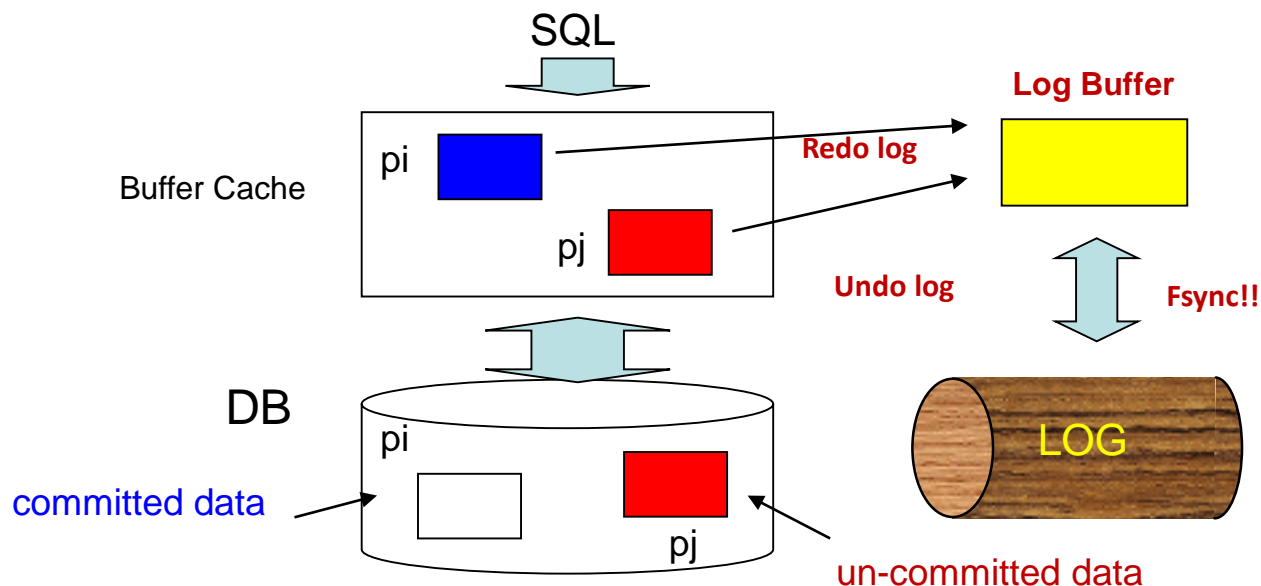
What You Should Know

- Log and LSN (Log Sequence Number)
 - Every update per page Pi will generate **its corresponding LOG** at DRAM log buffer and the LOG's LSN be stored as new **pageLSN** of Pi
- Interplay among buffer manager, database writer, log manager, and recovery manager
- Log-based UNDO/REDO recovery
 - Force-log-at-commit for **redo recovery**
 - Write-Ahead Log (WAL) protocol for **undo recovery**




Basic Idea: Logging

- Update-in-place + Steal + No-force → {redo + undo} log
 - “Page as unit of update propagation” makes recovery harder
- Log and LSN (Log Sequence Number)
 - Every update per page P_i will generate **its corresponding LOG** at DRAM log buffer and the LOG's LSN be stored as new **pageLSN** of P_i




Basic Idea: Logging (cont.)



- For every update, store REDO and UNDO info. in a *log*.
 - Sequential writes to log (put it on a separate disk).
 - **Minimal info** (diff) written to log, so multiple updates fit in a single log page.
- Log: An “ordered” list of REDO/UNDO actions
 - Log record contains:
<XID, pageID, offset, length, old data, new data>
 - and additional control info (which we’ll see soon).

Update and Log





- Page changes - **Insert**, **Update**, Delete
- Update log
 - E.g. (**LSN**, Tid, Pi, Offset, Length, **Before-value**, **After-value**)
 - ✓ **LSN** = Log Sequence Number

Three Types of Logs

- Physical(value) vs. Logical vs. Physio-logical log

Insert record r into table A



Table A		



Logical log record

insert, A , r

Table A		



Physiological log records

insert, A , page 508, r

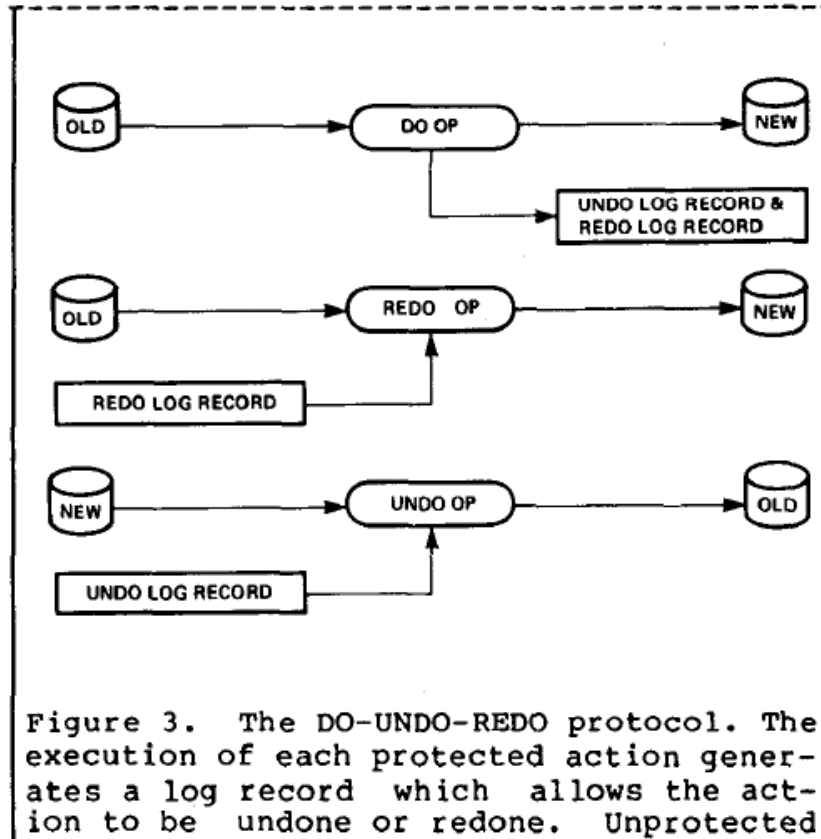
insert, B, page 72, s

insert, C, page 94, t

Log

- **Database** = Pages in disk + Log
 - Pages in DRAM: unreliable information
 - Log: a redundant information in reliable disk
- In spite of system failure,
 - Redo log guarantees durability for committed transactions, even though their updated pages are not forced out to disk
 - Undo log guarantee atomicity (i.e. “nothing”) for uncommitted or aborting transactions

DO-UNDO-REDO Model




- C.f.
 - “String” by Ariadne and Theseus (Greeks)
 - “Trail of bread crumbs” by Hensel and Gretel

From “The Transaction Concepts: Virtues and Limitations” VLDB 1981

Two Logging Protocols

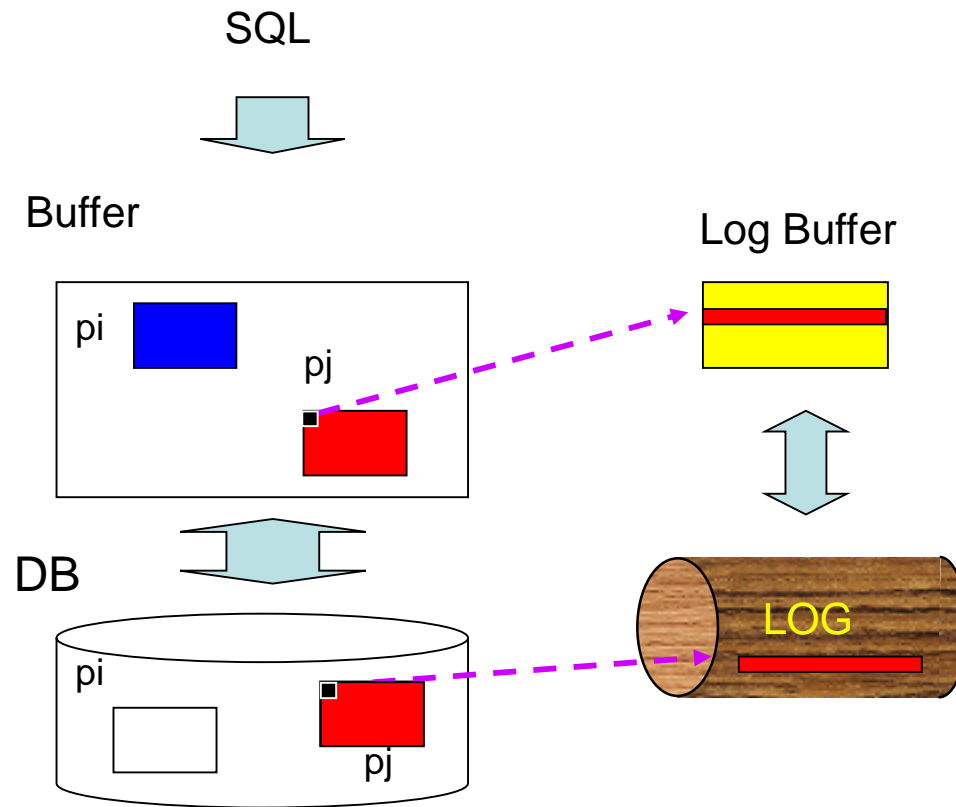
- WAL (Write-Ahead Log) protocol – for **undo**
 - Before (over)writing a page to disk, flush undo log in the log buffer
 - ✓ up to the page's pageLSN
 - DBWR calls LGWR to write the undo log (inter-process comm.) – in fact, the overhead (comm. and disk IO) is not big
- Log-Force-at-Commit – for **redo**
 - A transaction is truly committed only after the commit log record is persistently stored in log disk (for this, the cache in log disk is turned off)
 - Performance bottleneck
 - Group commit

Oracle Architecture



WAL Protocol (Chap 18.4)


- Before writing Pj to disk



Log Force at Commit

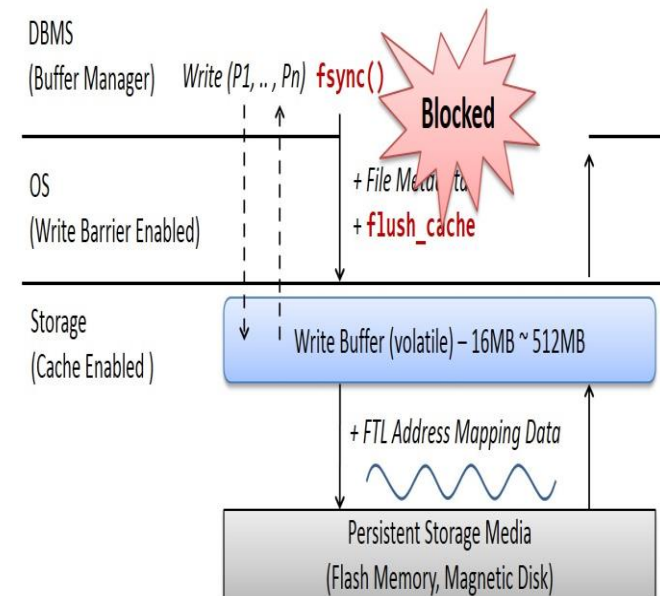
- When a transaction T_i commits

You Should Understand what happens upon normal Update SQL statements and Commit under Force-Log@Commit protocol




Write, Durability, Fsync, Write_barrier and Storage Cache

- Durability: changes made in DRAM are persistently stored in non-volatile media
- Common mechanism for D is to issue the `f(data)sync()` call
 - Flushes dirty pages from OS page cache to storage device
 - If `WRITE_BARRIER` is enabled, OS sends a `FLUSH_CACHE` command to storage device and flushes the write cache to persistent media
 - Modern storage device has its own cache
- **0-latency durability**
 - Latency for durability is a challenge for performance
 - HDD → SSD → NVRAM/NVDIMM @ Storage or Host
- What if NVRAM? Do we still need REDO/UNDO recovery? It still looks challenging even with NVRAM!
 - Given NVRAM instead of DRAM, {P1, P2} TX; If we meet crash after updating only P1 in place in NVRAM, can we undo it?



18.5 Checkpointing

- Due to no-force policy, many pages dirtified by already committed transactions remain unflushed in buffer cache.



- Thanks to redo log, those pages are recoverable upon failure
- But, it takes time to recover. And, the time is money!!
- Thus, it is very critical to minimize the recovery time upon failure.
- For this reason, those pages have to be periodically checkpointed.

Checkpointing (2)

- Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. Write to log:
 - begin_checkpoint** record: Indicates when chkpt began.
 - end_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a **`fuzzy checkpoint'**:
 - ✓ Other Xacts continue to run; so these tables accurate only as of the time of the **begin_checkpoint** record.
 - ✓ No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page. (So it's a good idea to periodically flush dirty pages to disk!)
 - Store LSN of chkpt record in a safe place (*master* record).



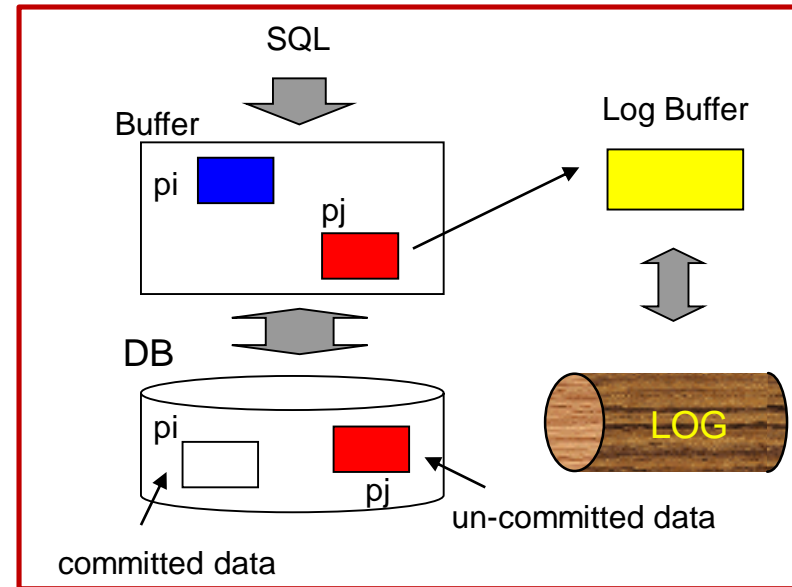
Summary – Normal-time Operations for Recovery

- WAL Protocol for Undo log
- Force-Write Log for Redo log
- Checkpointing



Recovery Summary – Disk

- Commit policy: **Force** vs. **No Force**
- Buffer mgt policy: **Steal** vs. **No Steal**
- Disk / RAM economics
 - “WAL, No Force + Steal” in normal exec.
 - “Undo/Redo” in recovery time
- SQLite vs. Enterprise-Class DBMS**



- TRADE-OFF: **Performance@NORMAL_TIME** vs. **Recovery@FAILURE**

	No Steal	Steal
No Force		Fastest
Force	Slowest	

Performance

	No Steal	Steal
No Force	REDO	UNDO REDO
Force	No-UNDO No-REDO	UNDO

Logging/Recovery

- Terribly complex
- Time-consuming

18. 1 & 18.6 ARIES Algorithm

(Algorithms for Recovery and *Isolation* Exploiting Semantics)


Golden Standard for Database Recovery



WAL protocol is known to achieves an appropriate **balance** in terms of the trade-off b/w normal-time performance and recovery time in the environment of **DRAM + harddisk.**

Research Questions: Does this assumption still hold for NVM / SSD??

Three Phases in ARIES



(Start from the last checkpoint – found via **master** record)

1. Analysis (forward sequential scan)
2. Redo (forward sequential scan)
 - “Repeat history”
3. Undo (backward sequential scan)